

Lecture 14:

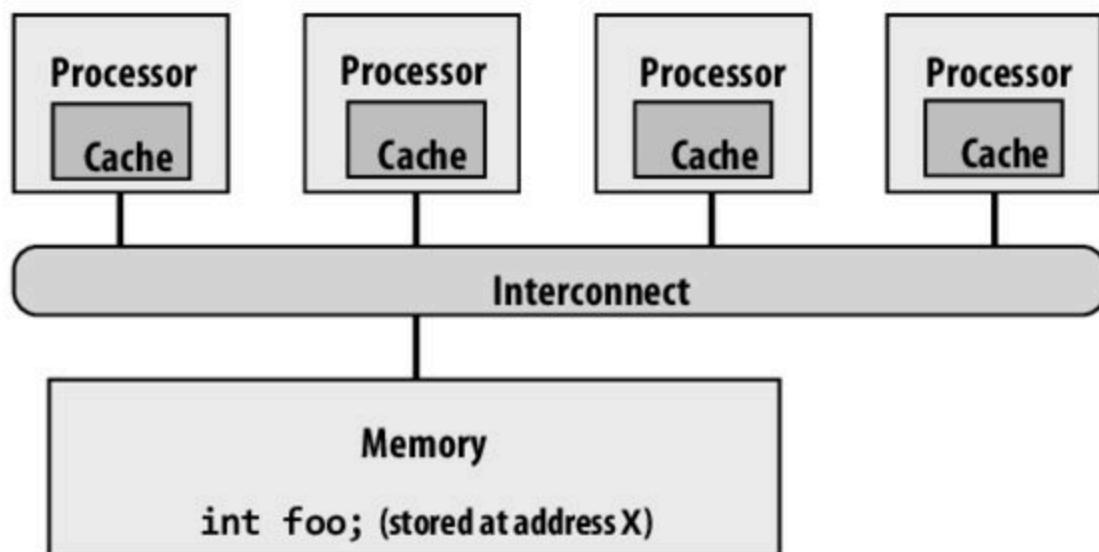
Memory Coherency and Consistency

**Parallel Computing
Stanford CS149, Fall 2024**

The cache coherence problem

Modern processors replicate contents of memory in local caches

Problem: processors can observe different values for the same memory location



Is this a mutual exclusion problem?

Can you fix the problem by adding locks to your program?

NO!

This is a problem created by replicating the data stored at address X in local caches

Action	P1 \$	P2 \$	P3 \$	P4 \$	mem[X]
					0
P1 load X	0 miss				0
P2 load X	0	0 miss			0
P1 store X	1	0			0
P3 load X	1	0	0 miss		0
P3 store X	1	0	2		0
P2 load X	1	0 hit	2		0
P1 load Y (assume this load causes eviction of X)	0	2			1

The chart at right shows the value of variable `foo` (stored at address X) in main memory and in each processor's cache

Assume the initial value stored at address X is 0

Assume write-back cache behavior

How could we fix this problem?

The memory coherence problem

Intuitive behavior for memory system: reading value at address X should return the last value written to address X by any processor.

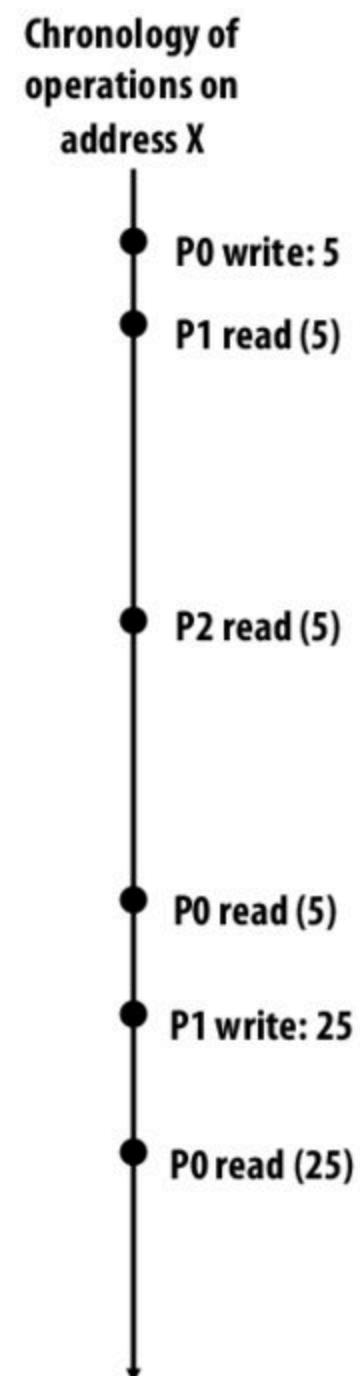
Memory coherence problem exists because there is both global storage (main memory) and per-processor local storage (processor caches) implementing the abstraction of a single shared address space.

Definition: Coherence

A memory system is coherent if:

The results of a parallel program's execution are such that for each memory location, there is a hypothetical **serial order** of all program operations (executed by all processors) to the location that is consistent with the results of execution, and:

1. Memory operations issued by any one processor occur in the order issued by the processor
2. The value returned by a read is the value written by the last write to the location... as given by the serial order



Implementation: Cache Coherence Invariants

For any memory address x, at any given time period (epoch):

Single-Writer, Multiple-Read (SWMR) Invariant

- **Read-write epoch:** there exists only a single processor that may write to x (and can also read it)
- **Read-Only- epoch:** some number of processors that may only read x

Data-Value Invariant (write serialization)

- **The value of the memory address at the start of an epoch is the same as the value of the memory location at the end of its last read-write epoch**



Implementing coherence

Software-based solutions (coarse grain: VM page)

- OS uses page-fault mechanism to propagate writes
- Can be used to implement memory coherence over clusters of workstations
- We won't discuss these solutions
- Big performance problem: false sharing (discussed later)

Hardware-based solutions (fine grain: cache line)

- "Snooping"-based coherence implementations (today)
- Directory-based coherence implementations (briefly)

Shared caches: coherence made easy

One single cache shared by all processors

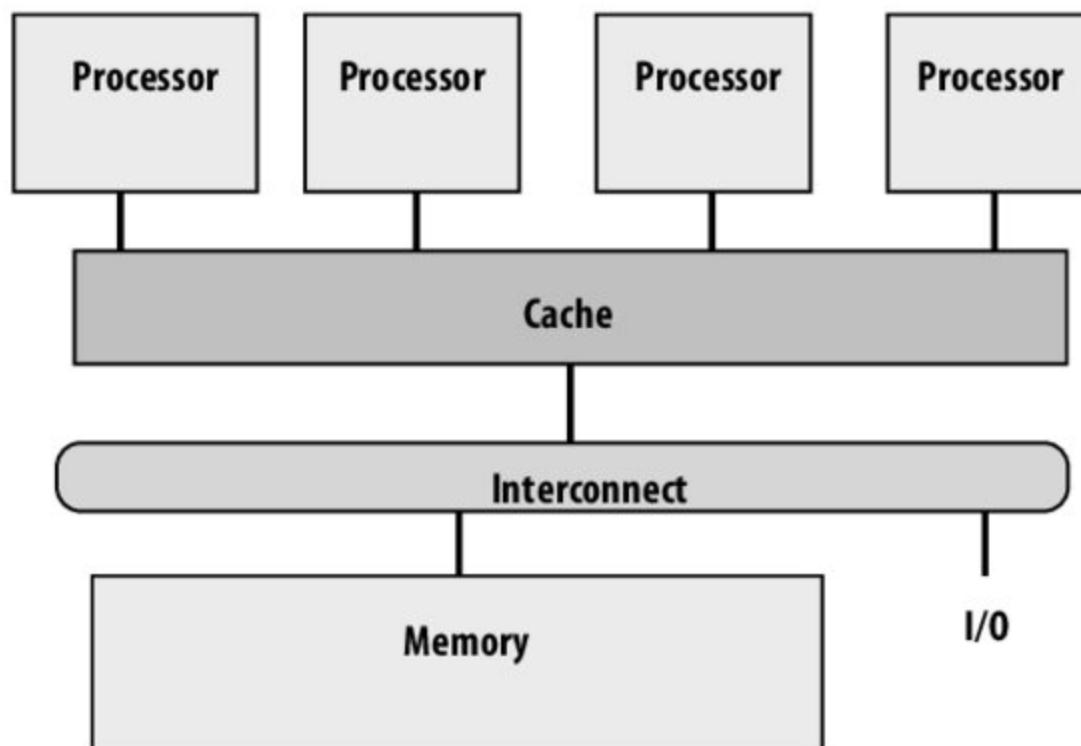
- Eliminates problem of replicating state in multiple caches

Obvious scalability problems (since the point of a cache is to be local and fast)

- Interference (conflict misses) / contention due to many clients (destructive)

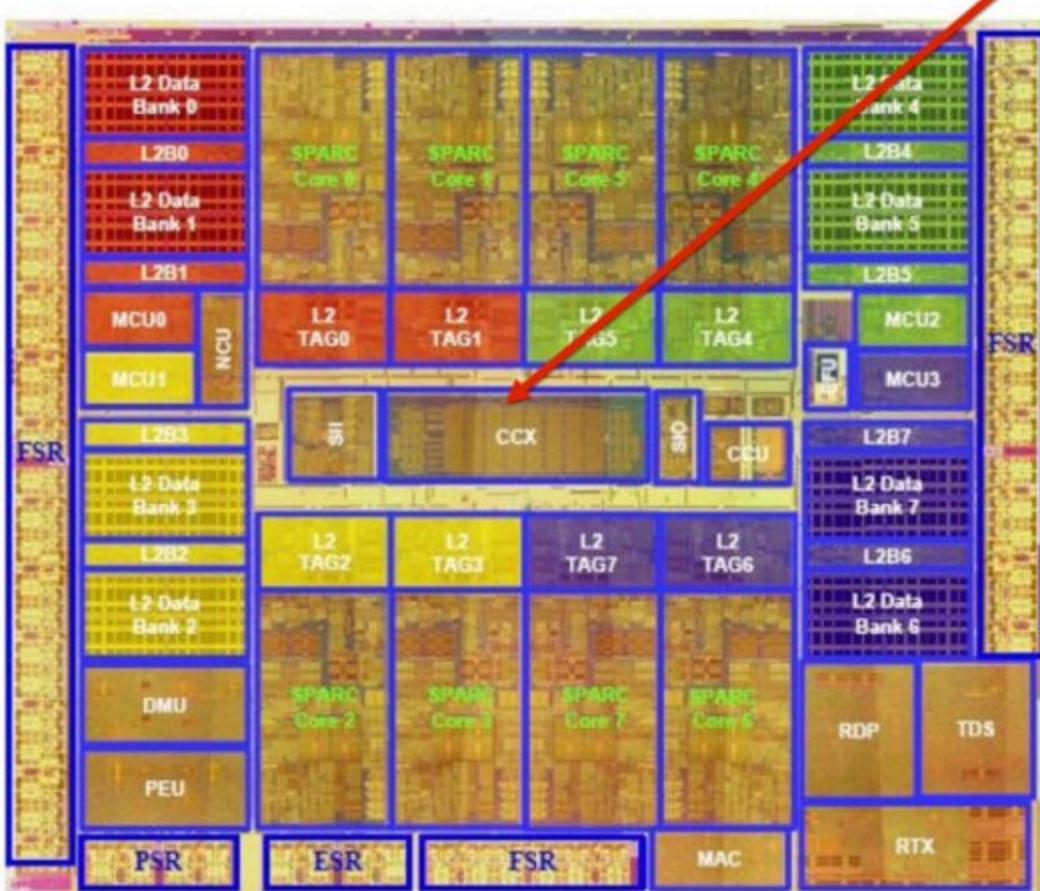
But shared caches can have benefits:

- Facilitates fine-grained sharing (overlapping working sets)
- Loads/stores by one processor might pre-fetch lines for another processor (constructive)

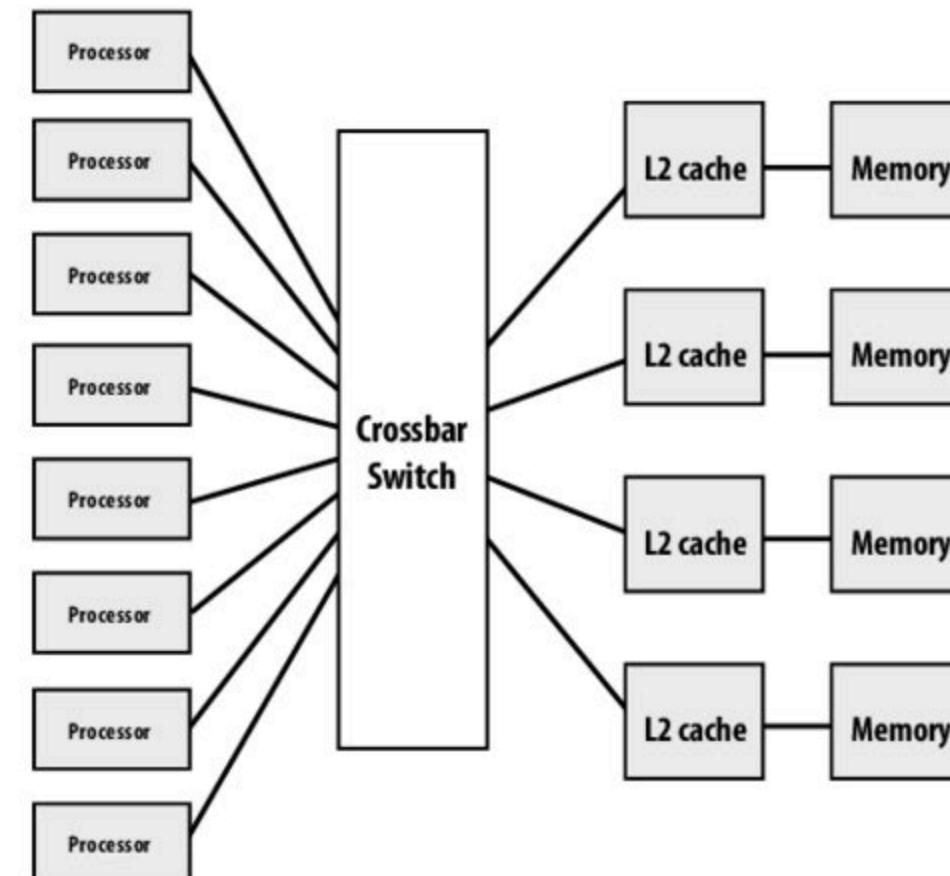


```
forall (i= 0; i++; i< N)  
    x[i] = y[i] + y[i+1] + y[i+2];
```

SUN Niagara 2 (UltraSPARC T2)



Note area of crossbar (CCX):
about same area as one core on chip



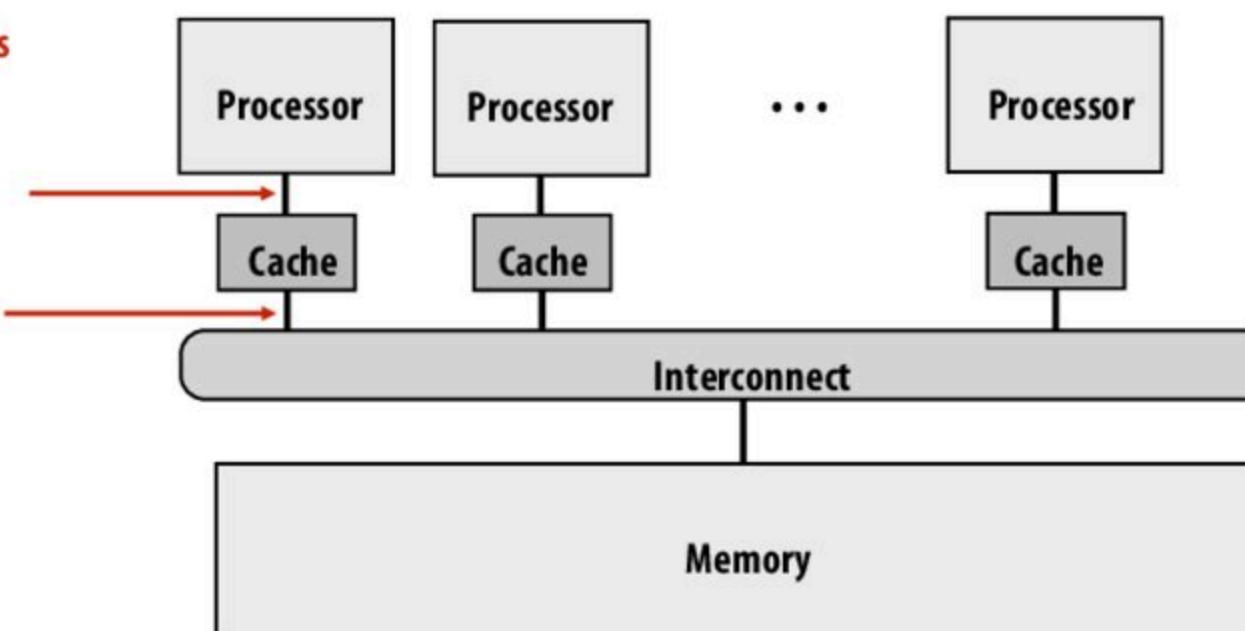
Snooping cache-coherence schemes

Main idea: all coherence-related activity is broadcast to all processors in the system
(more specifically: to the processor's cache controllers)

Cache controllers monitor (“they snoop”) memory operations, and follow **cache coherence protocol** to maintain memory coherence

Notice: now cache controller must respond to actions from “both ends”:

1. LD/ST requests from its local processor
2. Coherence-related activity broadcast over the chip’s interconnect



Very simple coherence implementation

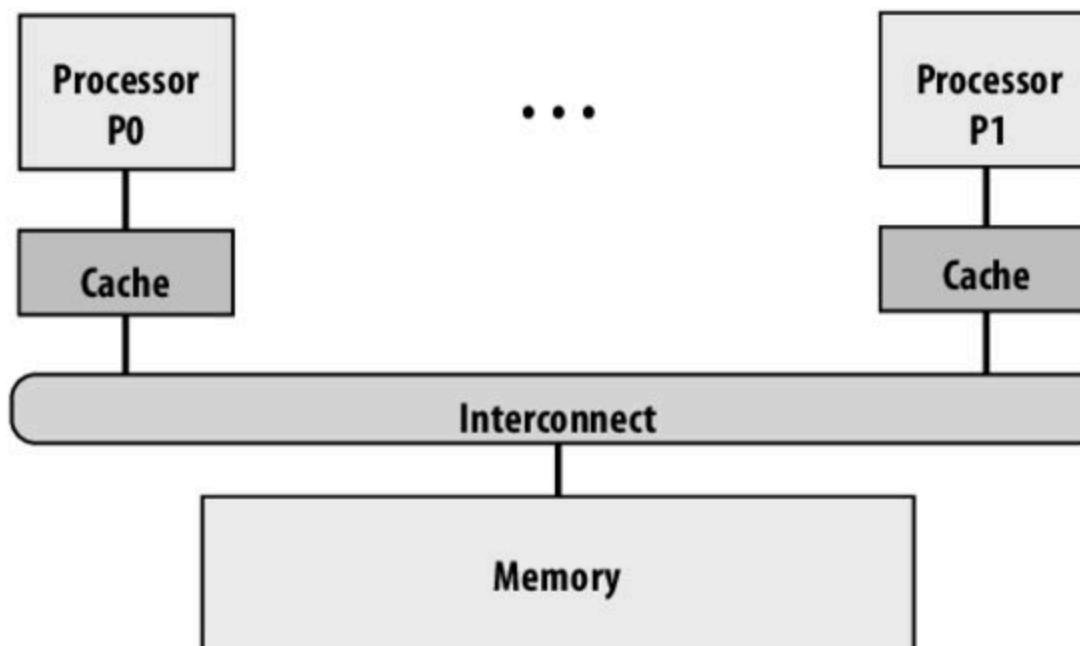
Let's assume:

1. Write-through caches
2. Granularity of coherence is cache line

Coherence Protocol:

- Upon write, cache controller broadcasts invalidation message
- As a result, the next read from other processors will trigger cache miss

(processor retrieves updated value from memory due to write-through policy)



Action	Interconnect activity	P0 \$	P1 \$	mem location X
				0
P0 load X	cache miss for X	0		0
P1 load X	cache miss for X	0	0	0
P0 write 100 to X	invalidation for X	100		100
P1 load X	cache miss for X	100	100	100

Write-through policy is inefficient

Every write operation goes out to memory

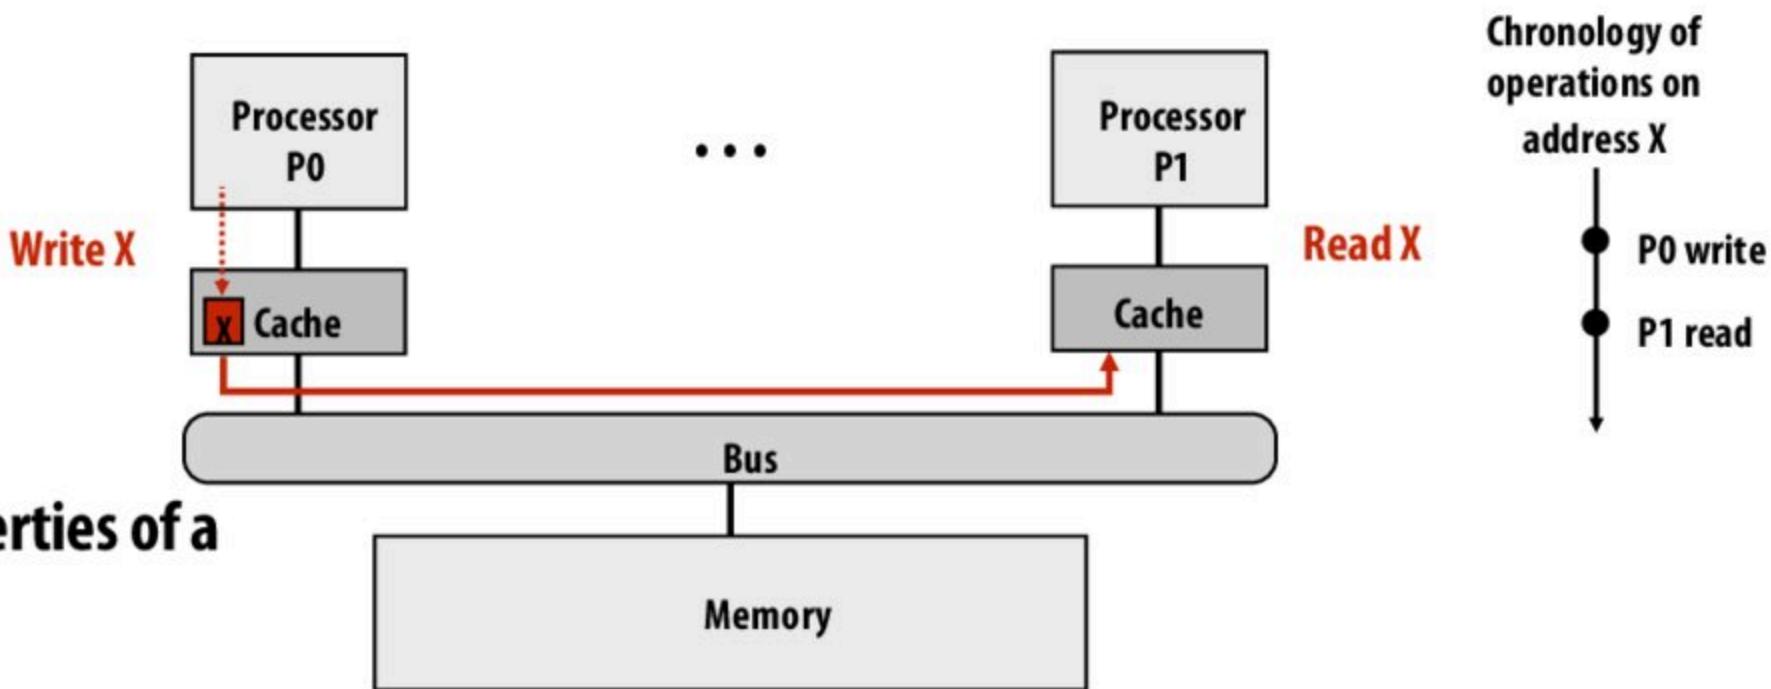
- **Very high bandwidth requirements**

Write-back caches absorb most write traffic as cache hits

- **Significantly reduces bandwidth requirements**
- **But now how do we maintain cache coherence invariants?**
- **This requires more sophisticated coherence protocols**

Cache coherence with write-back caches

What are two important properties of a bus?



Dirty state of cache line now indicates exclusive ownership (Read-Write Epoch)

- **Modified:** cache is only cache with a valid copy of line (it can safely be written to)
- **Owner:** cache is responsible for propagating information to other processors when they attempt to load it from memory (otherwise a load from another processor will get stale data from memory)

Invalidation-based write-back protocol

Key ideas:

A line in the “modified” state can be modified without notifying the other caches

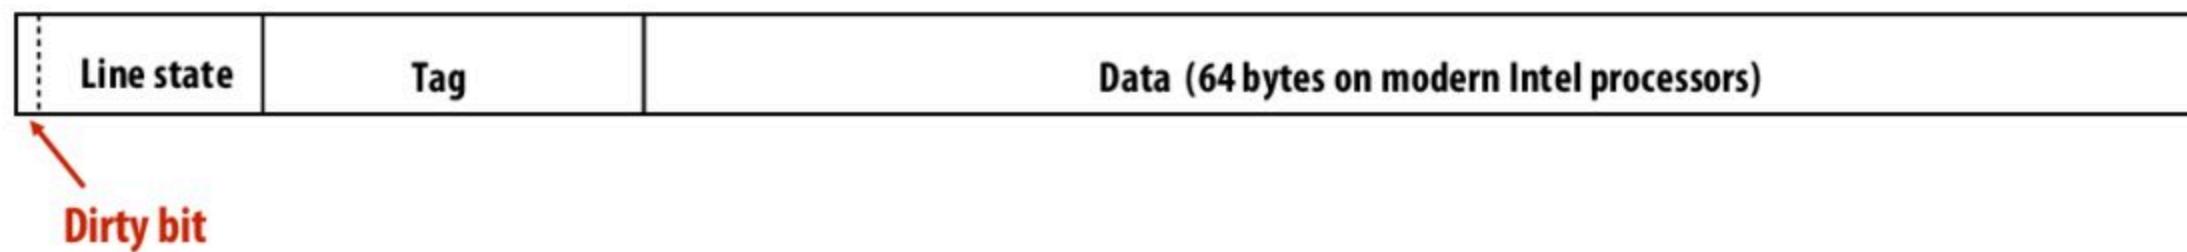
Processor can only write to lines in the modified state

- Need a way to tell other caches that processor wants exclusive access to the line
- We accomplish this by sending messages to all the other caches

When cache controller sees a request for modified access to a line it contains

- It must invalidate the line in its cache

Recall cache line state bits



MSI write-back invalidation protocol

Key tasks of protocol

- Ensuring processor obtains exclusive access for a write
- Locating most recent copy of cache line's data on cache miss

Three cache line states

- Invalid (I): same as meaning of invalid in uniprocessor cache
- Shared (S): line valid in one or more caches, memory is up to date
- Modified (M): line valid in exactly one cache (a.k.a. "dirty" or "exclusive" state)

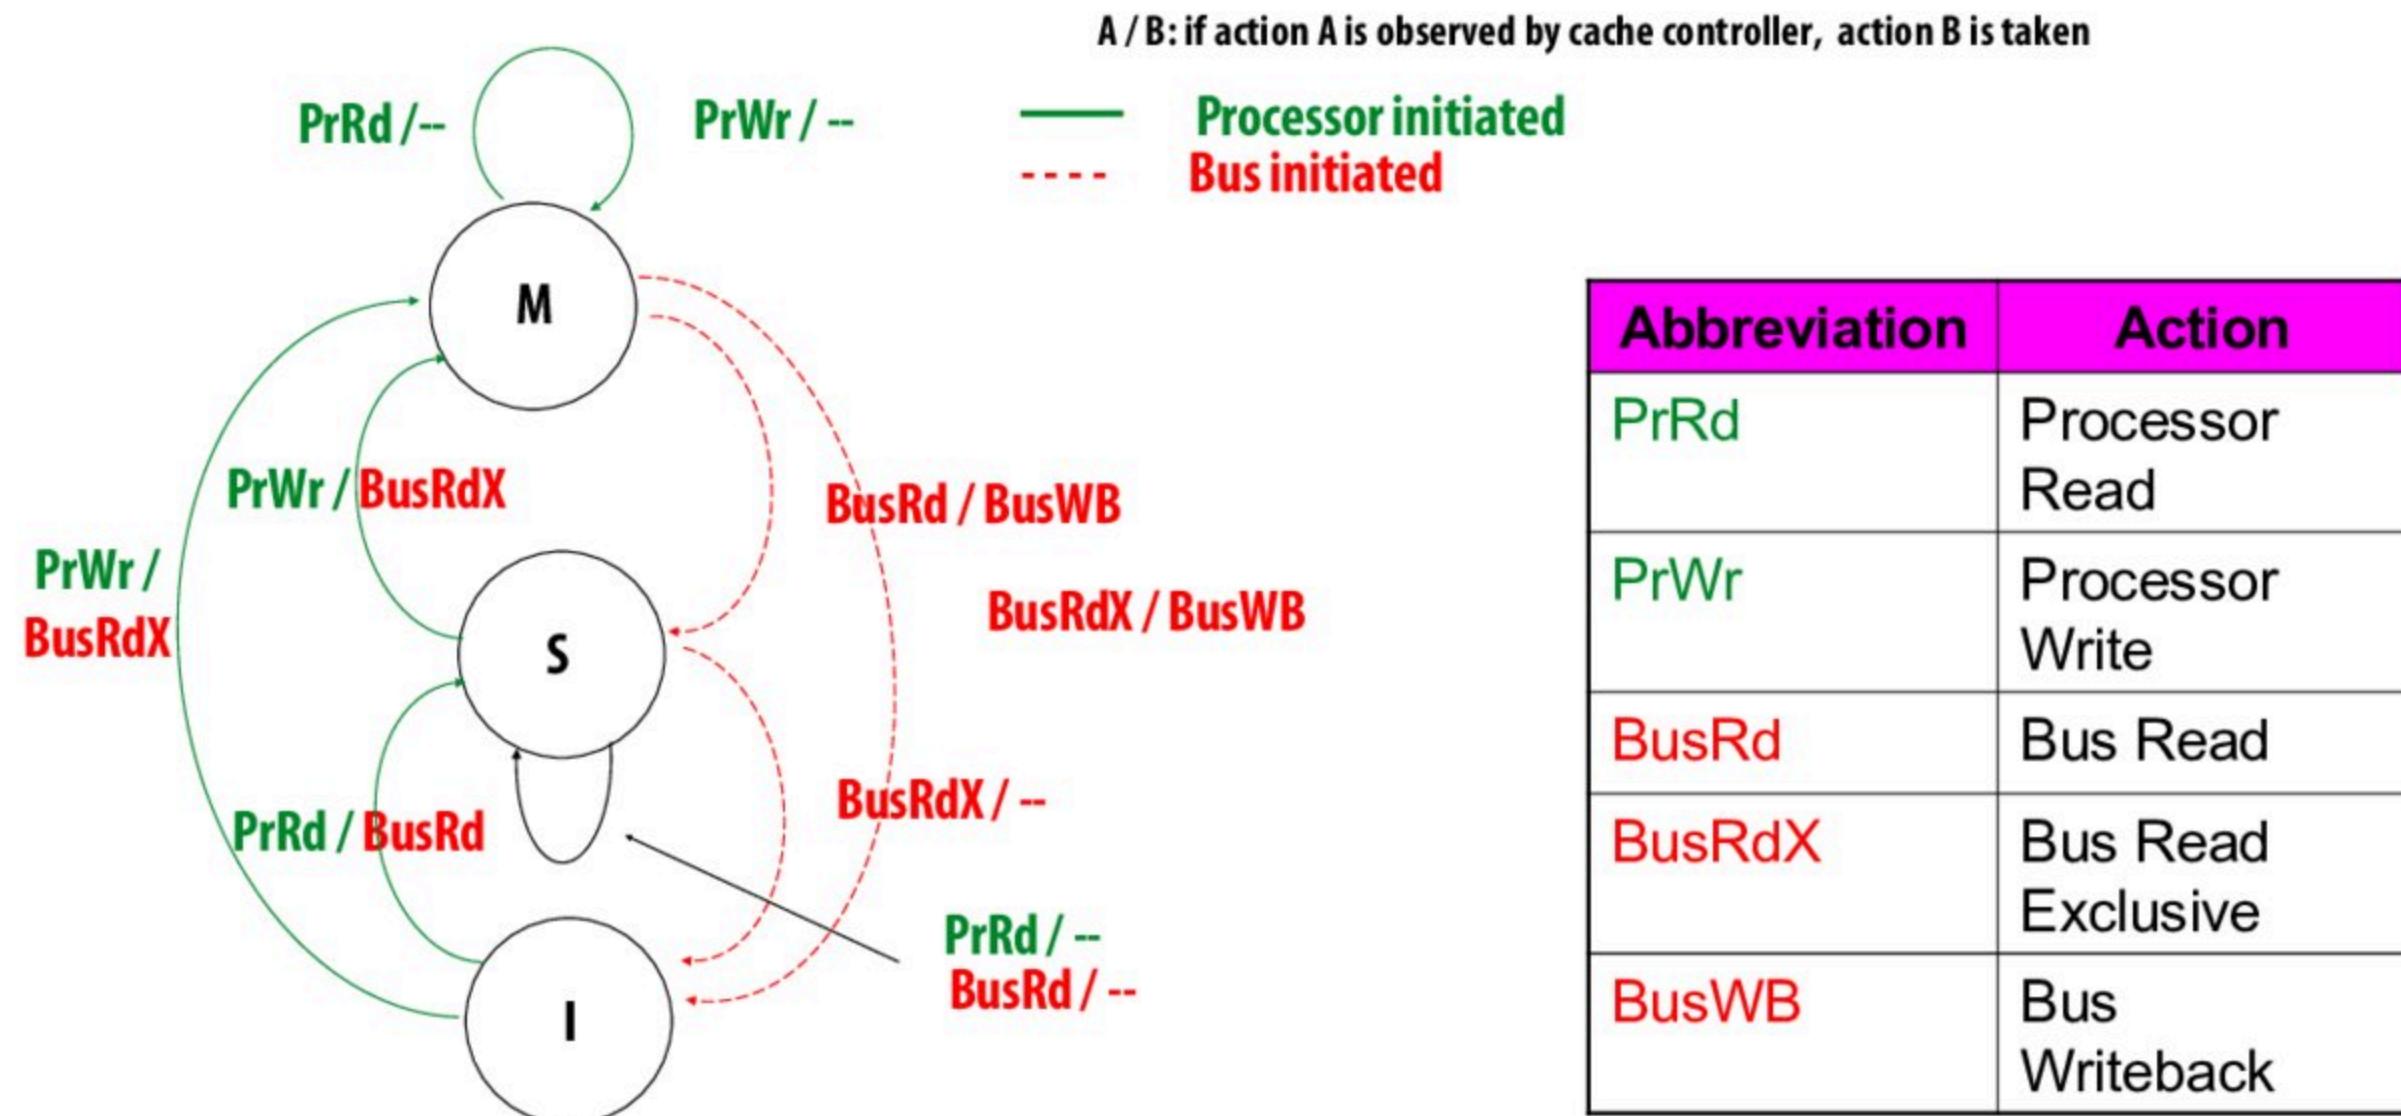
Two processor operations (triggered by local CPU)

- PrRd (read)
- PrWr (write)

Three coherence-related bus transactions (from remote caches)

- BusRd: obtain copy of line with no intent to modify
- BusRdX: obtain copy of line with intent to modify
- BusWB: write dirty line out to memory

Cache Coherence Protocol: MSI State Transition Diagram



MSI Invalidate Protocol

Read obtains block in “shared”

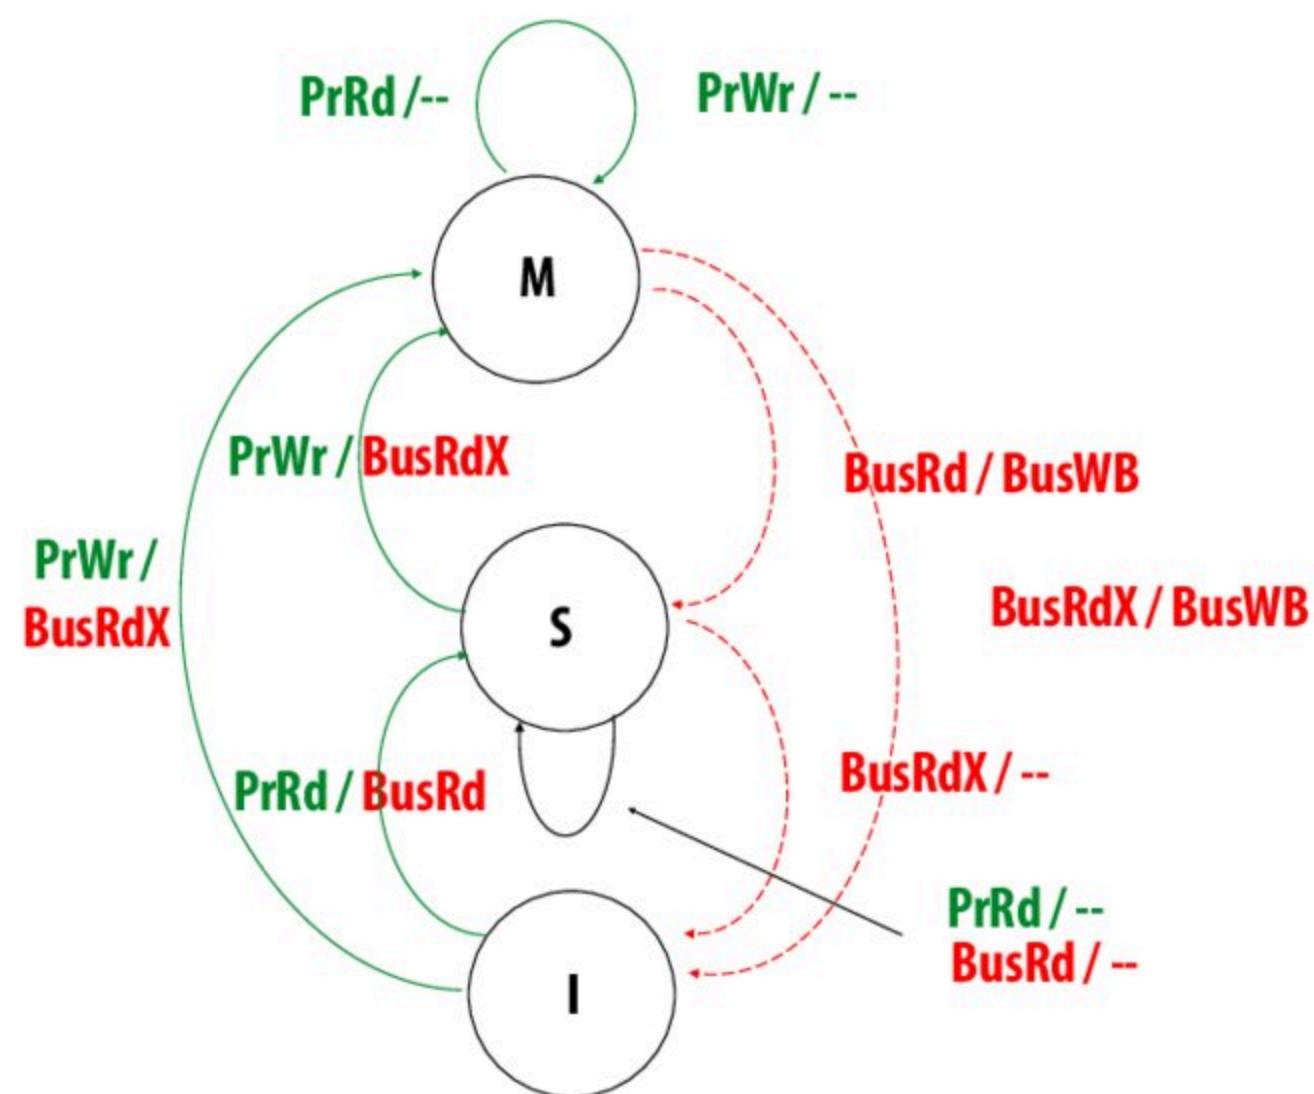
- even if only cached copy

Obtain exclusive ownership before writing

- BusRdX causes others to invalidate
- If M in another cache, will cause writeback
- BusRdX even if hit in S
 - promote to M (upgrade)

A / B: if action A is observed by cache controller, action B is taken

— Processor initiated
- - - Bus initiated



* Remember, all caches are carrying out this logic independently to maintain coherence

A Cache Coherence Example

Proc Action	P1 \$-state	P2 \$-state	P3 \$-state	Bus Trans	Data from
P1 read x	S	--	--	BusRd	Memory
P3 read x	S	--	S	BusRd	Memory
P3 write x	I	--	M	BusRdX	Memory
P1 read x	S	--	S	BusRd	P3 \$
P1 read x	S	--	S		P1 \$
P2 write x	I	M	I	BusRdX	Memory

Single writer, multiple reader protocol

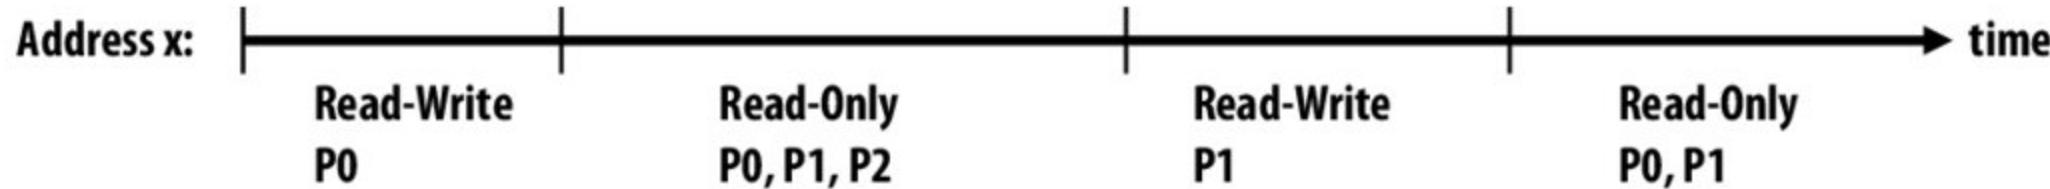
Why do you need Modified to Shared?

Communication increases memory latency

How Does MSI Satisfy Cache Coherence Invariants?

1. Single-Writer, Multiple-Read (SWMR) Invariant

2. Data-Value Invariant (write serialization)



Summary: MSI

A line in the M state can be modified without notifying other caches

- No other caches have the line resident, so other processors cannot read these values
- (without generating a memory read transaction)

Processor can only write to lines in the M state

- If processor performs a write to a line that is not exclusive in cache, cache controller must first broadcast a read-exclusive transaction to move the line into that state
- Read-exclusive tells other caches about impending write
("you can't read any more, because I'm going to write")
- Read-exclusive transaction is required even if line is valid (but not exclusive... it's in the S state) in processor's local cache (why?)
- Dirty state implies exclusive

When cache controller snoops a "read exclusive" for a line it contains

- Must invalidate the line in its cache
- Because if it didn't, then multiple caches will have the line
(and so it wouldn't be exclusive in the other cache!)

MESI invalidation protocol

MSI requires two interconnect transactions for the common case of reading an address, then writing to it

- Transaction 1: BusRd to move from I to S state
- Transaction 2: BusRdX to move from S to M state



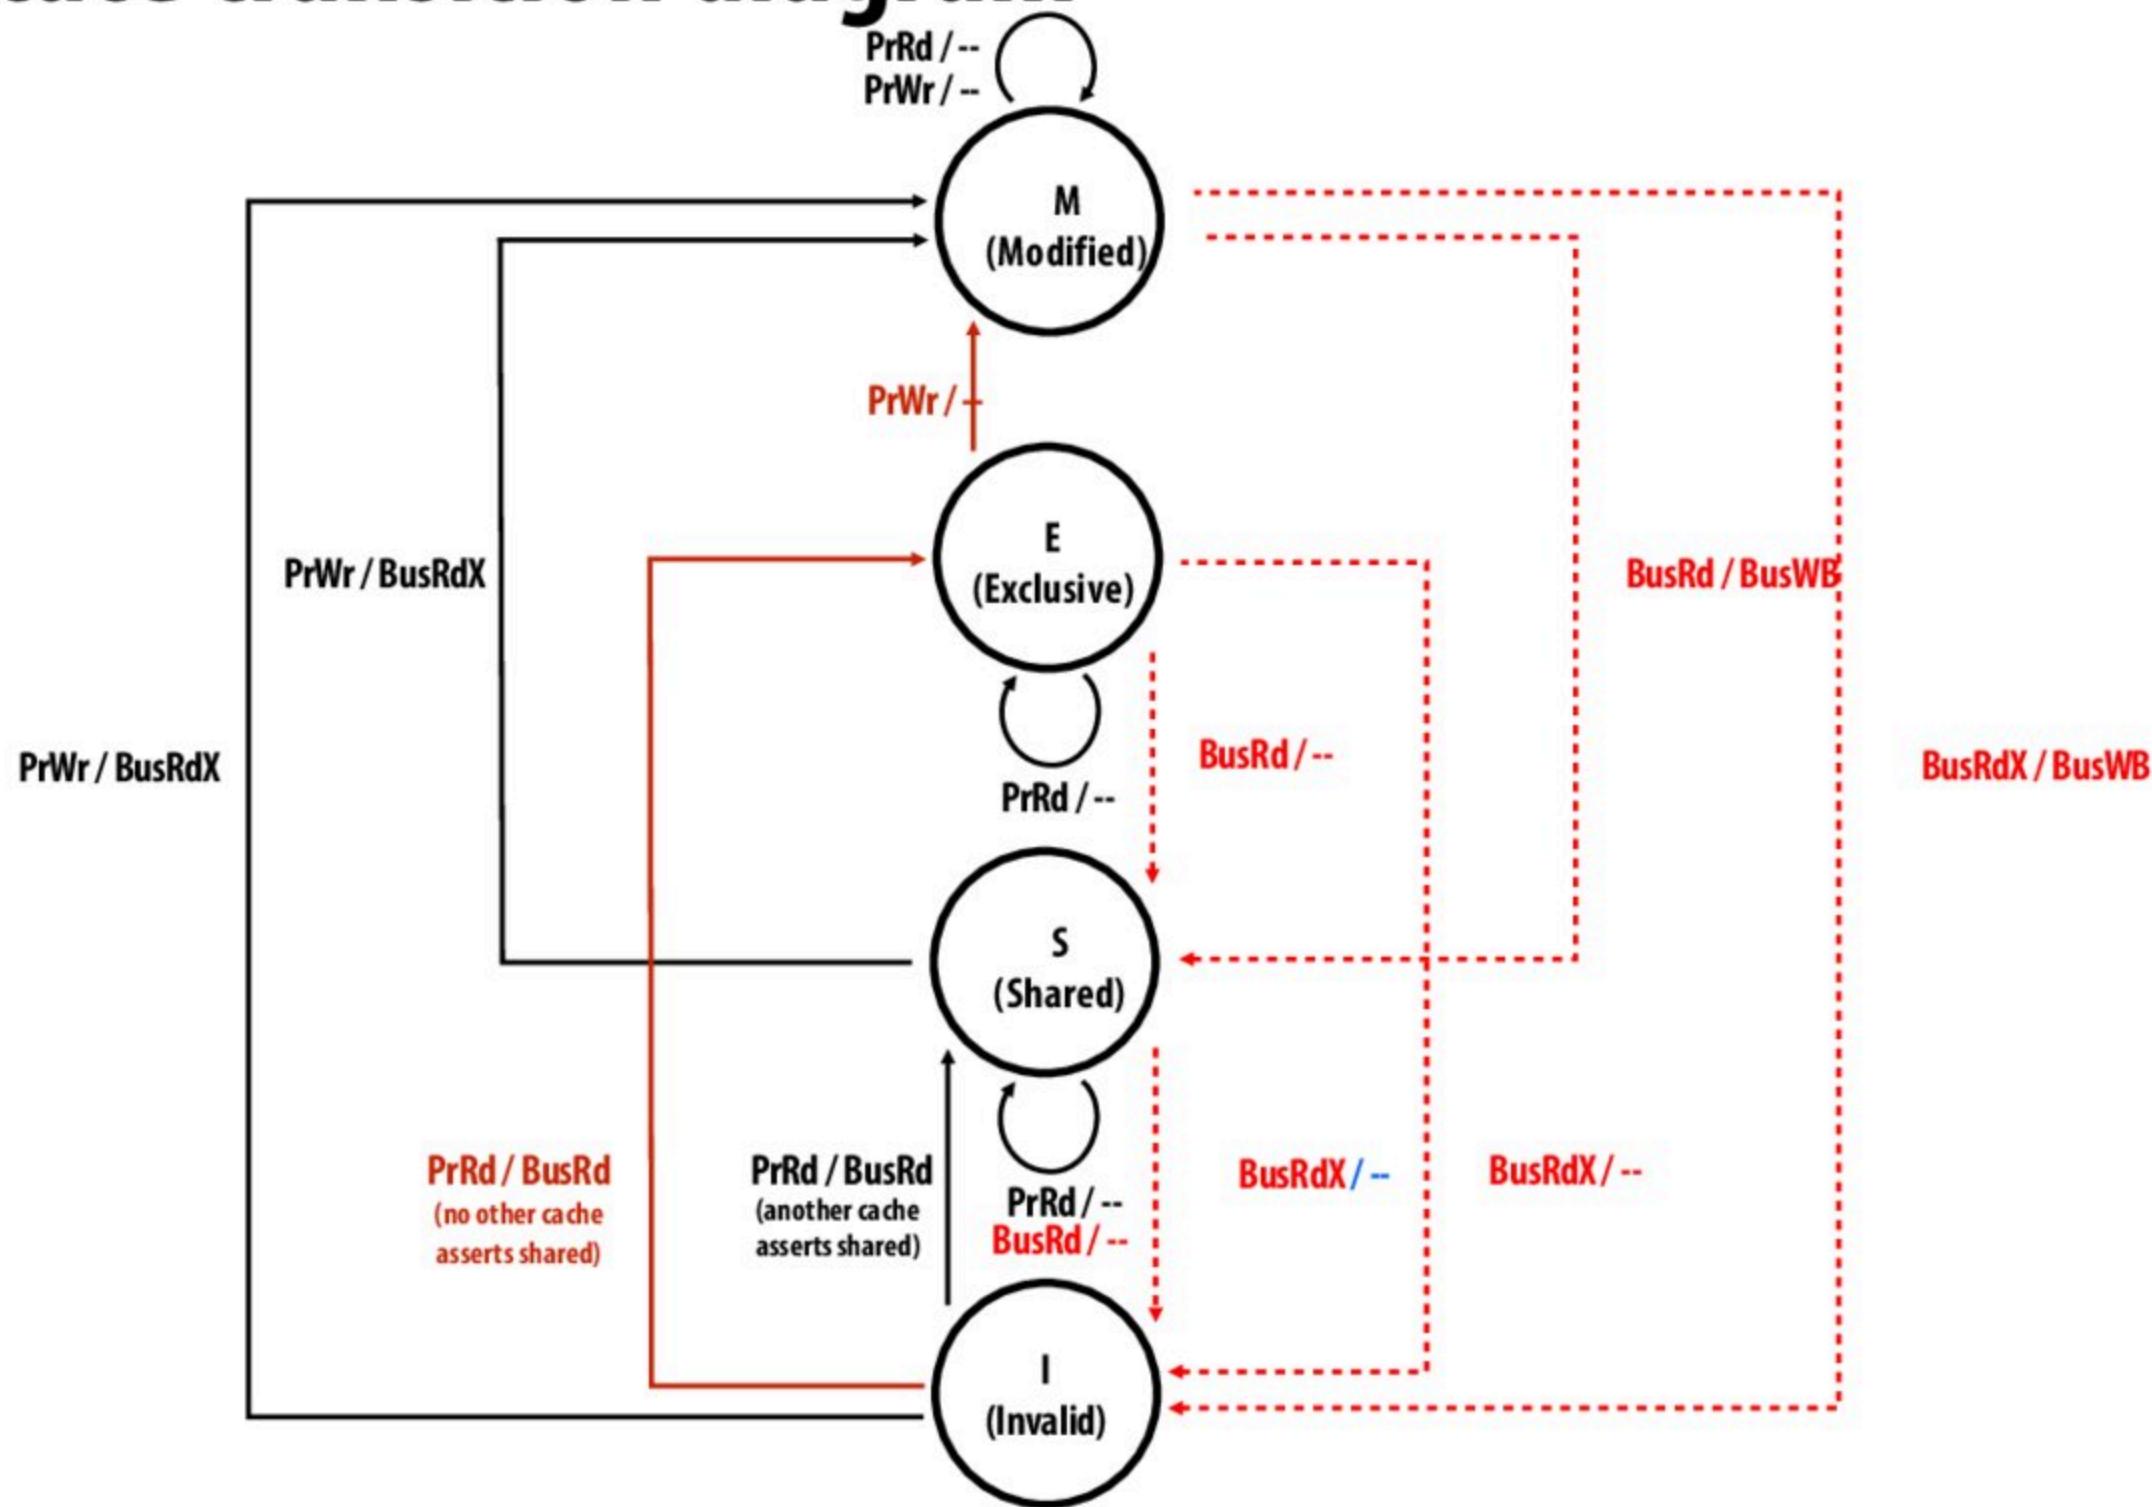
MESI, not Messi!

This inefficiency exists even if application has no sharing at all

Solution: add additional state E (“exclusive clean”)

- Line has not been modified, but only this cache has a copy of the line
- Decouples exclusivity from line ownership (line not dirty, so copy in memory is valid copy of data)
- Upgrade from E to M does not require an bus transaction

MESI state transition diagram



Scalable cache coherence using directories

Snooping schemes broadcast coherence messages to determine the state of a line in the other caches: not scalable and too restrictive

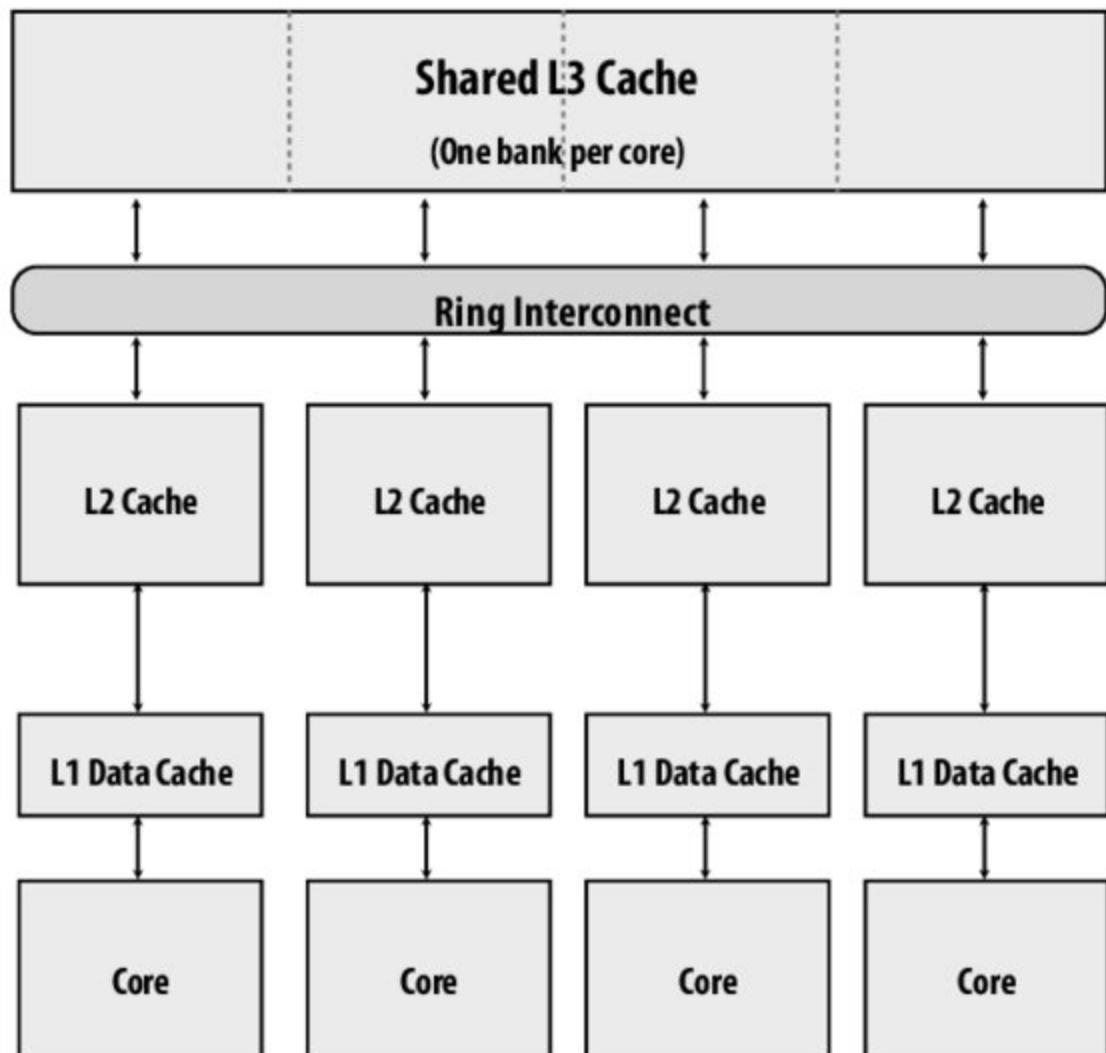
Alternative idea: avoid broadcast by storing information about the status of the line in one place: a “directory”

- The directory entry for a cache line contains information about the state of the cache line in all caches.
- Caches look up information from the directory as necessary
- Cache coherence is maintained by point-to-point messages between the caches on a “need to know” basis (not by broadcast mechanisms)

■ Still need to maintain invariants

- SWMR
- Write serialization

Directory coherence in Intel Core i7 CPU



L3 serves as centralized directory for all lines in the L3 cache

- **Serialization point**

(Since L3 is an inclusive cache, any line in L2 is guaranteed to also be resident in L3)

**Directory maintains list of L2 caches containing line
Instead of broadcasting coherence traffic to all L2's, only send coherence messages to L2's that contain the line**

(Core i7 interconnect is a ring, it is not a bus)

Directory dimensions:

- **P=4**
- **M = number of L3 cache lines**

Implications of cache coherence to the programmer

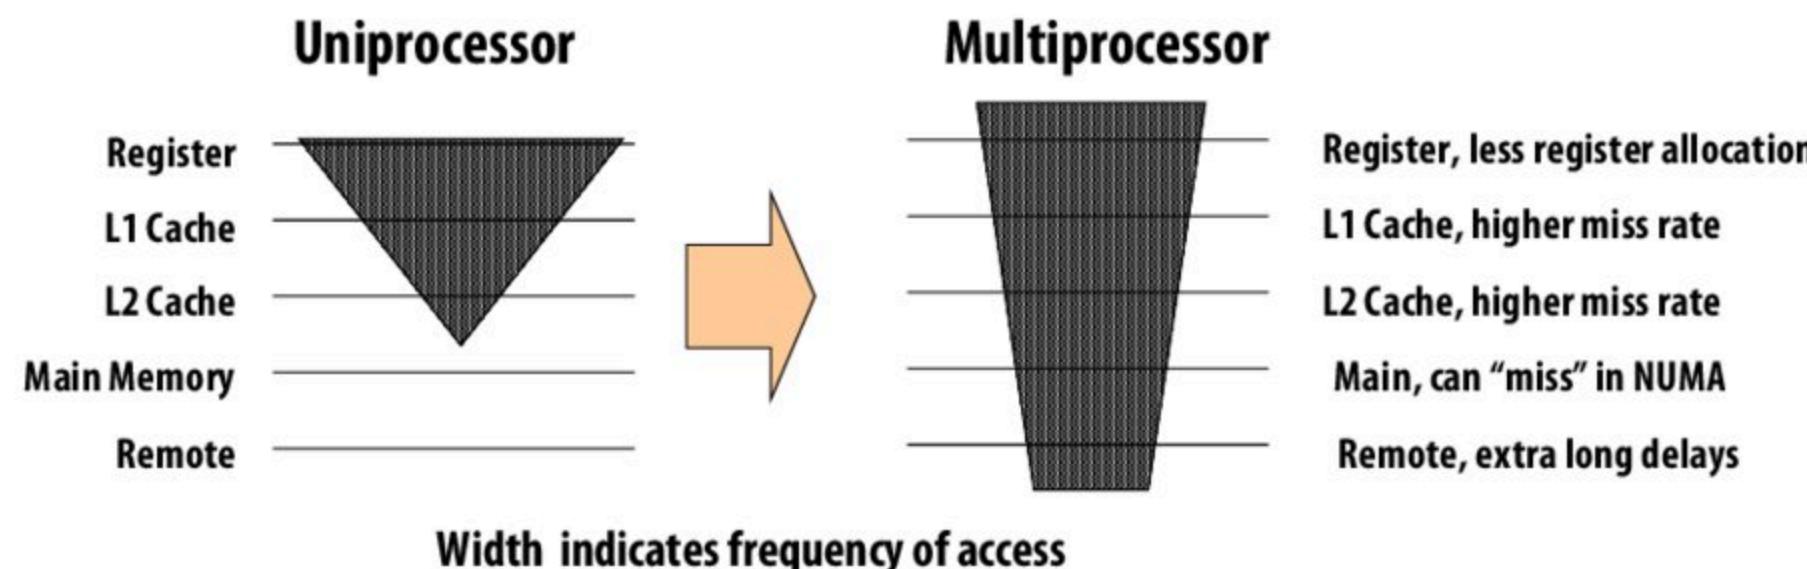
Communication Overhead

Communication time is a key parallel overhead

- Appears as increased memory access time in multiprocessor
 - Extra main memory accesses in UMA systems
 - Must determine increase in cache miss rate vs. uniprocessor
 - Some accesses have higher latency in NUMA systems
 - Only a fraction of a % of these can be significant!

$$\text{AMAT}_{\text{Multiprocessor}} > \text{AMAT}_{\text{Uniprocessor}}$$

Average Memory Access Time (AMAT) = $\sum_0^n \text{frequency of access} \times \text{latency of access}$



Core i7 Xeon 5500 Series Data Source Latency (approx.)	
L1 hit,	~4 cycles
L2 hit,	~10 cycles
L3 hit, line unshared	~40 cycles
L3 hit, shared line in another core	~65 cycles
L3 hit, modified in another core	~75 cycles remote
Local DRAM	~30 ns (~120 cycles)
Remote DRAM	~100 ns (~400 cycles)

Use system tools to optimize cache performance

Memory Access Analysis for Cache Misses and High Bandwidth Issues

Use the Intel® VTune™ Profiler's Memory Access analysis to identify memory-related issues, like NUMA problems and bandwidth-limited accesses, and attribute performance events to memory objects (data structures), which is provided due to instrumentation of memory allocations/de-allocations and getting static/global variables from symbol information.

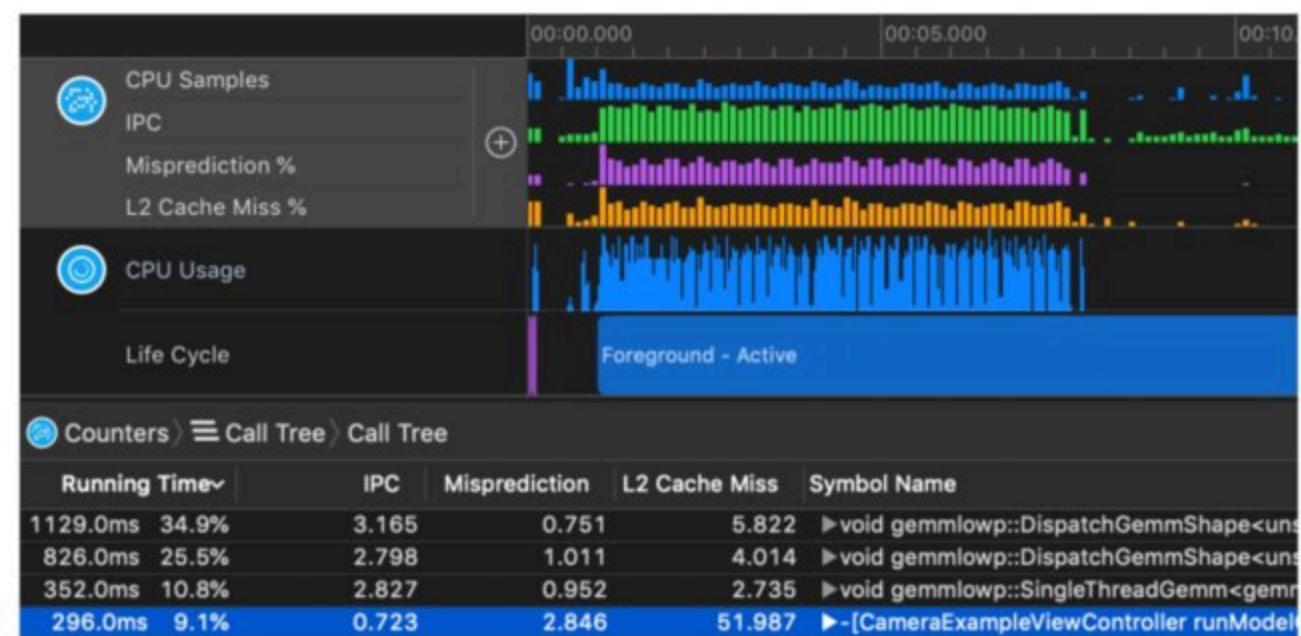
NOTE:
Intel® VTune™ Profiler is a new renamed version of the Intel® VTune™ Amplifier.

How It Works

Grouping:	Bandwidth Domain / Bandwidth Utilization Type / Function / Call Stack	%				
Bandwidth Domain / Bandwidth Utilization Type / Function / Call Stack	CPU Time	Memory Bound	Loads	Stores	LLC Miss Count	Average Latency (cycles)
DRAM, GB/sec	9.708	64.3%	4,517.0	4,141.26	191,811,508	90
High	4.258	56.0%	2,345.0	2,111.22	119,007,140	115
Medium	4.258	54.6%	2,310.0	2,046.82	119,007,140	115
F_>_mem	0.175	100.0%	175,000	63,000,945	0	0
F_>_softf	0.012	0.0%	0	0	0	0
F_>_timer_softf	0.000	0.0%	0	0	0	0
F_>_page_softf	0.000	0.0%	0	0	0	0
Numa_migrate_prep	0.000	0.0%	0	0	0	0
Host_coutine	0.0	0.0%	0	0	0	0
Medium	2.480	70.3%	2,765.0	982,414...	52,853,171	83

Memory Access analysis type uses hardware event-based sampling to collect data for the following metrics:

- **Loads and Stores** metrics that show the total number of loads and stores
- **LLC Miss Count** metric that shows the total number of last-level cache misses
 - **Local DRAM Access Count** metric that shows the total number of LLC misses serviced by the local memory
 - **Remote DRAM Access Count** metric that shows the number of accesses to the remote socket memory
 - **Remote Cache Access Count** metric that shows the number of accesses to the remote socket cache
- **Memory Bound** metric that shows a fraction of cycles spent waiting due to demand load or store instructions
 - **L1 Bound** metric that shows how often the machine was stalled without missing the L1 data cache
 - **L2 Bound** metric that shows how often the machine was stalled on L2 cache
 - **L3 Bound** metric that shows how often the CPU was stalled on L3 cache, or contended with a sibling core
 - **L3 Latency** metric that shows a fraction of cycles with demand load accesses that hit the L3 cache under unloaded scenarios (possibly L3 latency limited)
 - **NUMA: % of Remote Accesses** metric shows percentage of memory requests to remote DRAM. The lower its value is, the better.
 - **DRAM Bound** metric that shows how often the CPU was stalled on the main memory (DRAM). This metric enables you to identify **DRAM Bandwidth Bound**, **UPI Utilization Bound** issues, as well as **Memory Latency** issues with the following metrics:
 - **Remote / Local DRAM Ratio** metric that is defined by the ratio of remote DRAM loads to local DRAM loads
 - **Local DRAM** metric that shows how often the CPU was stalled on loads from the local memory
 - **Remote DRAM** metric that shows how often the CPU was stalled on loads from the remote memory
 - **Remote Cache** metric that shows how often the CPU was stalled on loads from the remote cache in other sockets
- **Average Latency** metric that shows an average load latency in cycles



Apple Xcode Instruments

Intel VTune

Unintended communication via false sharing

What is the potential performance problem with this code?

```
// allocate per-thread variable for local per-thread accumulation
int myPerThreadCounter[NUM_THREADS];
```

Why might this code be more performant?

```
// allocate per thread variable for local accumulation
struct PerThreadState {
    int myPerThreadCounter;
    char padding[CACHE_LINE_SIZE - sizeof(int)];
};
PerThreadState myPerThreadCounter[NUM_THREADS];
```

Demo: false sharing

```
void* worker(void* arg) {  
    volatile int* counter = (int*)arg;  
    for (int i=0; i<MANY_ITERATIONS; i++)  
        (*counter)++;  
  
    return NULL;  
}
```

```
void test1(int num_threads) {  
  
    pthread_t threads[MAX_THREADS];  
    int counter[MAX_THREADS];  
  
    for (int i=0; i<num_threads; i++)  
        pthread_create(&threads[i], NULL,  
                      &worker, &counter[i]);  
  
    for (int i=0; i<num_threads; i++)  
        pthread_join(threads[i], NULL);  
}
```

**Execution time with
num_threads=8 on 4-core system:
14.2 sec**

**threads update a per-thread counter
many times**

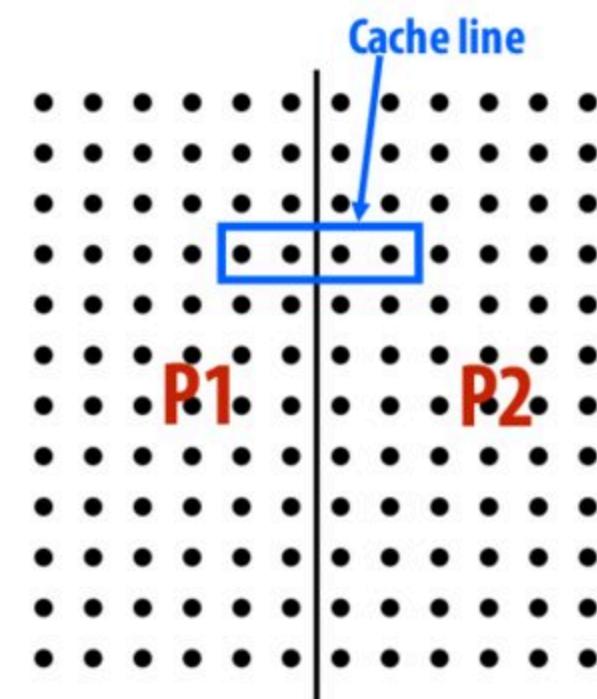
```
struct padded_t {  
    int counter;  
    char padding[CACHE_LINE_SIZE - sizeof(int)];  
};  
  
void test2(int num_threads) {  
  
    pthread_t threads[MAX_THREADS];  
    padded_t counter[MAX_THREADS];  
  
    for (int i=0; i<num_threads; i++)  
        pthread_create(&threads[i], NULL,  
                      &worker, &(counter[i].counter));  
  
    for (int i=0; i<num_threads; i++)  
        pthread_join(threads[i], NULL);  
}
```

**Execution time with
num_threads=8 on 4-core system:
4.7 sec**

False sharing

Condition where two processors write to different addresses, but addresses map to the same cache line

Cache line “ping-pongs” between caches of writing processors, generating significant amounts of communication due to the coherence protocol

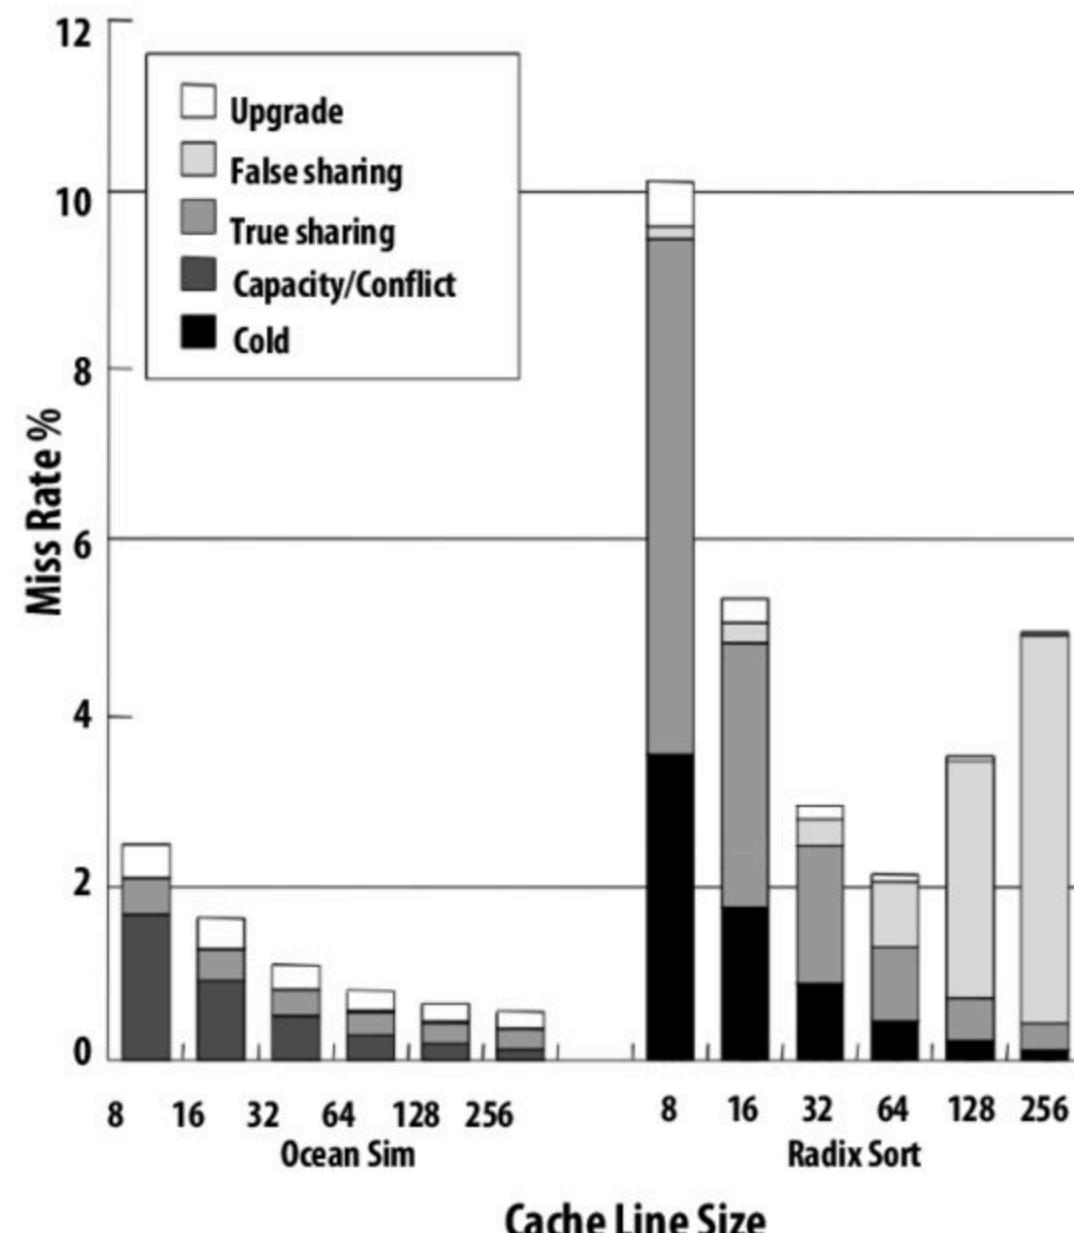
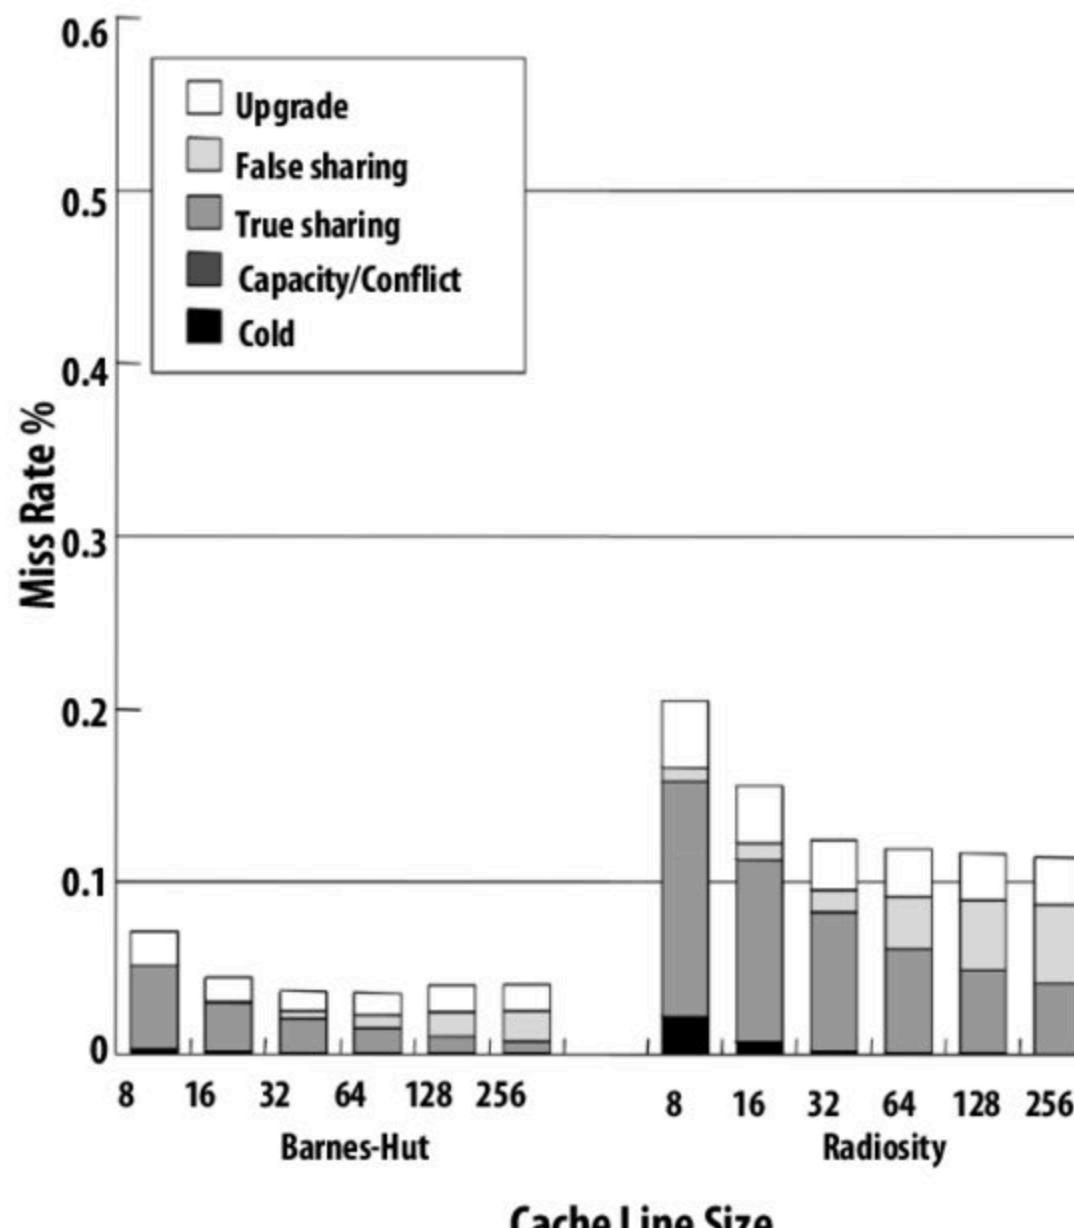


No inherent communication, this is entirely artifactual communication (cachelines > 4B)

False sharing can be a factor in when programming for cache-coherent architectures

Impact of cache line size on miss rate

Results from simulation of a 1 MB cache (four example applications)



* Note: I separated the results into two graphs because of different Y-axis scales

Figure credit: Culler, Singh, and Gupta

Stanford CS149, Fall 2024

Summary: Cache coherence

The cache coherence problem exists because the abstraction of a single shared address space is not implemented by a single storage unit

- Storage is distributed among main memory and local processor caches
- Data is replicated in local caches for performance

Main idea of snooping-based cache coherence: whenever a cache operation occurs that could affect coherence, the cache controller **broadcasts a notification to all other cache controllers in the system**

- Challenge for HW architects: minimizing overhead of coherence implementation
- Challenge for SW developers: be wary of artifactual communication due to coherence protocol (e.g., false sharing)

Scalability of snooping implementations is limited by ability to broadcast coherence messages to all caches!

- Scaling cache coherence via directory-based approaches

Lecture 14+:

Memory Consistency

Parallel Computing
Stanford CS149, Fall 2024

Shared Memory Behavior

Intuition says loads should return latest value written

- **What is latest?**
- **Coherence: only one memory location**
- **Consistency: apparent ordering for all locations**
 - Order in which memory operations performed by one thread become visible to other threads

Affects

- **Programmability: how programmers reason about program behavior**
 - Allowed behavior of multithreaded programs executing with shared memory
- **Performance: limits HW/SW optimizations that can be used**
 - Reordering memory operations to hide latency

Today: who should care

Anyone who:

- Wants to implement a synchronization library
- Will ever work a job in kernel (or driver) development
- Seeks to implement lock-free data structures *

* Topic of a later lecture

Memory coherence vs. memory consistency

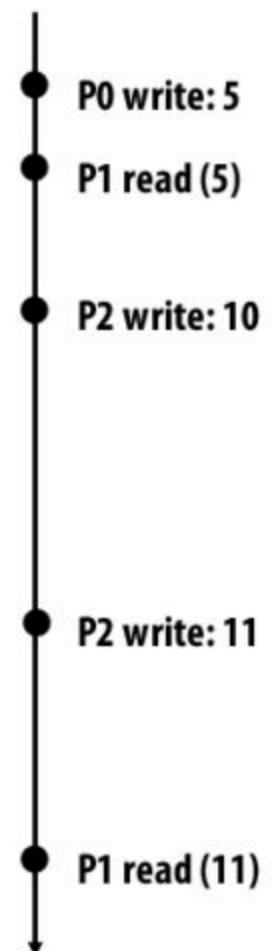
Memory coherence defines requirements for the observed behavior of reads and writes to the same memory location

- All processors must agree on the order of reads/writes to X
- In other words: it is possible to put all operations involving X on a timeline such that the observations of all processors are consistent with that timeline

Memory consistency defines the behavior of reads and writes to different locations (as observed by other processors)

- Coherence only guarantees that writes to address X will eventually propagate to other processors
- Consistency deals with when writes to X propagate to other processors, relative to reads and writes to other addresses

Observed chronology of operations on address X



Coherence vs. Consistency

(said again, perhaps more intuitively this time)

The goal of cache coherence is to ensure that the memory system in a parallel computer behaves as if the caches were not there

- Just like how the memory system in a uni-processor system behaves as if the cache was not there

A system without caches would have no need for cache coherence

Memory consistency defines the allowed behavior of loads and stores to different addresses in a parallel system

- The allowed behavior of memory should be specified whether or not caches are present (and that's what a memory consistency model does)

Memory Consistency

TL;DR:

- Multiprocessors reorder memory operations in unintuitive and strange ways
- This behavior is necessary for performance
- Application programmers rarely see this behavior
- Systems (OS and compiler) developers see it all the time

Memory operation ordering

A program defines a sequence of loads and stores
(this is the “program order” of the loads and stores)

Four types of memory operation orderings

- $W_X \rightarrow R_Y$: write to X must commit before subsequent read from Y *
- $R_X \rightarrow R_Y$: read from X must commit before subsequent read from Y
- $R_X \rightarrow W_Y$: read to X must commit before subsequent write to Y
- $W_X \rightarrow W_Y$: write to X must commit before subsequent write to Y

* To clarify: “write must commit before subsequent read” means:

When a write comes before a read in program order, the write must commit (its results are visible) by the time the read occurs.

Multiprocessor Execution

Initially $A = B = 0$

Proc 0

- (1) $A = 1$
- (2) **print B**

Proc 1

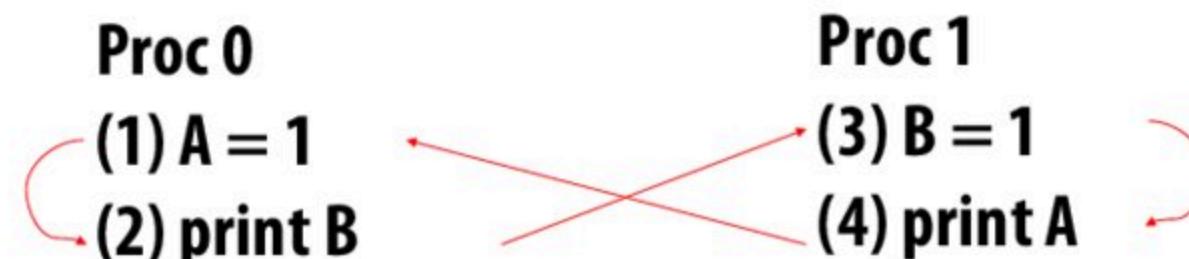
- (3) $B = 1$
- (4) **print A**

What can be printed?

- “01”?
- “10”?
- “11”?
- “00”?

Orderings That Should Not Happen

Initially $A = B = 0$



The program should not print “00” or “10”

A “happens-before” graph shows the order in which events must execute to get a desired outcome

If there’s a cycle in the graph, an outcome is impossible—an event must happen before itself!

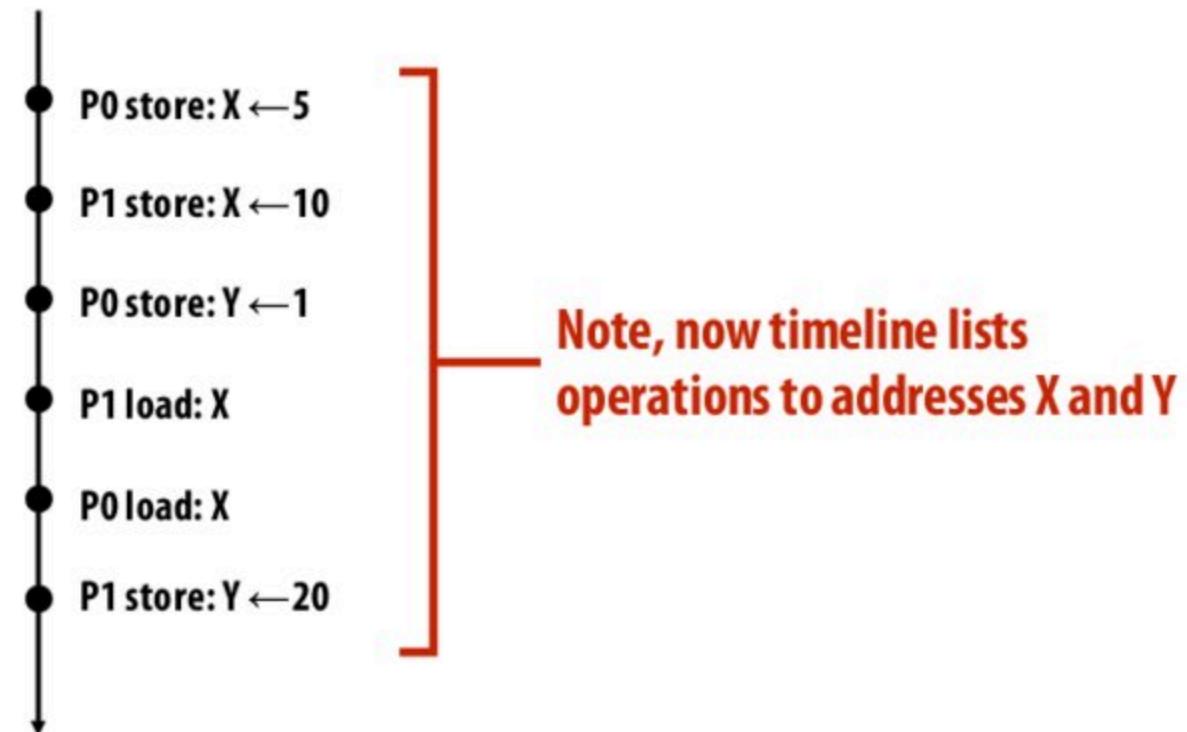
What Should Programmers Expect

Sequential Consistency

- Lamport 1976 (Turing Award 2013)
- All operations executed in some sequential order
 - As if they were manipulating a single shared memory
- Each thread's operations happen in program order

A sequentially consistent memory system maintains all four memory operation orderings ($W_X \rightarrow R_Y$, $R_X \rightarrow R_Y$, $R_X \rightarrow W_Y$, $W_X \rightarrow W_Y$)

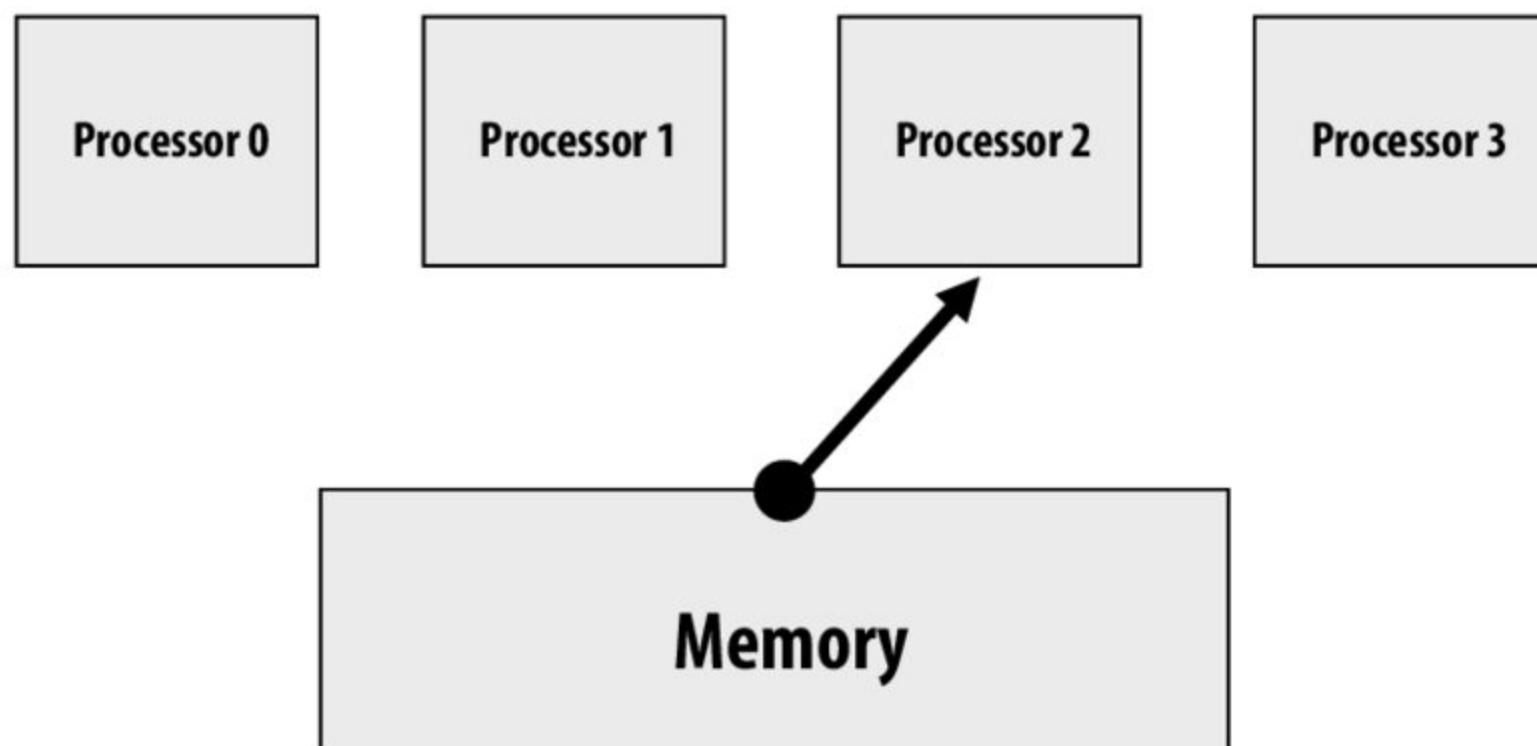
There is a chronology of all memory operations that is consistent with observed values



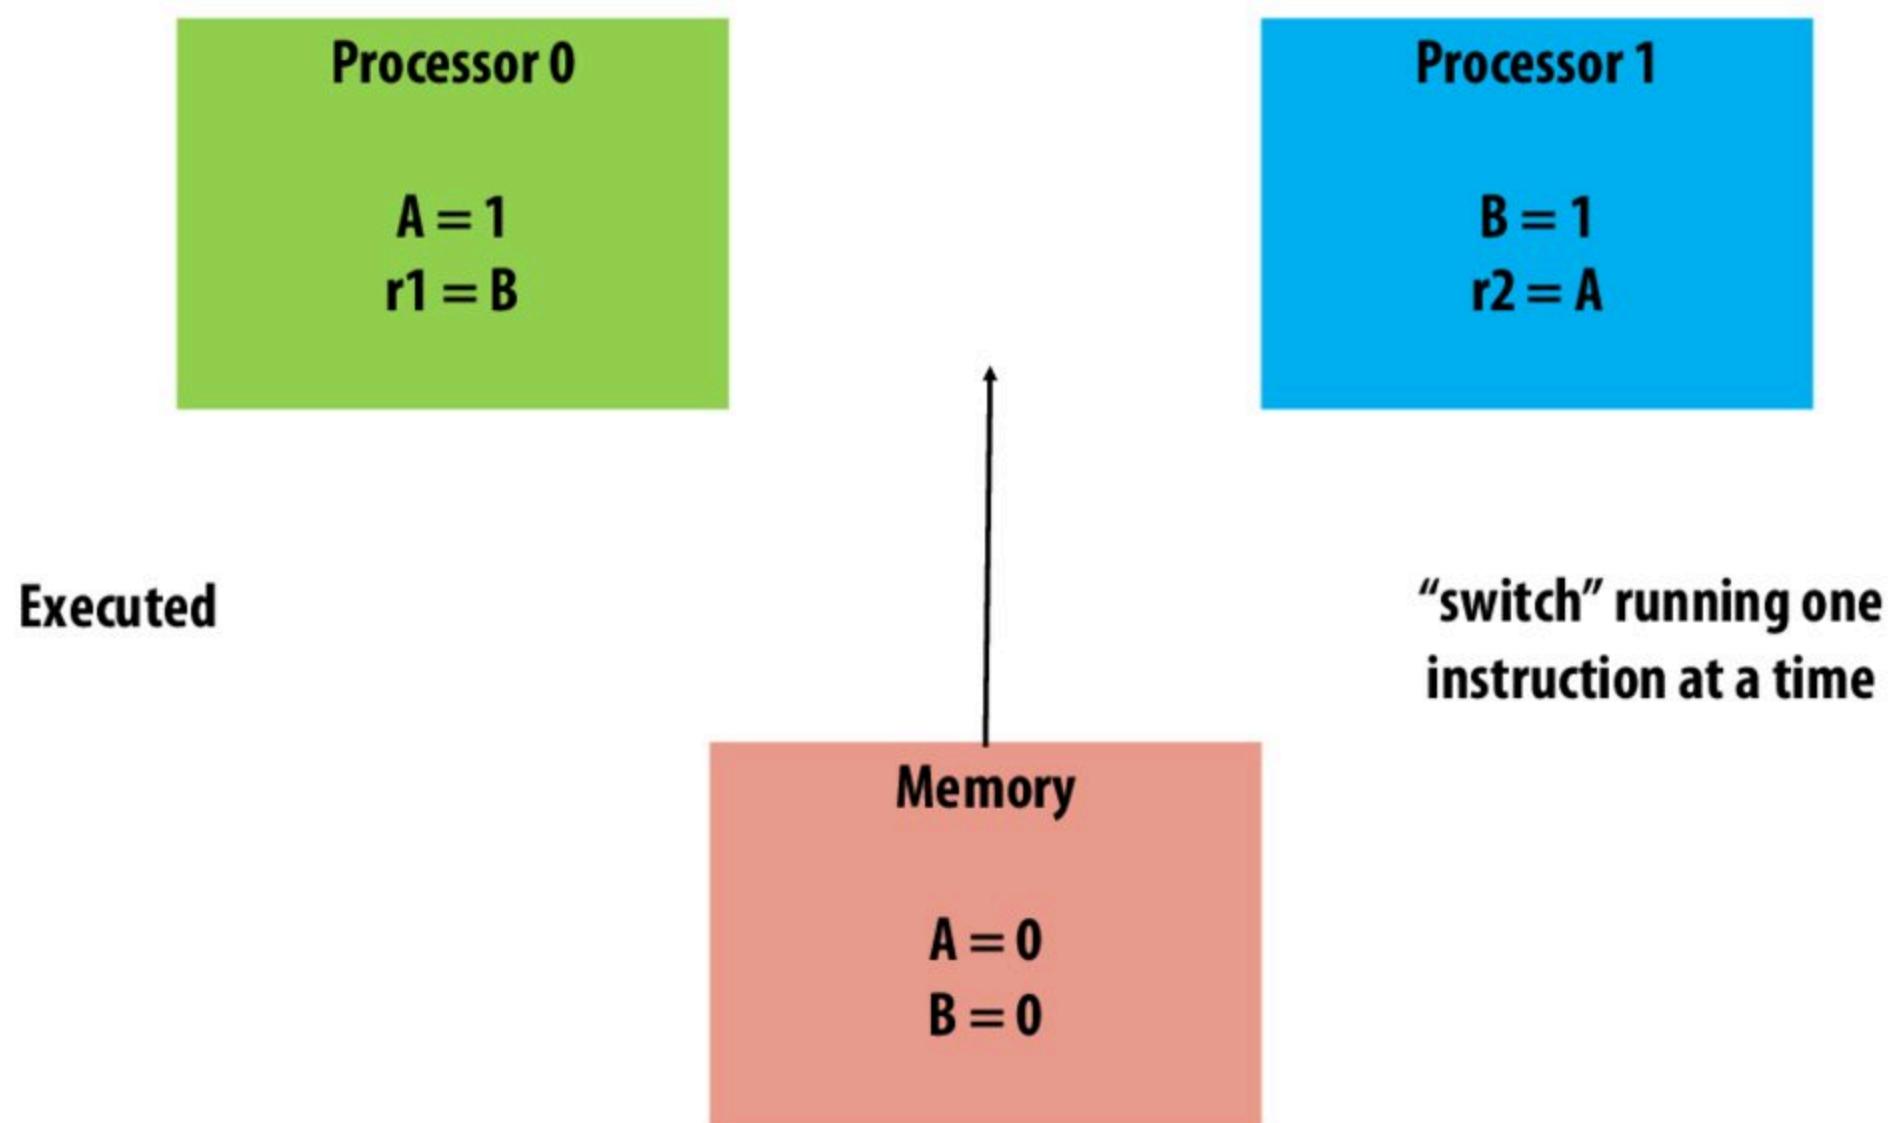
Sequential consistency (switch metaphor)

All processors issue loads and stores in program order

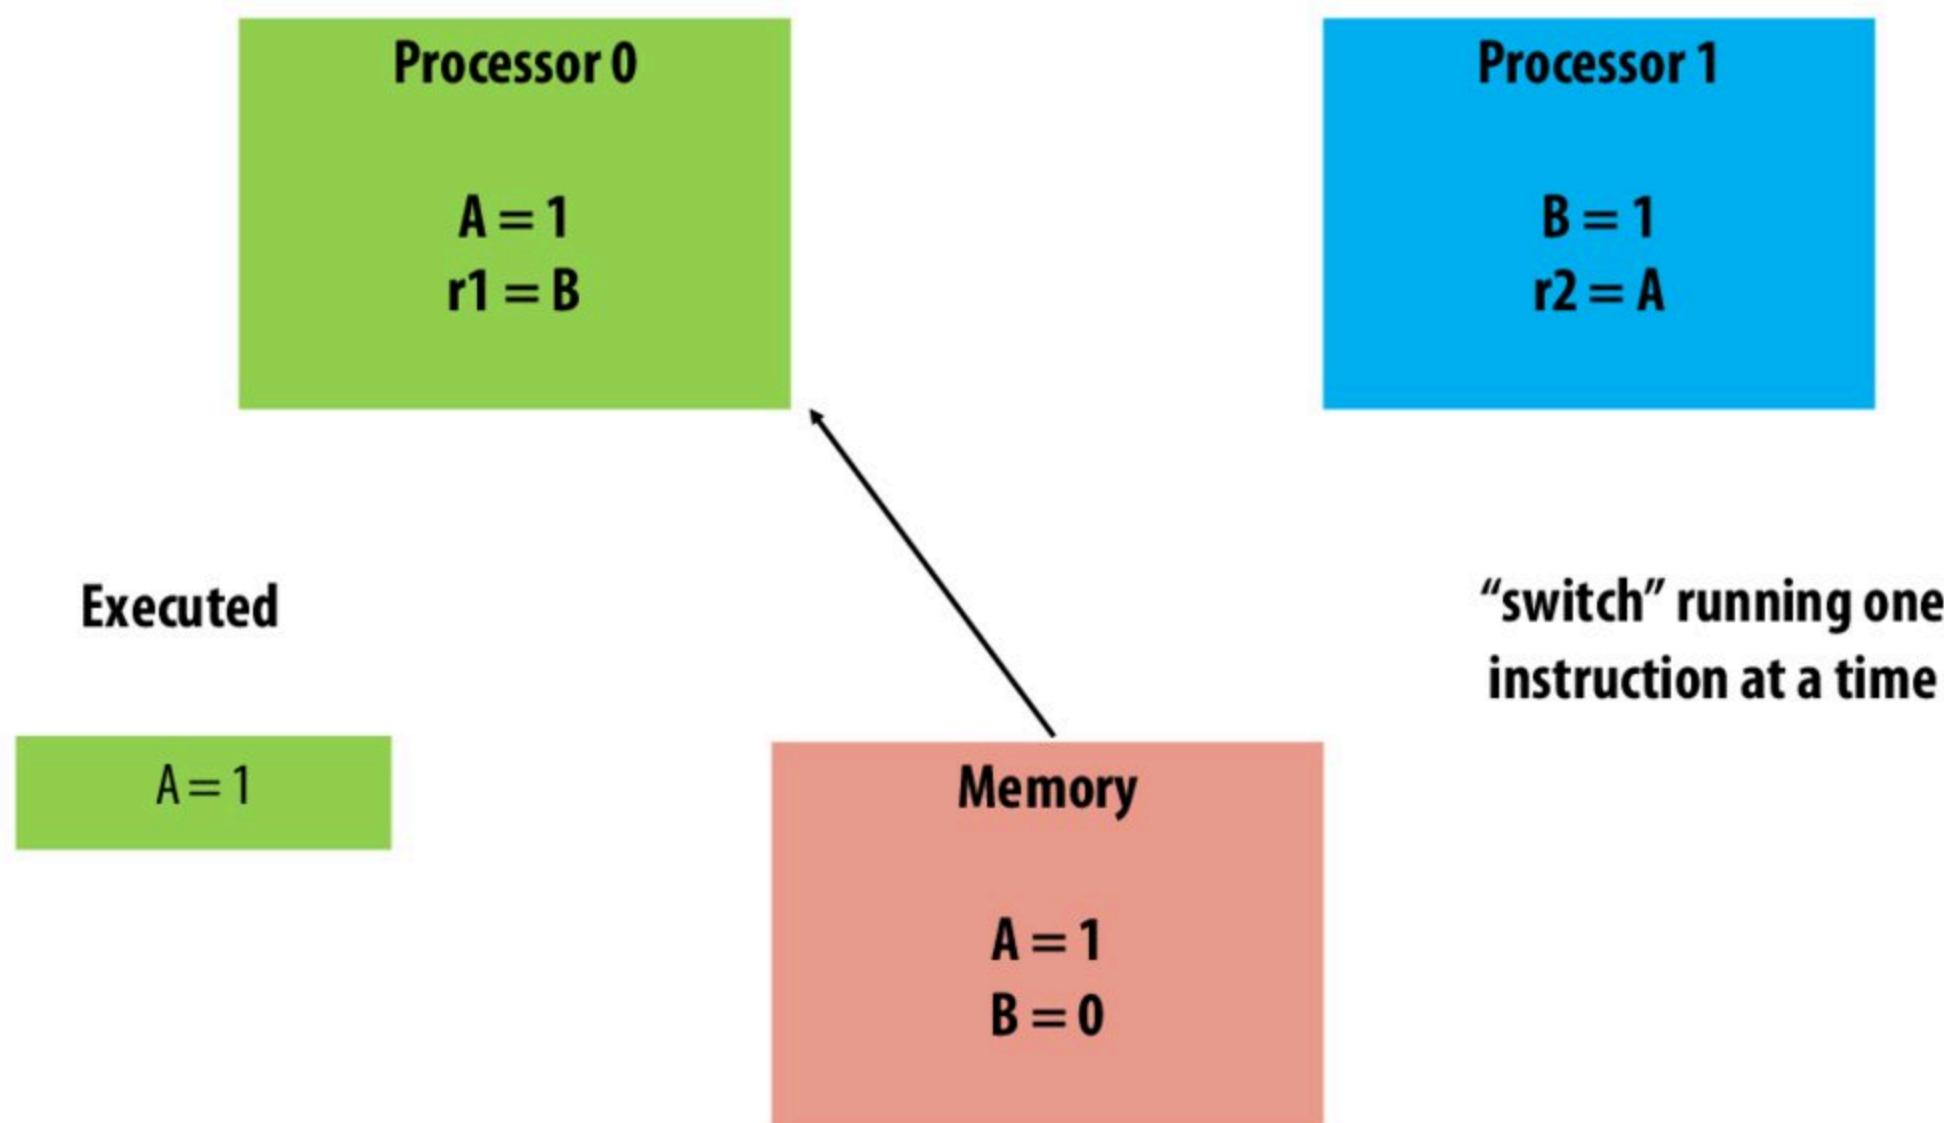
Memory chooses a processor at random, performs a memory operation to completion, then chooses another processor, ...



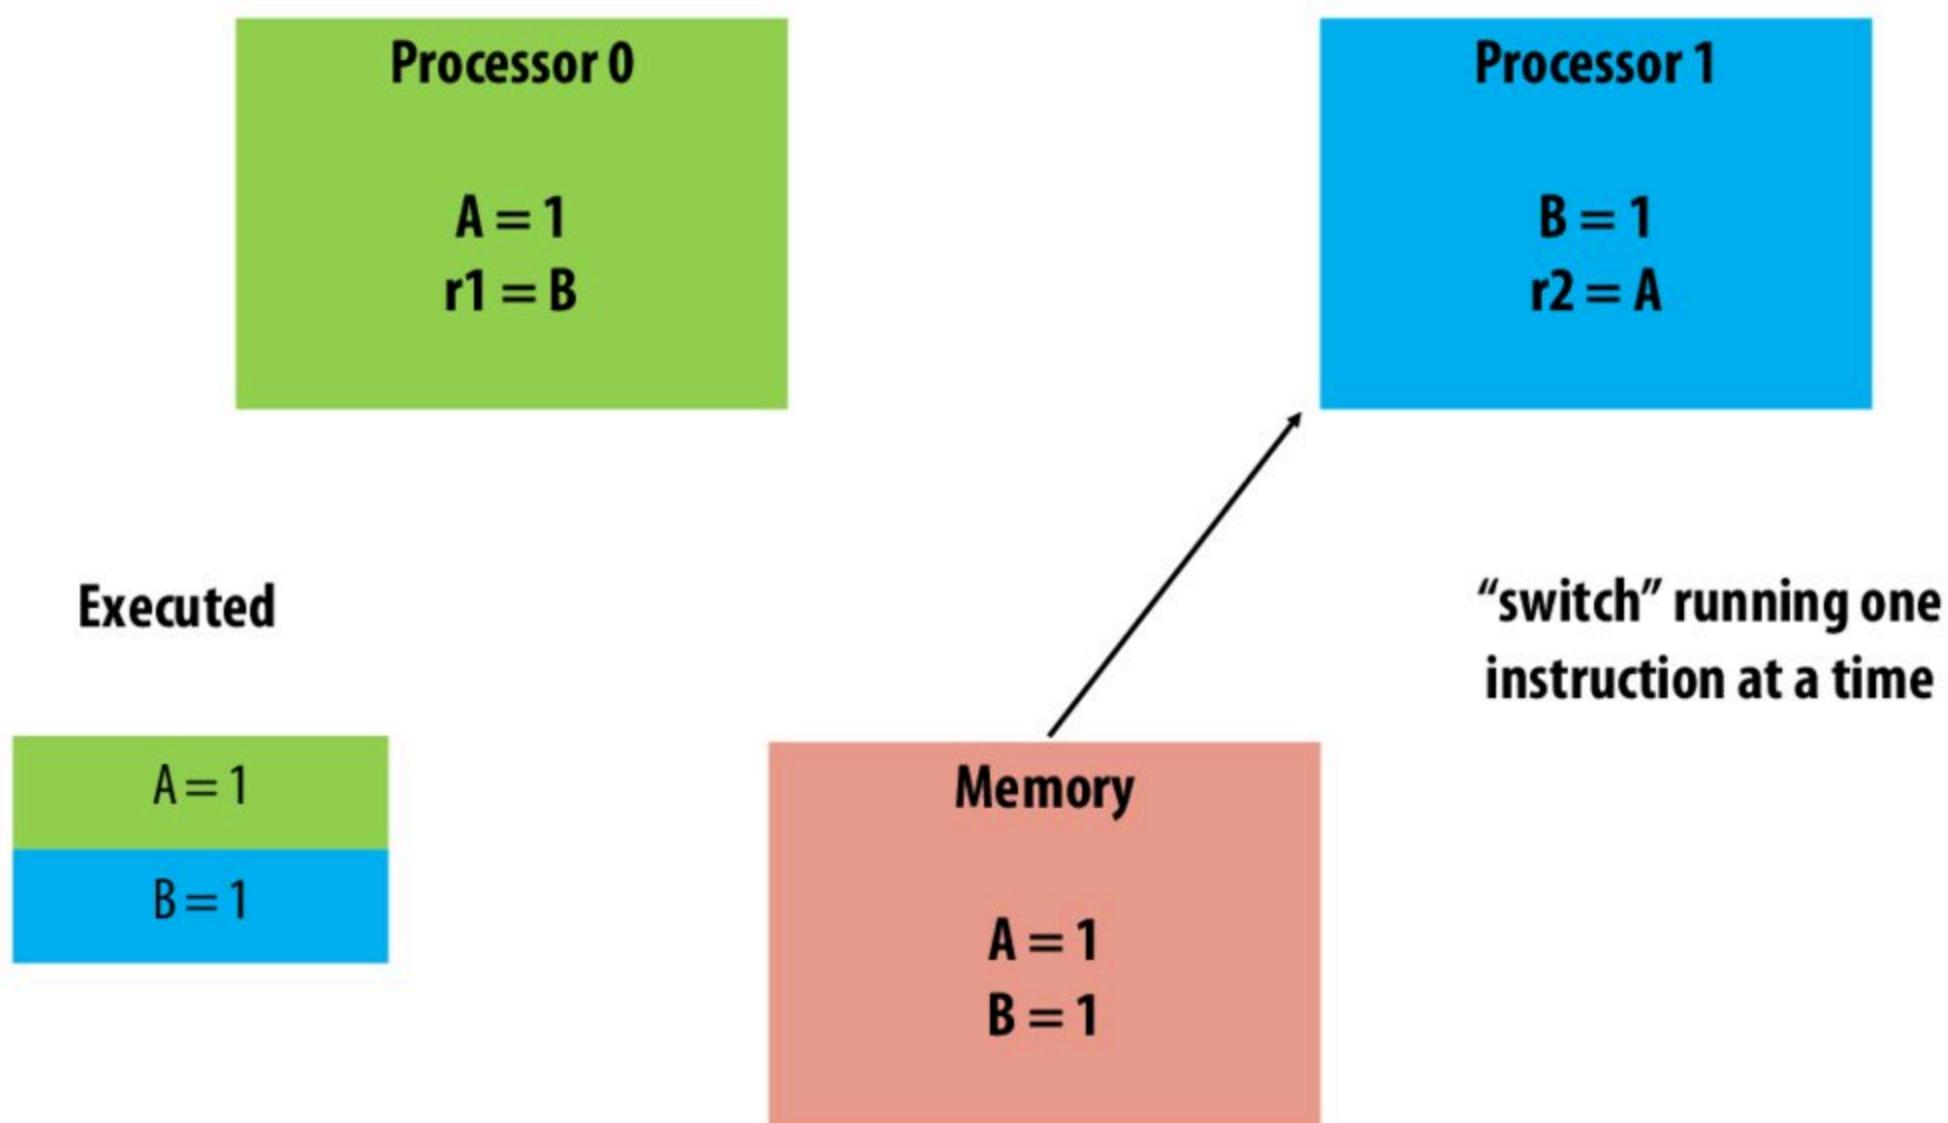
Sequential Consistency Example



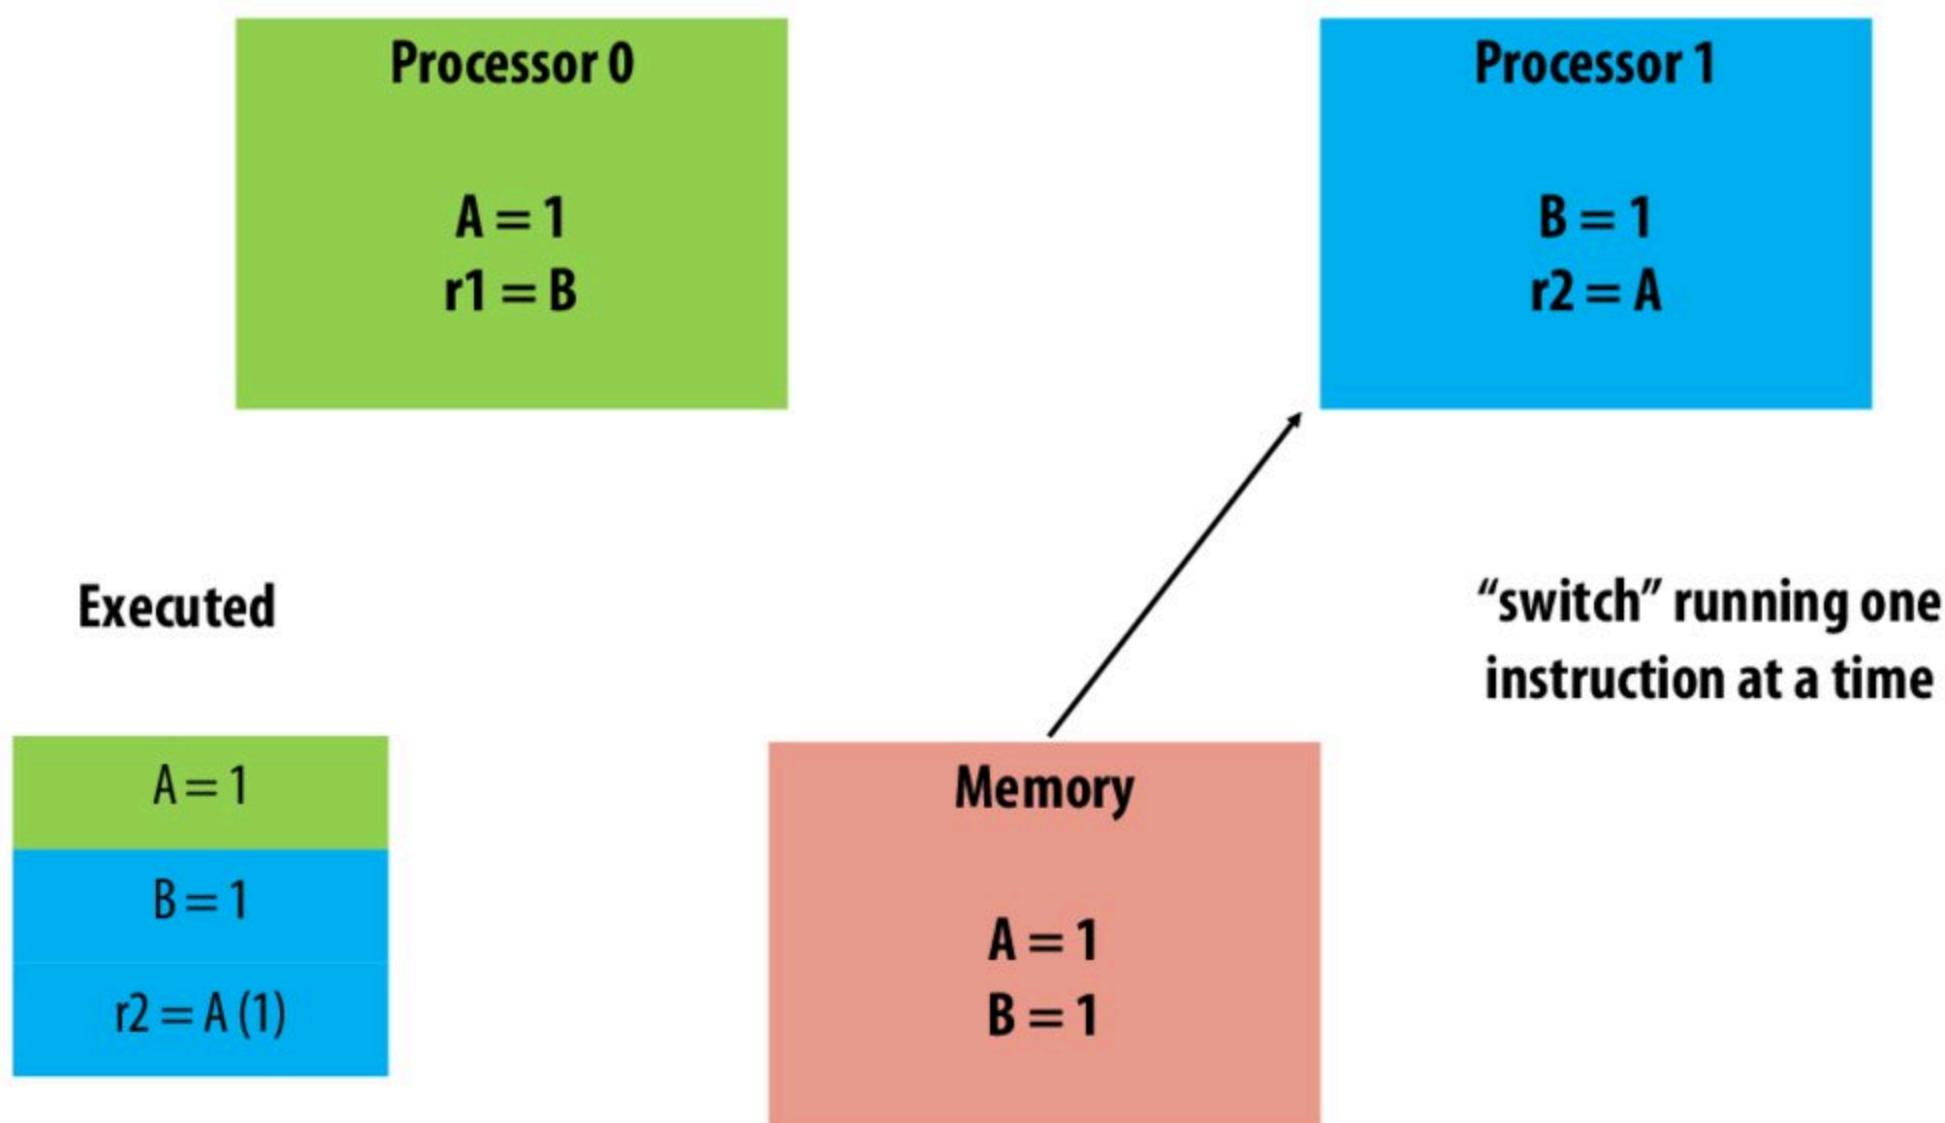
Sequential Consistency Example



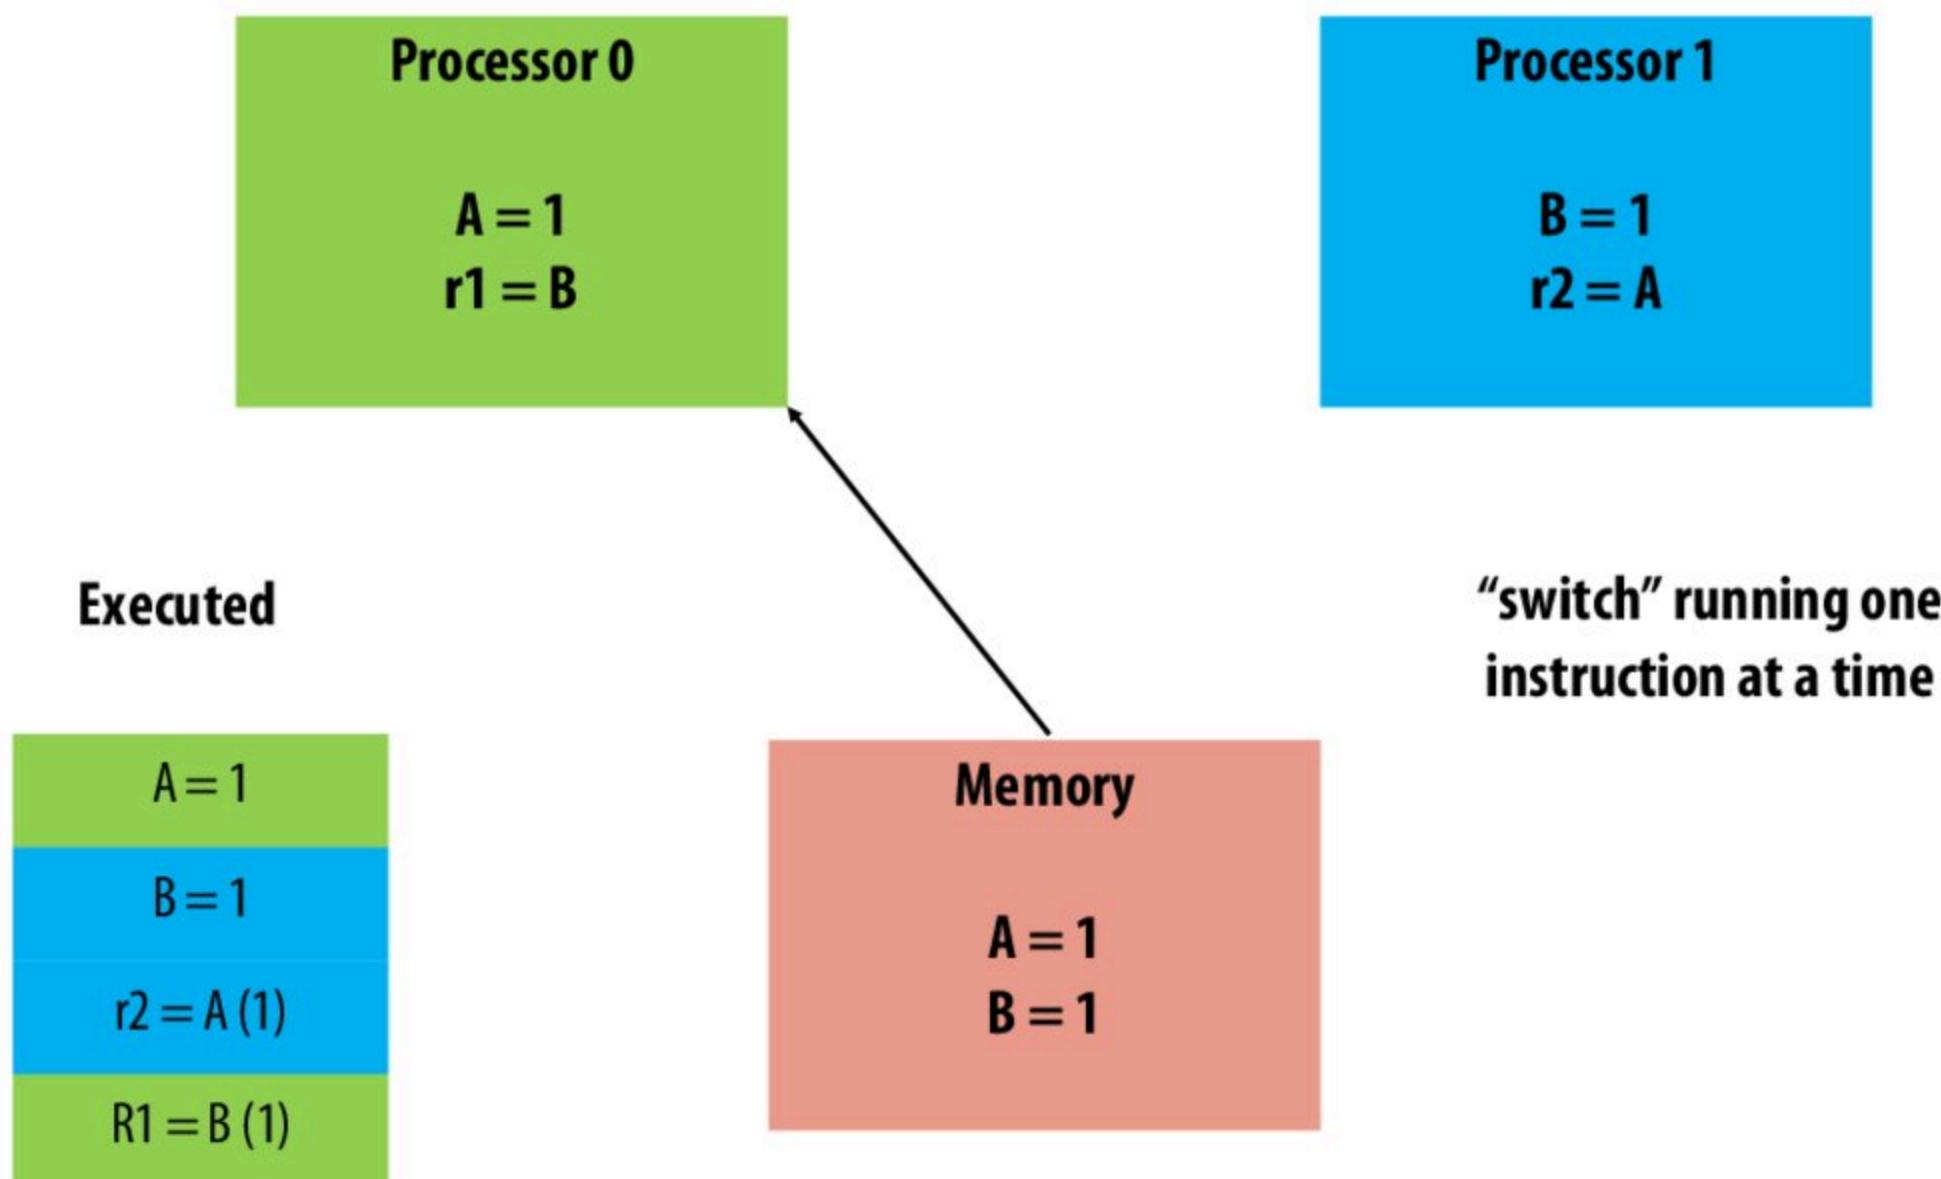
Sequential Consistency Example



Sequential Consistency Example



Sequential Consistency Example



Relaxing memory operation ordering

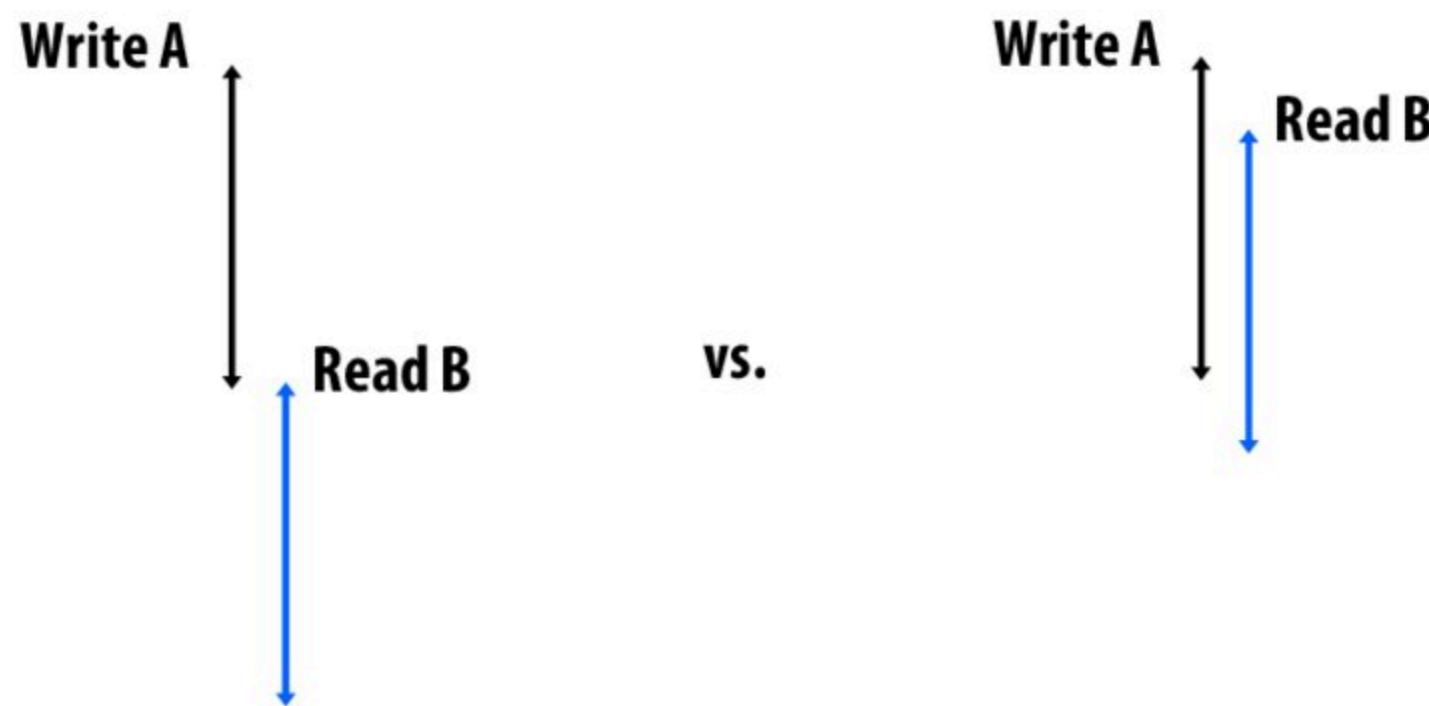
A sequentially consistent memory system maintains all four memory operation orderings ($W_X \rightarrow R_Y$, $R_X \rightarrow R_Y$, $R_X \rightarrow W_Y$, $W_X \rightarrow W_Y$)

Relaxed memory consistency models allow certain orderings to be violated

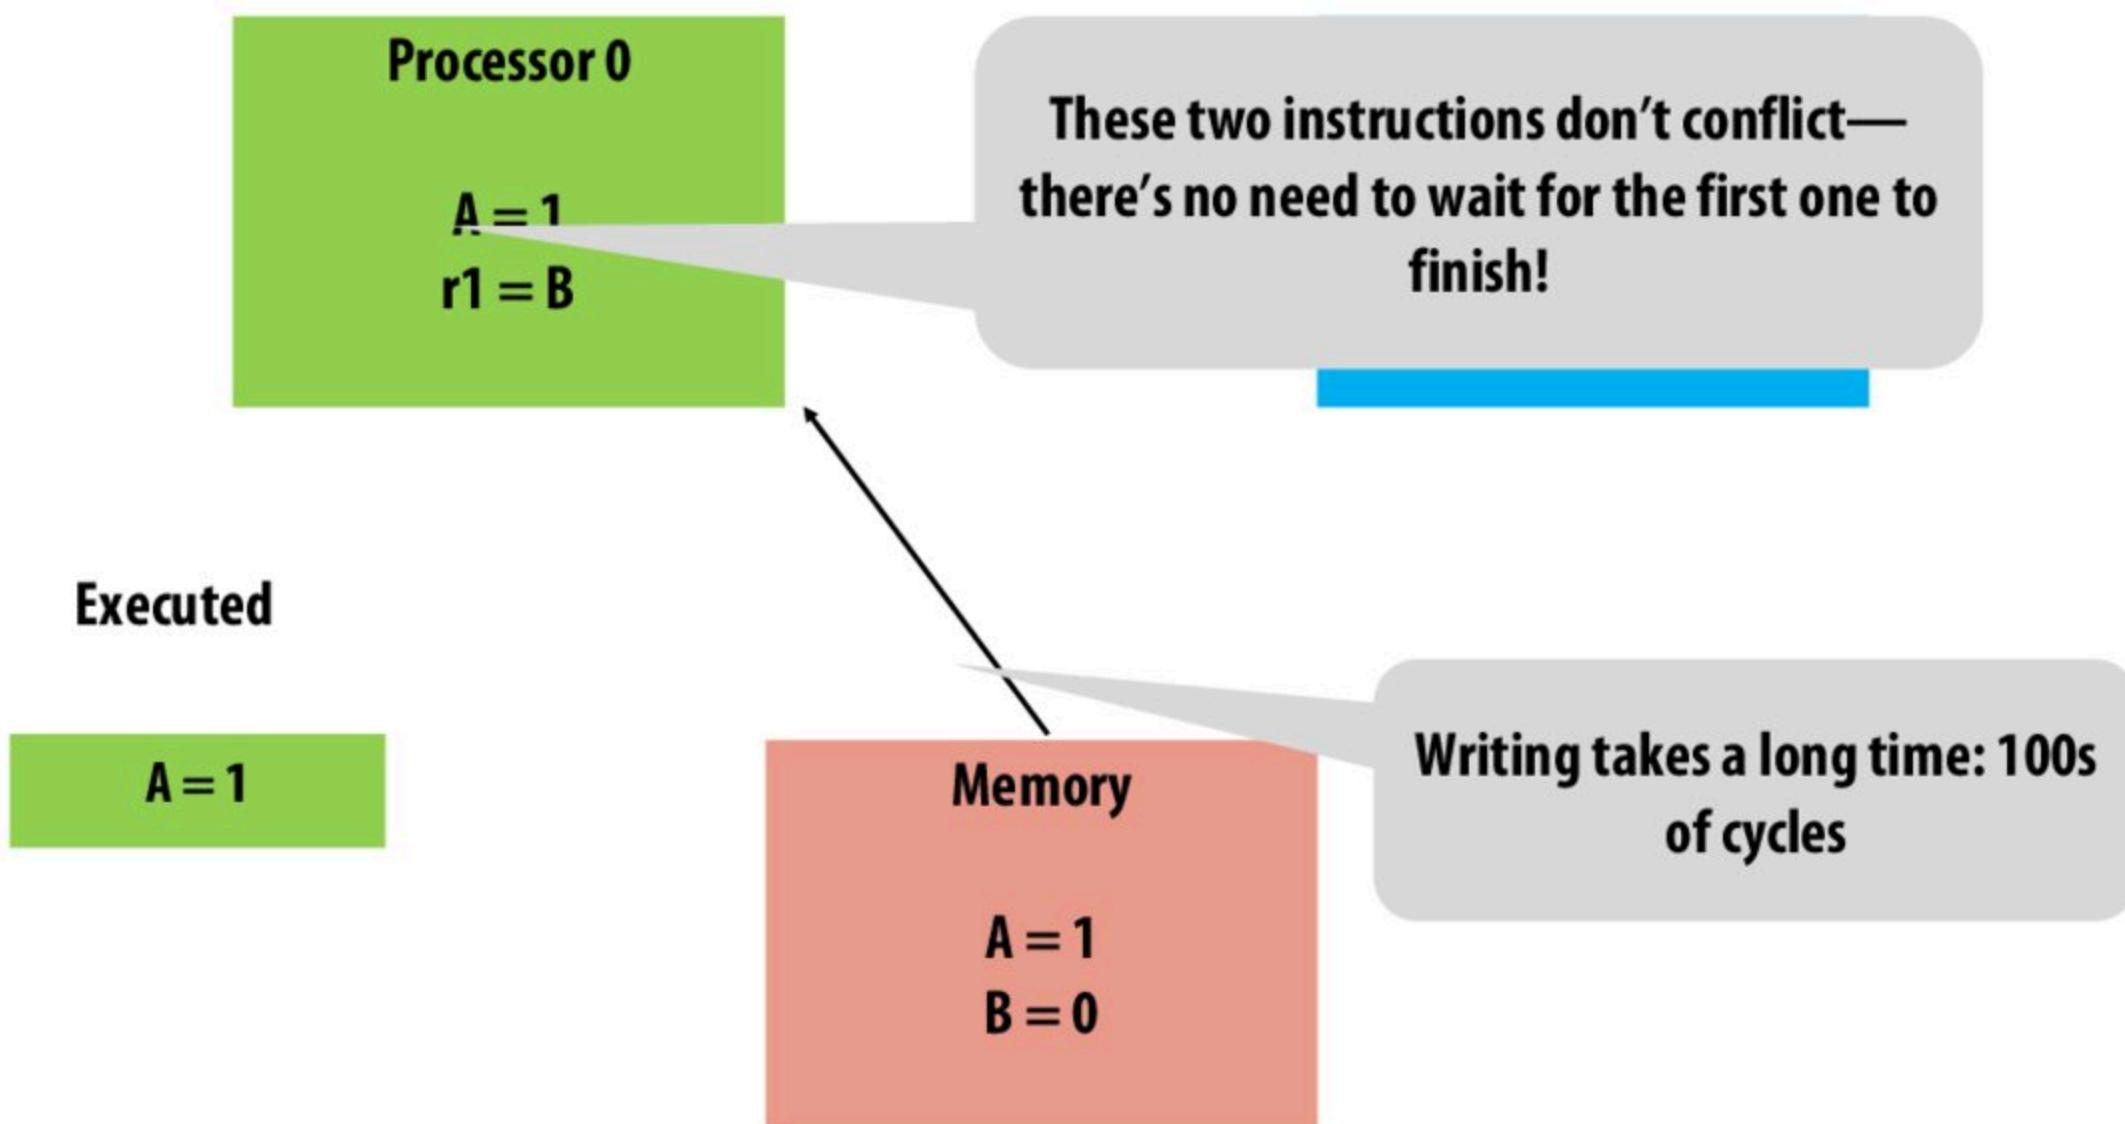
Motivation for relaxed consistency: hiding latency

Why are we interested in relaxing ordering requirements?

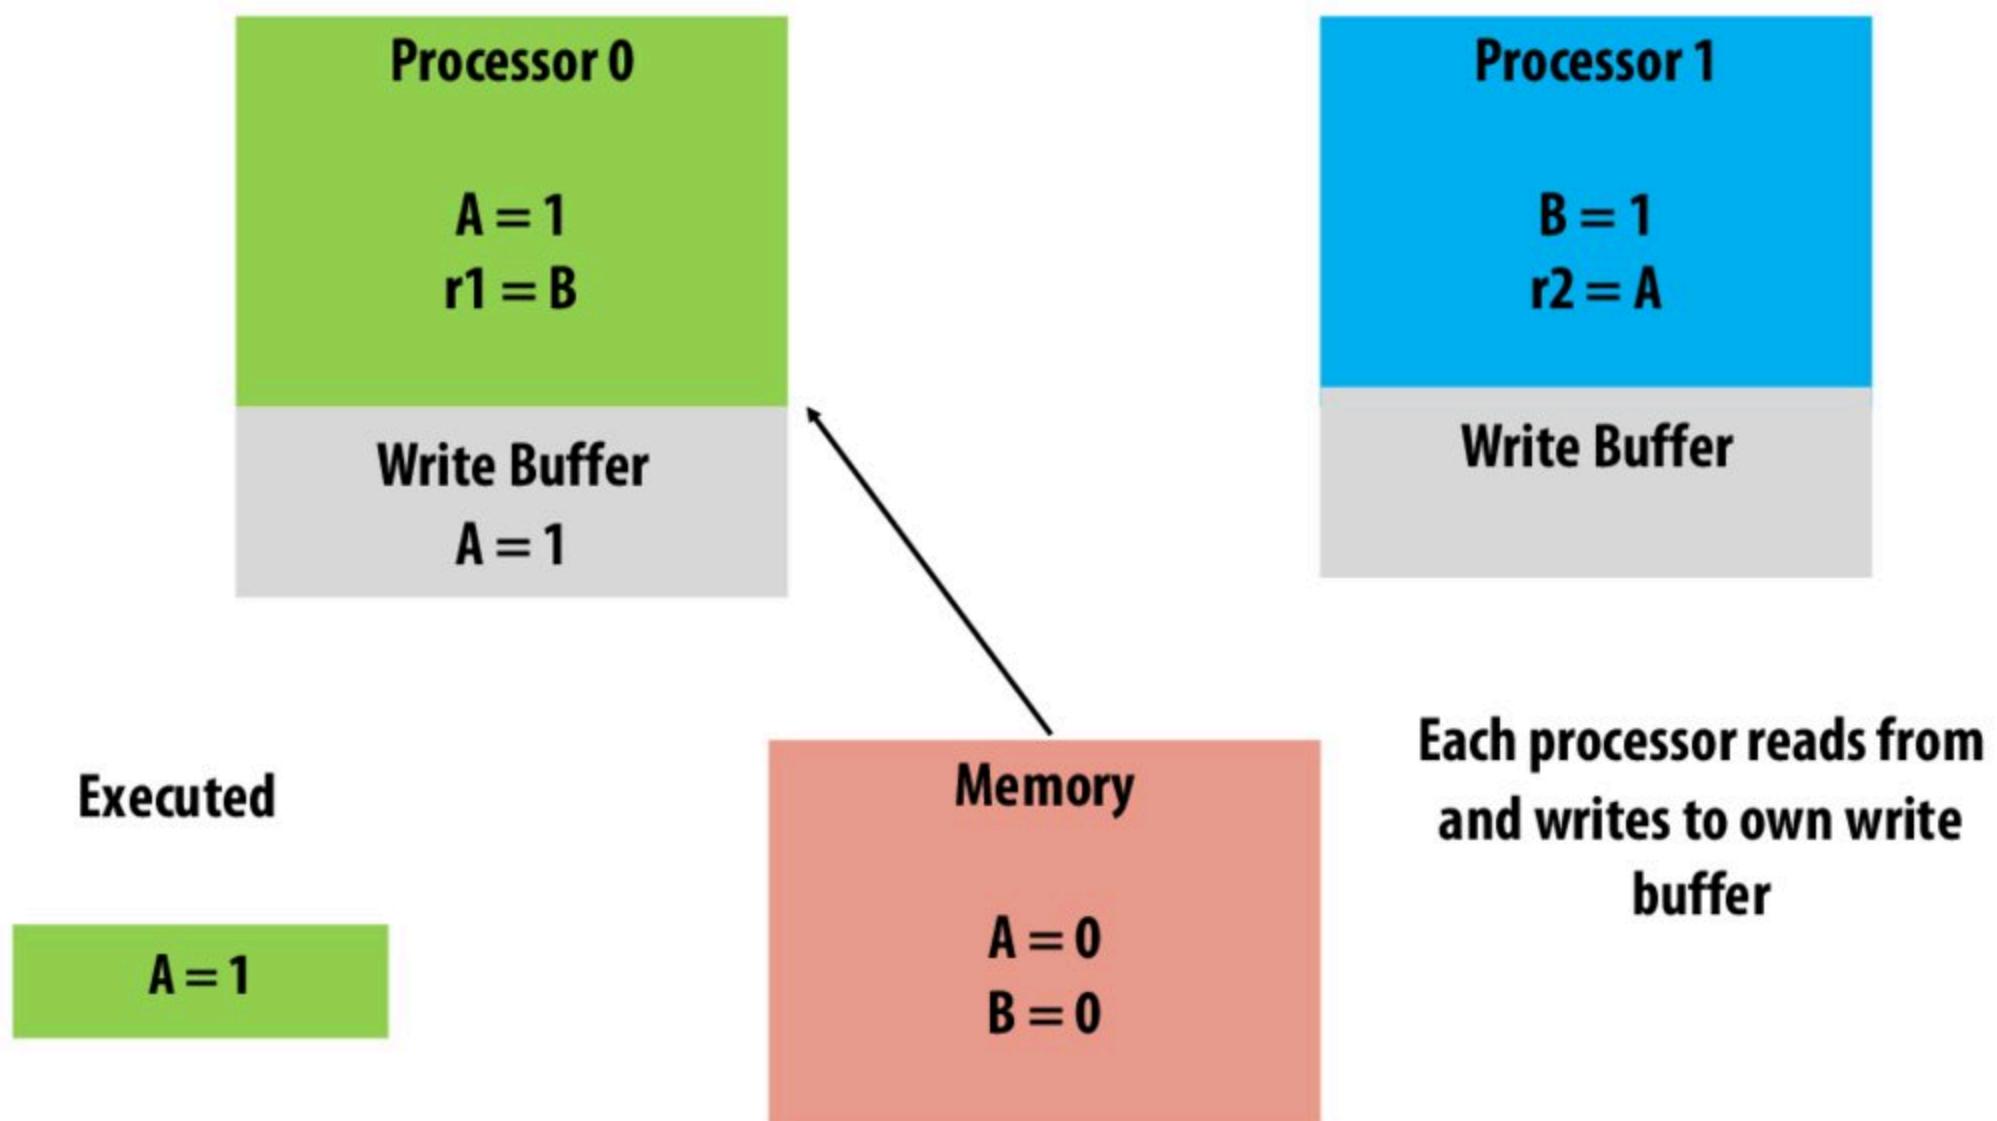
- To gain performance
- Specifically, hiding memory latency: overlap memory access operations with other operations when they are independent
- Remember, memory access in a cache coherent system may entail much more work than simply reading bits from memory (finding data, sending invalidations, etc.)



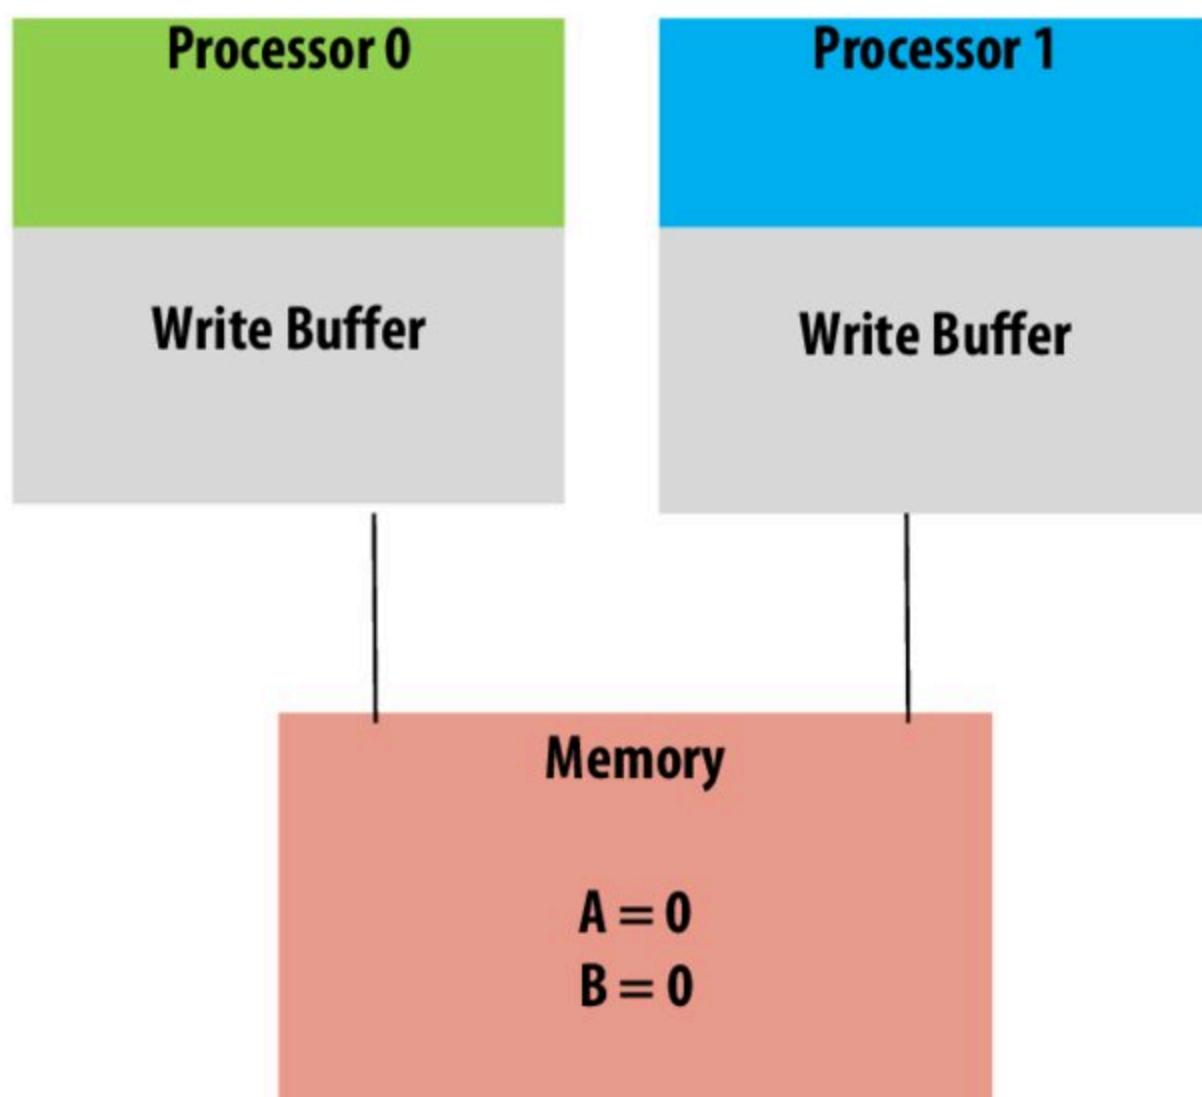
Problem with SC



Optimization: Write Buffer



Write Buffers Change Memory Behavior



Initially $A = B = 0$

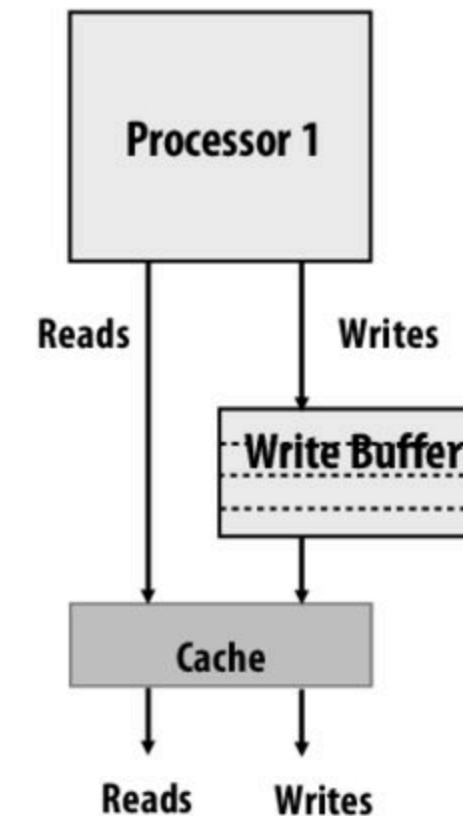
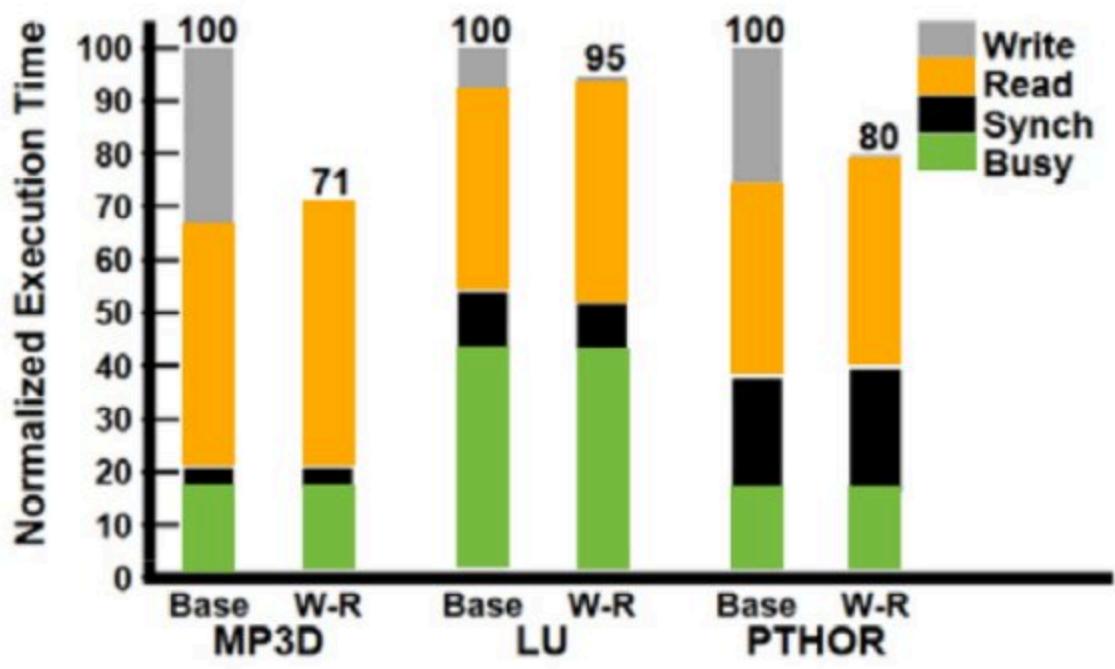
Proc 0	Proc 1
(1) $A = 1$	(3) $B = 1$
(2) $r1 = B$	(4) $r2 = A$

Can $r1 = r2 = 0$?

SC: No

Write buffers:

Write Buffer Performance



Base: Sequentially consistent execution. Processor issues one memory operation at a time, stalls until completion

W-R: relaxed W→R ordering constraint (write latency almost fully hidden)

Write Buffers: Who Cares?

Performance improvement

Every modern processor uses them

- **Intel x86, ARM, SPARC**

Need a weaker memory model

- **TSO: Total Store Order**
- **Slightly harder to reason about than SC**
- **x86 uses an incompletely specified form of TSO**

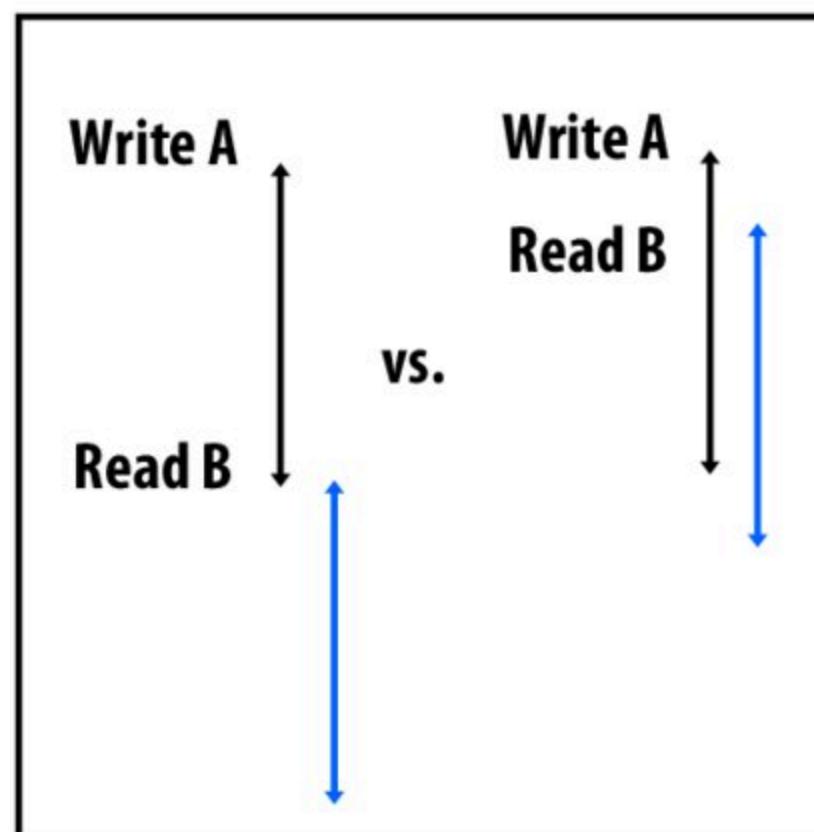
Allowing reads to move ahead of writes

Four types of memory operation orderings

- ~~$W_x \rightarrow R_y$: write must complete before subsequent read~~
- $R_x \rightarrow R_y$: read must complete before subsequent read
- $R_x \rightarrow W_y$: read must complete before subsequent write
- $W_x \rightarrow W_y$: write must complete before subsequent write

Allow processor to hide latency of writes

- **Total Store Ordering (TSO)**
- **Processor Consistency (PC)**



Allowing reads to move ahead of writes

Total store ordering (TSO)

- Processor P can read B before its write to A is seen by all processors
(processor can move its own reads in front of its own writes)
- Reads by other processors cannot return new value of A until the write to A is observed by all processors

Processor consistency (PC)

- Any processor can read new value of A before the write is observed by all processors

In TSO and PC, only $W_x \rightarrow R_y$ order is relaxed. The $W_x \rightarrow W_y$ constraint still exists. Writes by the same thread are not reordered (they occur in program order)

Clarification (make sure you get this!)

The cache coherency problem exists because hardware implements the optimization of duplicating data in multiple processor caches. The copies of the data must be kept coherent.

Relaxed memory consistency issues arise from the optimization of reordering memory operations. (Consistency is unrelated to whether or not caches exist in the system)

Allowing writes to be reordered

Four types of memory operation orderings

- ~~$W_x \rightarrow R_y$: write must complete before subsequent read~~
- $R_x \rightarrow R_y$: read must complete before subsequent read
- $R_x \rightarrow W_y$: read must complete before subsequent write
- ~~$W_x \rightarrow W_y$: write must complete before subsequent write~~

Partial Store Ordering (PSO)

- Execution may not match sequential consistency on program 1
(P2 may observe change to flag before change to A)

Thread 1 (on P1)

```
A = 1;  
flag = 1;
```

Thread 2 (on P2)

```
while (flag == 0);  
print A;
```

Why might it be useful to allow more aggressive memory operation reorderings?

$W_X \rightarrow W_Y$: processor might reorder write operations in a write buffer (e.g., one is a cache miss while the other is a hit)

$R_X \rightarrow W_Y, R_X \rightarrow R_Y$: processor might reorder independent instructions in an instruction stream (out-of-order execution)

Keep in mind these are all valid optimizations if a program consists of a single instruction stream

Allowing all reorderings

Four types of memory operation orderings

- ~~$W_x \rightarrow R_y$: write must complete before subsequent read~~
- ~~$R_x \rightarrow R_y$: read must complete before subsequent read~~
- ~~$R_x \rightarrow W_y$: read must complete before subsequent write~~
- ~~$W_x \rightarrow W_y$: write must complete before subsequent write~~

No guarantees about operations on data!

- **Everything can be reordered**

Motivation is increased performance

- **Overlap multiple reads and writes in the memory system**
- **Execute reads as early as possible and writes as late as possible to hide memory latency**

Examples:

- **Weak ordering (WO)**
- **Release Consistency (RC)**

Synchronization to the Rescue

Memory reordering seems like a nightmare (it is!)

Every architecture provides synchronization primitives to make memory ordering stricter

reorderable reads
and writes here

...

MEMORY FENCE

...

reorderable reads
and writes here

...

MEMORY FENCE

Fence (memory barrier) instructions prevent reorderings, but are expensive

- All memory operations complete before any memory operation after it can begin

Other synchronization primitives (per address):

- read-modify-write/compare-and-swap, transactional memory, ...

Example: expressing synchronization in relaxed models

Intel x86/x64 ~ total store ordering

- Provides sync instructions if software requires a specific instruction ordering not guaranteed by the consistency model**
 - mm_lfence ("load fence": wait for all loads to complete)
 - mm_sfence ("store fence": wait for all stores to complete)
 - mm_mfence ("mem fence": wait for all me operations to complete)

ARM processors: very relaxed consistency model

A cool post on the role of memory fences in x86:

<http://bartoszmilewski.com/2008/11/05/who-ordered-memory-fences-on-an-x86/>

ARM has some great examples in their programmer's reference:

http://infocenter.arm.com/help/topic/com.arm.doc.genc007826/Barrier_Litmus_Tests_and_Cookbook_A08.pdf

A great list of academic papers:

<http://www.cl.cam.ac.uk/~pes20/weakmemory/>

Problem: Data Races

Every example so far has involved a data race

- **Two accesses to the same memory location**
- **At least one is a write**
- **Unordered by synchronization operations**

Conflicting data accesses

Two memory accesses by different processors conflict if...

- They access the same memory location
- At least one is a write

Unsynchronized program

- Conflicting accesses not ordered by synchronization (e.g., a fence, operation with release/acquire semantics, barrier, etc.)
- Unsynchronized programs contain data races: the output of the program depends on relative speed of processors (non-deterministic program results)

Synchronized Programs

Synchronized programs yield SC results on non-SC systems

- **Synchronized programs are data-race-free**

If there are no data races, reordering behavior doesn't matter

- **Accesses are ordered by synchronization, and synchronization forces sequential consistency**

In practice, most programs you encounter will be synchronized (via locks, barriers, etc. implemented in synchronization libraries)

- **Rather than via ad-hoc reads/writes to shared variables like in the example programs**

Summary: Relaxed Consistency

Motivation: obtain higher performance by allowing reordering of memory operations (reordering is not allowed by sequential consistency)

One cost is software complexity: programmer or compiler must correctly insert synchronization to ensure certain specific operation orderings when needed

- But in practice complexities encapsulated in libraries that provide intuitive primitives like lock/unlock, barrier (or lower-level primitives like fence)
- Optimize for the common case: most memory accesses are not conflicting, so don't design a system that pays the cost as if they are

Relaxed consistency models differ in which memory ordering constraints they ignore

Languages Need Memory Models Too

Thread 1

```
X = 0  
for i=0 to 100:  
    X = 1  
    print X
```

compiler

Thread 1

```
X = 1  
for i=0 to 100:  
    print X
```

Languages Need Memory Models Too

Optimization not visible to programmer

Thread 1

```
X = 0  
for i=0 to 100:  
    X = 1  
    print X
```

11111111111...

Thread 1

```
X = 1  
for i=0 to 100:  
    print X
```

11111111111...

Languages Need Memory Models Too

Optimization is visible to programmer

Thread 1

```
X = 0  
for i=0 to 100:  
    X = 1  
    print X
```

111111111111...

111110111111...

Thread 2

```
X = 0
```

Thread 1

```
X = 1  
for i=0 to 100:  
    print X
```

111111111111...

111110000000...

Thread 2

```
X = 0
```

Provide a contract to programmers about how their memory operations will be reordered by the compiler e.g. no reordering of shared memory operations

Language Level Memory Models

Modern (C11, C++11) and not-so-modern (Java 5) languages guarantee sequential consistency for data-race-free programs ("SC for DRF")

- **Compilers will insert the necessary synchronization to cope with the hardware memory model**

No guarantees if your program contains data races!

- **The intuition is that most programmers would consider a racy program to be buggy**

Use a synchronization library!

Memory Consistency Models Summary

Define the allowed reorderings of memory operations by hardware and compilers

A contract between hardware or compiler and application software

Weak models required for good performance?

- SC can perform well with many more resources**

Details of memory model can be hidden in synchronization library

- Requires data race free (DRF) programs**