

Lecture 18:

Transactional Memory II

Parallel Computing
Stanford CS149, Fall 2024

Transactional Memory (TM) Review

Memory transaction

- An atomic and isolated sequence of memory accesses
- Inspired by database transactions

Atomicity (all or nothing)

- Upon transaction commit, all memory writes in transaction take effect at once
- On transaction abort, none of the writes appear to take effect (as if transaction never happened)

Isolation

- No other processor can observe writes before transaction commits

Serializability

- Transactions appear to commit in a single serial order
- But the exact order of commits is not guaranteed by semantics of transaction

Advantages (promise) of transactional memory

Easy to use synchronization construct

- It is difficult for programmers to get synchronization right
- Programmer declares need for atomicity, system implements it well
- Claim: transactions are as easy to use as coarse-grain locks

Often performs as well as fine-grained locks

- Provides automatic read-read concurrency and fine-grained concurrency
- Performance portability: locking scheme for four CPUs may not be the best scheme for 64 CPUs
- Productivity argument for transactional memory: system support for transactions can achieve 90% of the benefit of expert programming with fine-grained locks, with 10% of the development time

Failure atomicity and recovery

- No lost locks when a thread fails
- Failure recovery = transaction abort + restart

Composability

- Safe and scalable composition of software modules

Implementing transactional memory

TM implementation basics

TM systems must provide atomicity and isolation

- While maintaining concurrency as much as possible

Two key implementation questions

- Data versioning policy: How does the system manage uncommitted (new) and previously committed (old) versions of data for concurrent transactions?
 - Eager versioning (undo-log based)
 - Lazy versioning (write-buffer based)
- Conflict detection policy: how/when does the system determine that two concurrent transactions conflict?
 - Pessimistic detection: check on every memory access
 - Optimistic detection: check on transaction commit

TM implementation space (examples)

Software TM systems

- Lazy + optimistic (rd/wr): Sun TL2
- Lazy + optimistic (rd)/pessimistic (wr): MS OSTM
- Eager + optimistic (rd)/pessimistic (wr): Intel STM
- Eager + pessimistic (rd/wr): Intel STM

Hardware TM systems

- Lazy + optimistic: Stanford TCC
- Lazy + pessimistic: MIT LTM, Intel VTM
- Eager + pessimistic: Wisconsin LogTM (easiest with conventional cache coherence)

Optimal design remains an open question

- May be different for HW, SW, and hybrid

Software Transactional Memory

```
atomic {  
    a.x = t1  
    a.y = t2  
    if (a.z == 0) {  
        a.x = 0  
        a.z = t3  
    }  
}
```



```
tmTxnBegin()  
tmWr(&a.x, t1)  
tmWr(&a.y, t2)  
if (tmRd(&a.z) != 0) {  
    tmWr(&a.x, 0);  
    tmWr(&a.z, t3)  
}  
tmTxnCommit()
```

- Software barriers (STM function call) for TM bookkeeping
 - Versioning, read/write-set tracking, commit, ...
 - Using locks, timestamps, data copying, ...
- Requires function cloning or dynamic translation
 - Function used inside and outside of transaction

STM Runtime Data Structures

Transaction descriptor (per-thread)

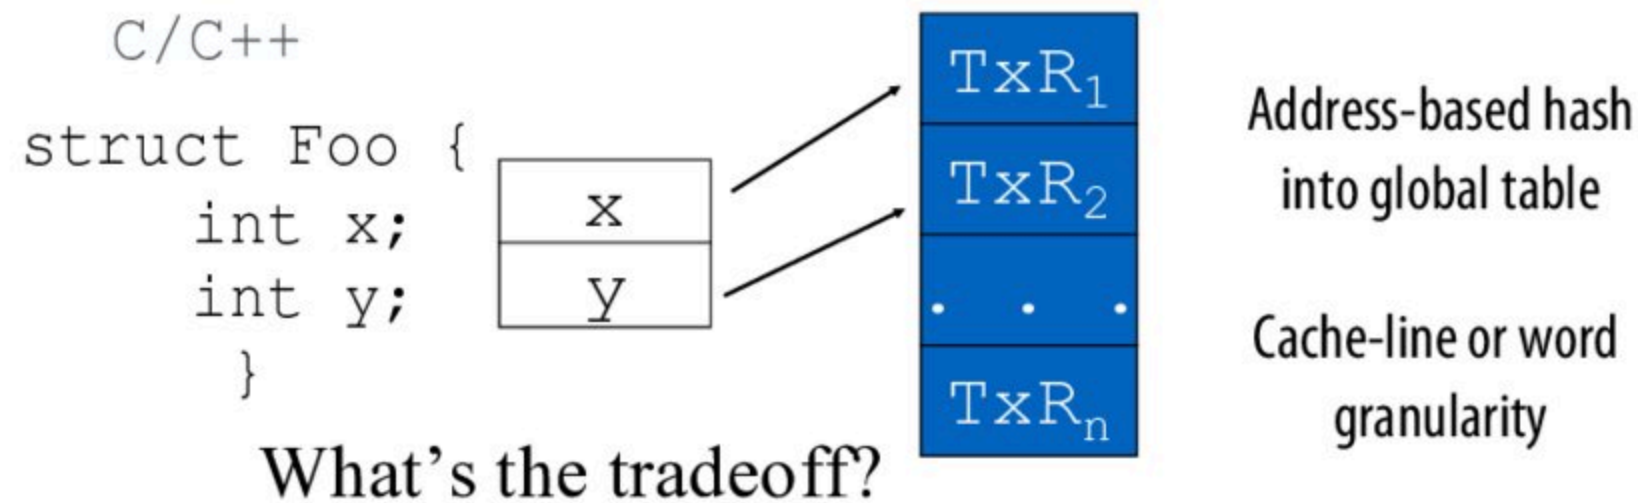

- Used for conflict detection, commit, abort, ...
- Includes the read set, write set, undo log or write buffer

Transaction record (per data)

- Pointer-sized record guarding shared data
- Tracks transactional state of data
 - Shared: accessed by multiple readers
 - Using version number or shared reader lock
 - Exclusive: access by one writer
 - Using writer lock that points to owner
 - BTW: same way that HW cache coherence works

Mapping Data to Transaction Records

Every data item has an associated transaction record



Conflict Detection Granularity

Object granularity

- Low overhead mapping operation
- Exposes optimization opportunities
- False conflicts (e.g. Txn 1 and Txn 2)

Element/field granularity (word)


- Reduces false conflicts
- Improves concurrency (e.g. Txn 1 and Txn 2)
- Increased overhead (time/space)

Cache line granularity (multiple words)

- Matches hardware TM
- Reduces storage overhead of transactional records
- Hard for programmer & compiler to analyze

Mix & match per type basis

- E.g., element-level for arrays, object-level for non-arrays



An Example STM Algorithm

Based on Intel's McRT STM [PPoPP' 06, PLDI' 06, CGO' 07]

- Eager versioning, optimistic reads, pessimistic writes

Based on timestamp for version tracking

- Global timestamp
 - Incremented when a writing xaction commits
- Local timestamp per xaction
 - Global timestamp value when xaction last validated

Transaction record (32-bit)

- LS bit: 0 if writer-locked, 1 if not locked
- MS bits
 - Timestamp (version number) of last commit if not locked
 - Pointer to owner xaction if locked

STM Operations

STM read (optimistic)

- Direct read of memory location (eager)
- Validate read data
 - Check if unlocked and data version \leq local timestamp
 - If not, validate all data in read set for consistency
- Insert in read set
- Return value

STM write (pessimistic)

- Validate data
 - Check if unlocked and data version \leq local timestamp
- Acquire lock
- Insert in write set
- Create undo log entry
- Write data in place (eager)

STM Operations (cont)


Read-set validation

- Get global timestamp
- For each item in the read set
 - If locked by other or data version $>$ local timestamp, abort
- Set local timestamp to global timestamp from initial step

STM commit

- Atomically increment global timestamp by 2 (LSb used for write-lock)
- If preincremented (old) global timestamp $>$ local timestamp, validate read-set
 - Check for recently committed transactions
- For each item in the write set
 - Release the lock and set version number to global timestamp

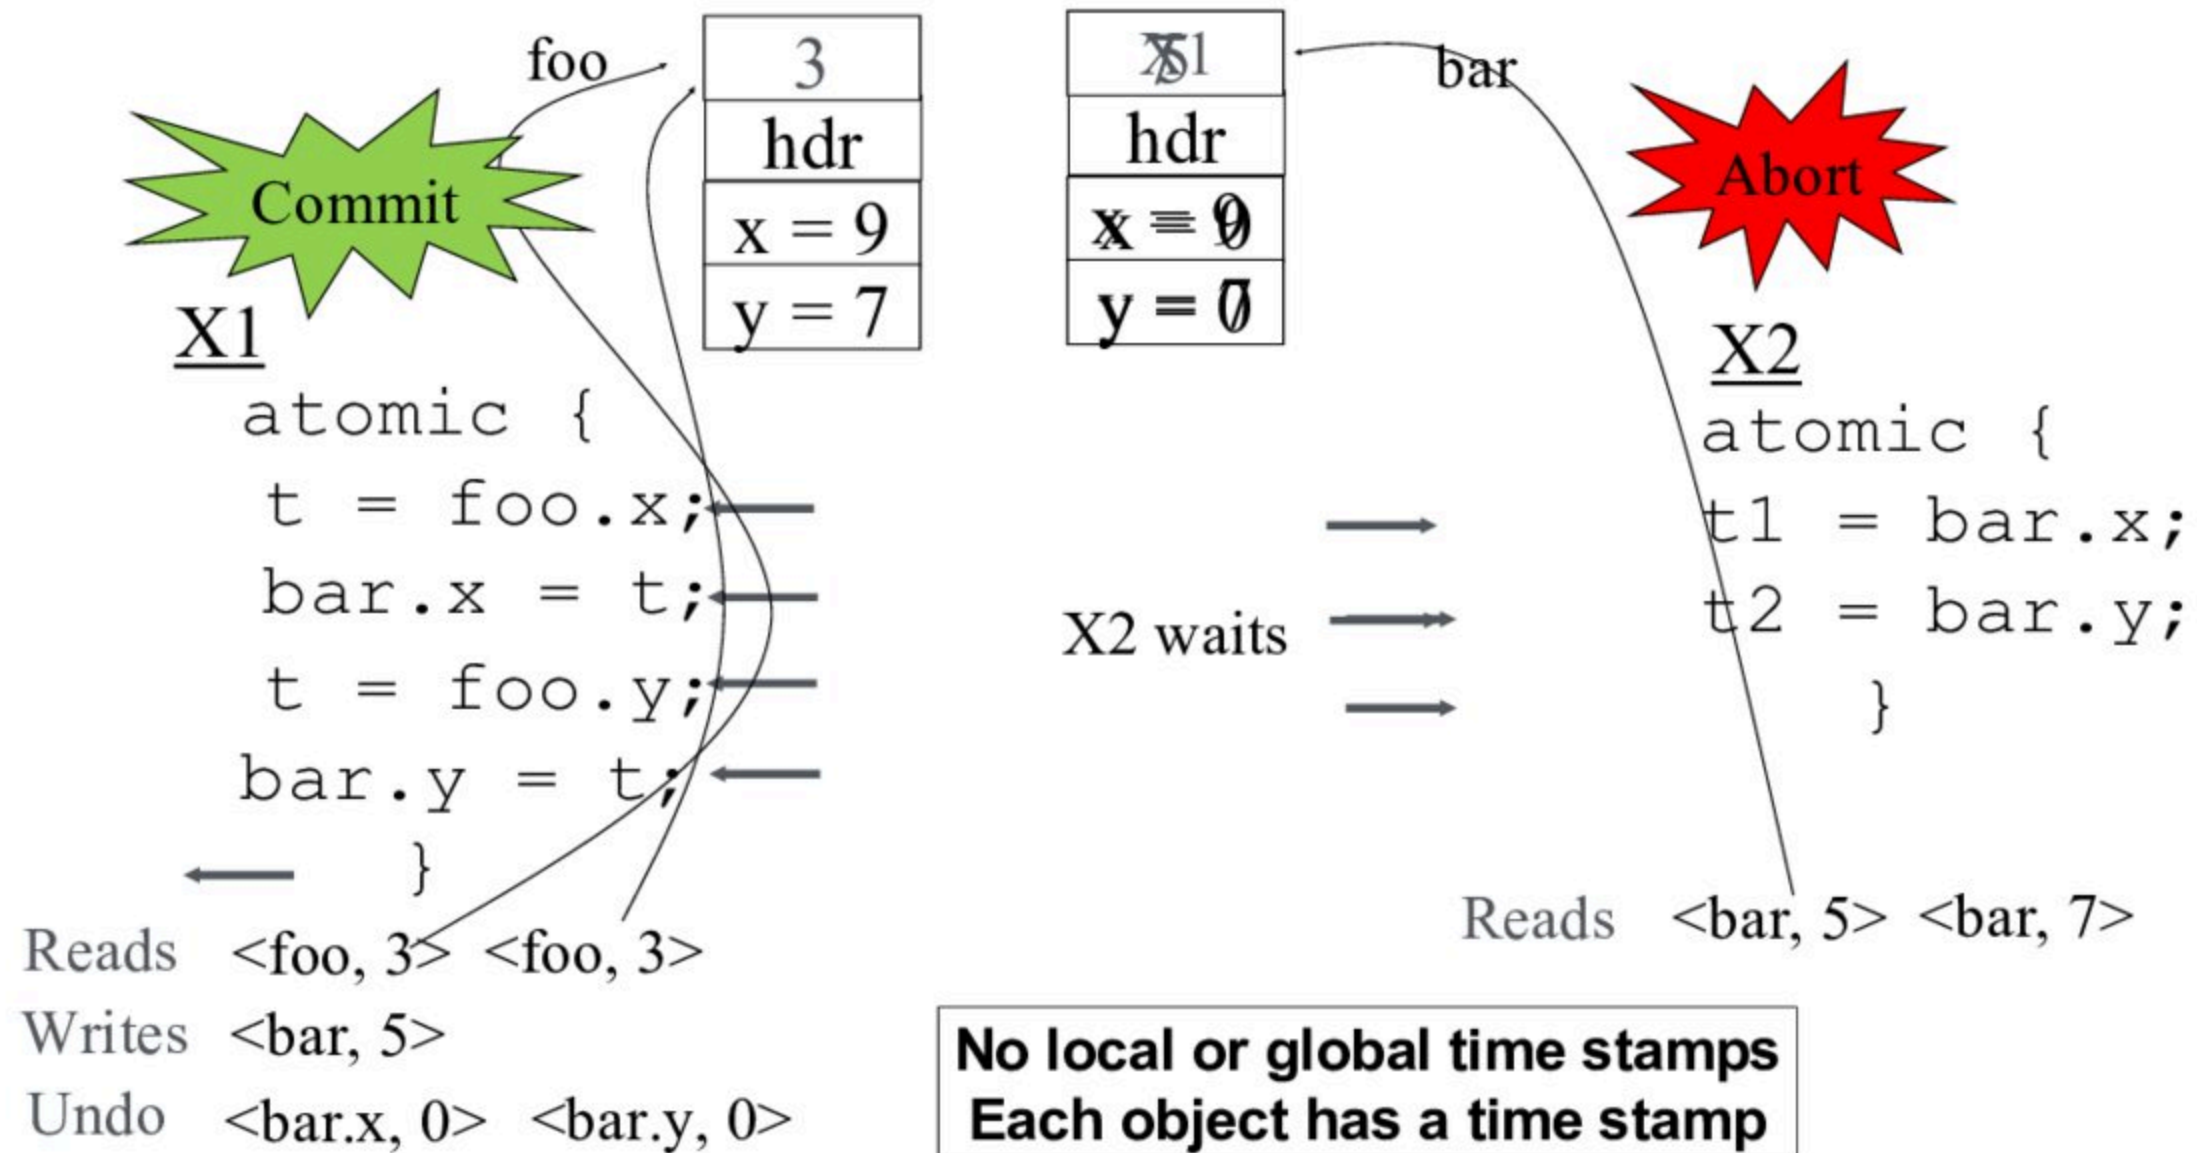
STM Example



X1 copies object foo into object bar

X2 should read bar as [0,0] or [9,7]

STM Example



TM Implementation Summary 1

TM implementation


- Data versioning: eager or lazy
- Conflict detection: optimistic or pessimistic
 - Granularity: object, word, cache-line, ...

Software TM systems

- Compiler adds code for versioning & conflict detection
 - Note: STM barrier = instrumentation code
- Basic data-structures
 - Transactional descriptor per thread (status, rd/wr set, ...)
 - Transactional record per data (locked/version)

Effect of Compiler Optimizations


1 thread overheads over thread-unsafe baseline



With compiler optimizations

- **<40% over no concurrency control**
- **<30% over lock-based synchronization**

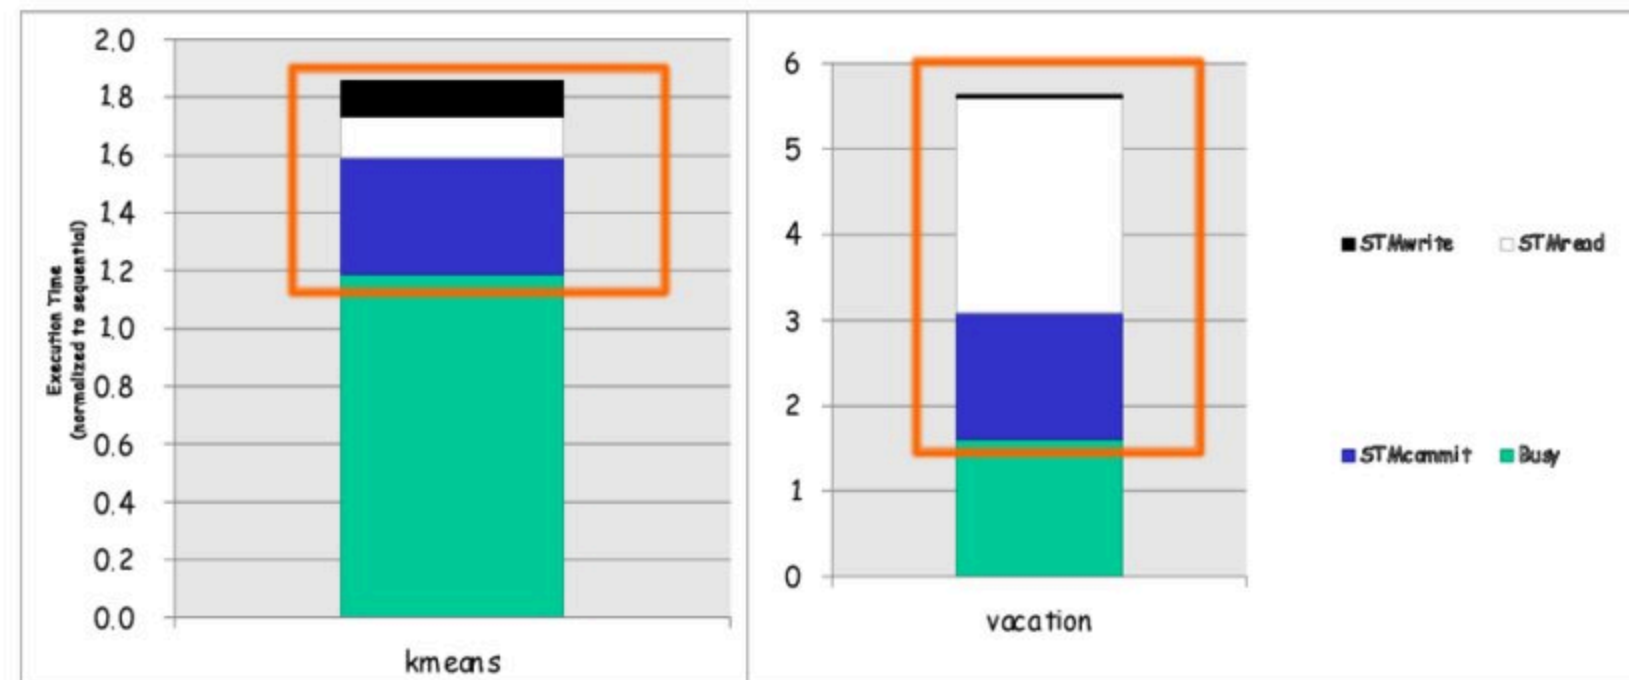
Motivation for Hardware Support



- STM slowdown: 2-8x per thread overhead due to barriers
 - Short term issue: demotivates parallel programming
 - Long term issue: energy wasteful
- Lack of strong atomicity
 - Costly to provide purely in software

Why is STM Slow?

Measured single-thread STM performance



1.8x – 5.6x slowdown over sequential

Most time goes in read barriers & commit

- Most apps read more data than they write

Types of Hardware Support

Hardware-accelerated STM systems (HASTM, SigTM, USTM, ...)

- Start with an STM system & identify key bottlenecks
- Provide (simple) HW primitives for acceleration, but keep SW barriers

Hardware-based TM systems (TCC, LTM, VTM, LogTM, ...)

- Versioning & conflict detection directly in HW
- No SW barriers

Hybrid TM systems (Sun Rock, ...)

- Combine an HTM with an STM by switching modes when needed
 - Based on xaction characteristics available resources, ...

	HTM	STM	HW-STM
Write versioning	HW	SW	SW
Conflict detection	HW	SW	HW

Hardware transactional memory (HTM)

Data versioning is implemented in caches

- Cache the write buffer or the undo log
- Add new cache line metadata to track transaction read set and write set

Conflict detection through cache coherence protocol

- Coherence lookups detect conflicts between transactions
- Works with snooping and directory coherence

Note:

- Register checkpoint must also be taken at transaction begin (to restore execution context state on abort)

HTM design

Cache lines annotated to track read set and write set

- R bit: indicates data read by transaction (set on loads)
- W bit: indicates data written by transaction (set on stores)
 - R/W bits can be at word or cache-line granularity
- R/W bits gang-cleared on transaction commit or abort

MESI state bit for line (e.g., M state)

This illustration tracks read and write set at cache line granularity

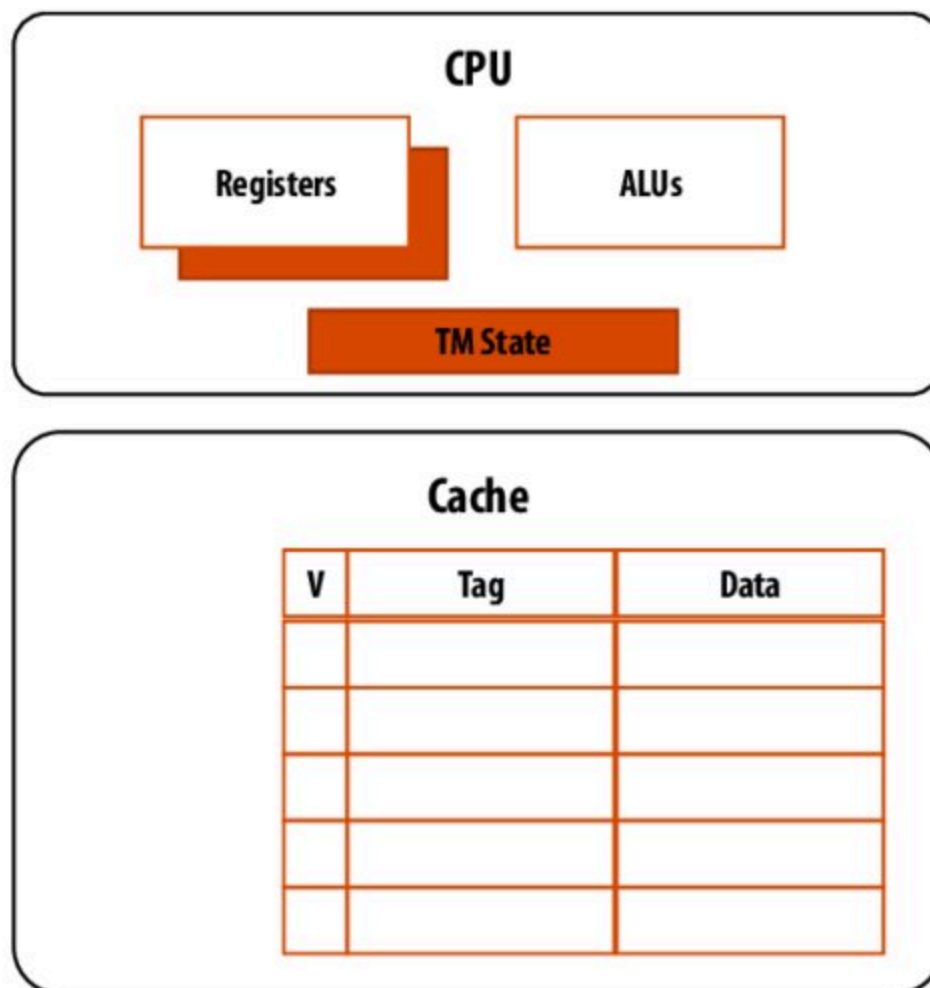


- For eager versioning, need a 2nd cache write for undo log

Coherence requests check R/W bits to detect conflicts

- Observing shared request to W-word is a read-write conflict
- Observing exclusive (intent to write) request to R-word is a write-read conflict
- Observing exclusive (intent to write) request to W-word is a write-write conflict

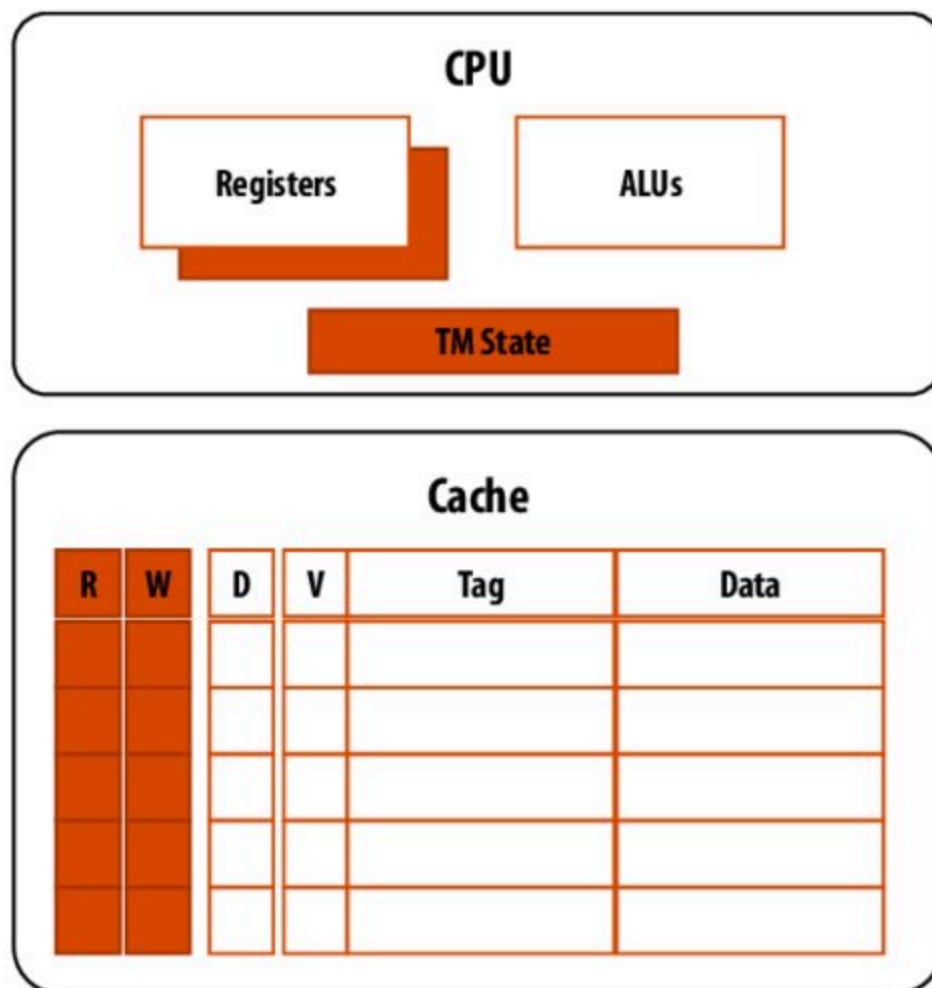
Example HTM implementation: lazy-optimistic



CPU changes

- Ability to checkpoint register state (available in many CPUs)
- TM state registers (status, pointers to abort handlers, ...)

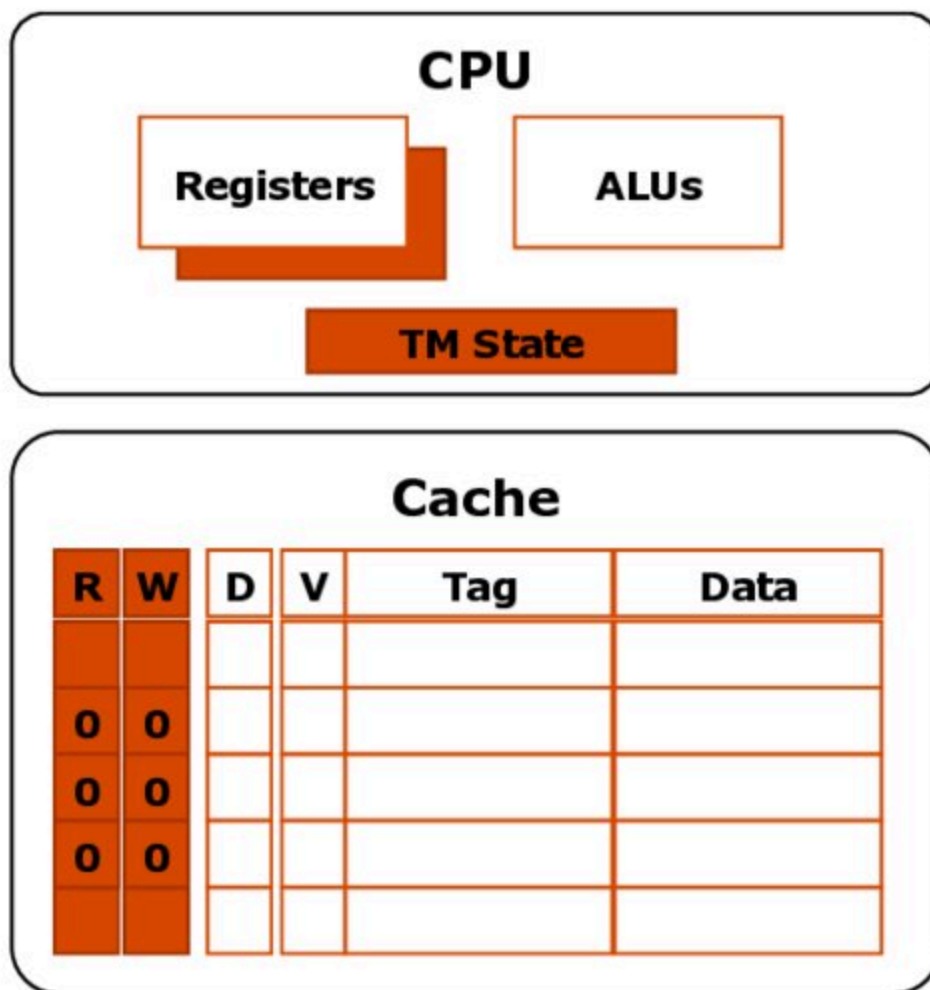
Example HTM implementation: lazy-optimistic



Cache changes

- R bit indicates membership to read set
- W bit indicates membership to write set

HTM transaction execution

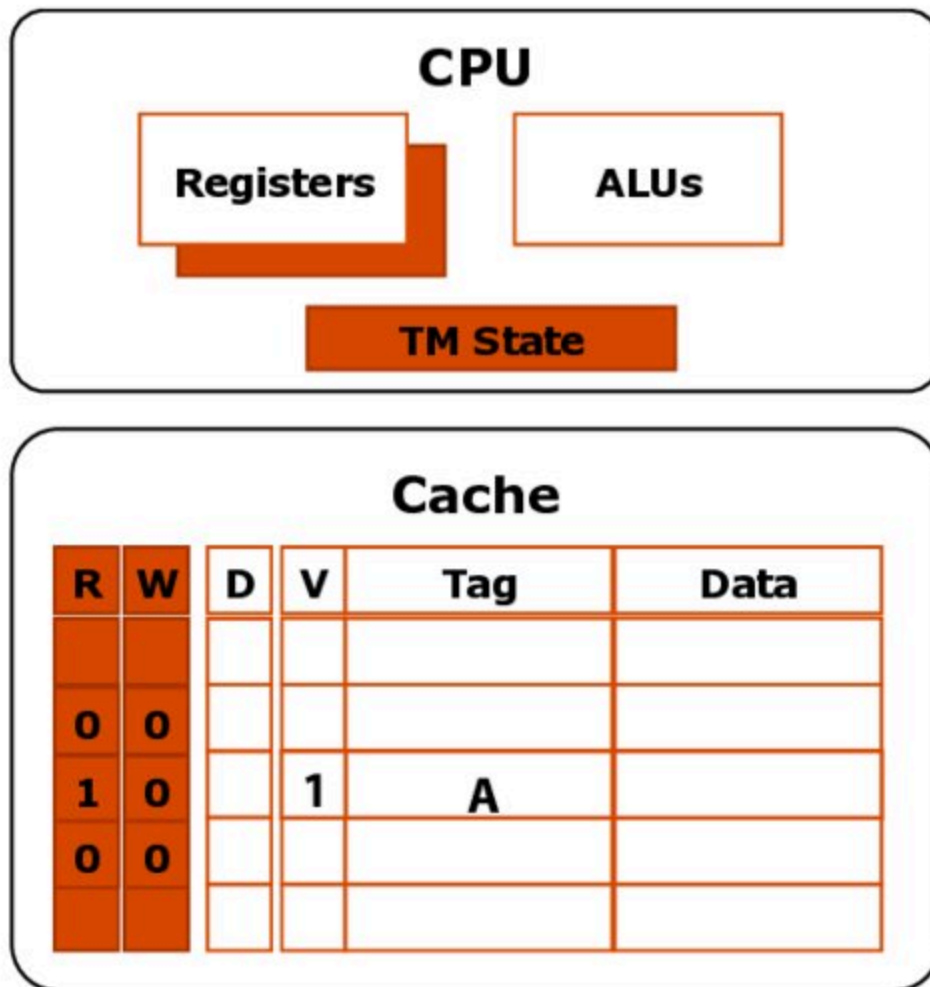


Xbegin ←
Load A
Load B
Store C ← 5
Xcommit

Transaction begin

- Initialize CPU and cache state
- Take register checkpoint

HTM transaction execution



Load operation

- Serve cache miss if needed
- Mark data as part of read set

Xbegin

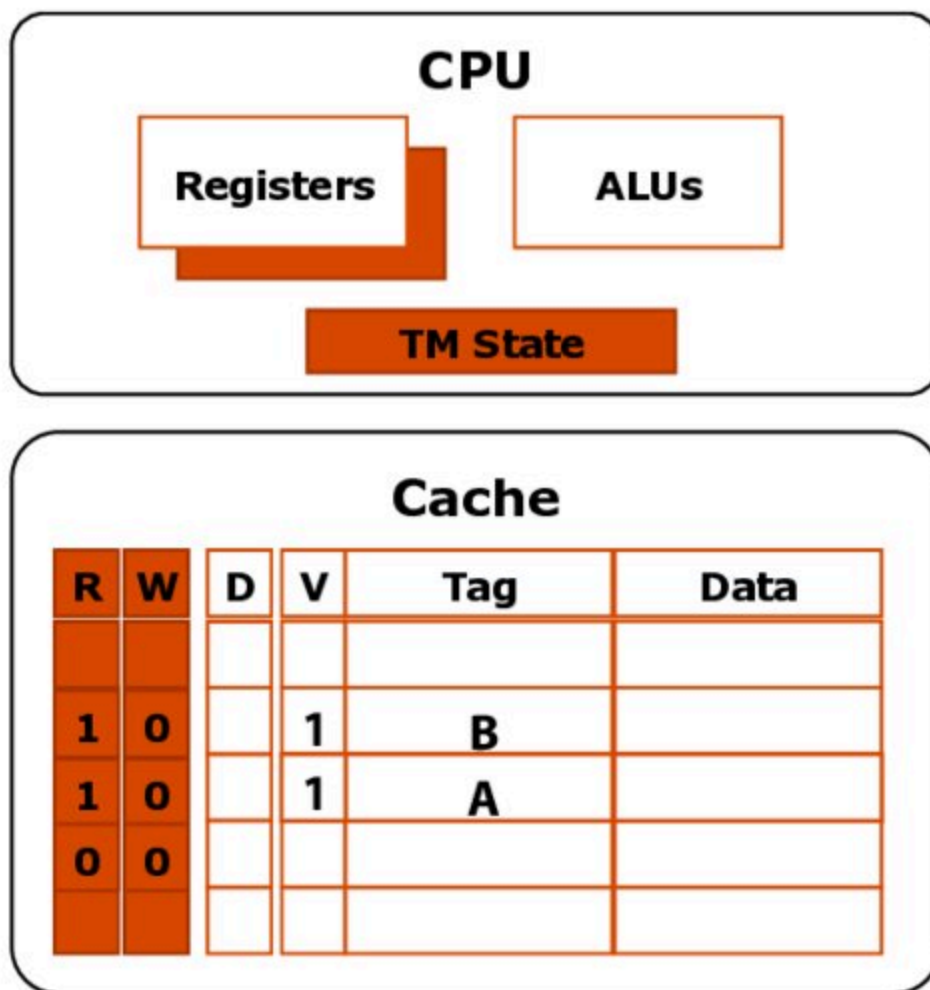
Load A ←

Load B

Store C ← 5

Xcommit

HTM transaction execution



Xbegin

Load A

Load B ←

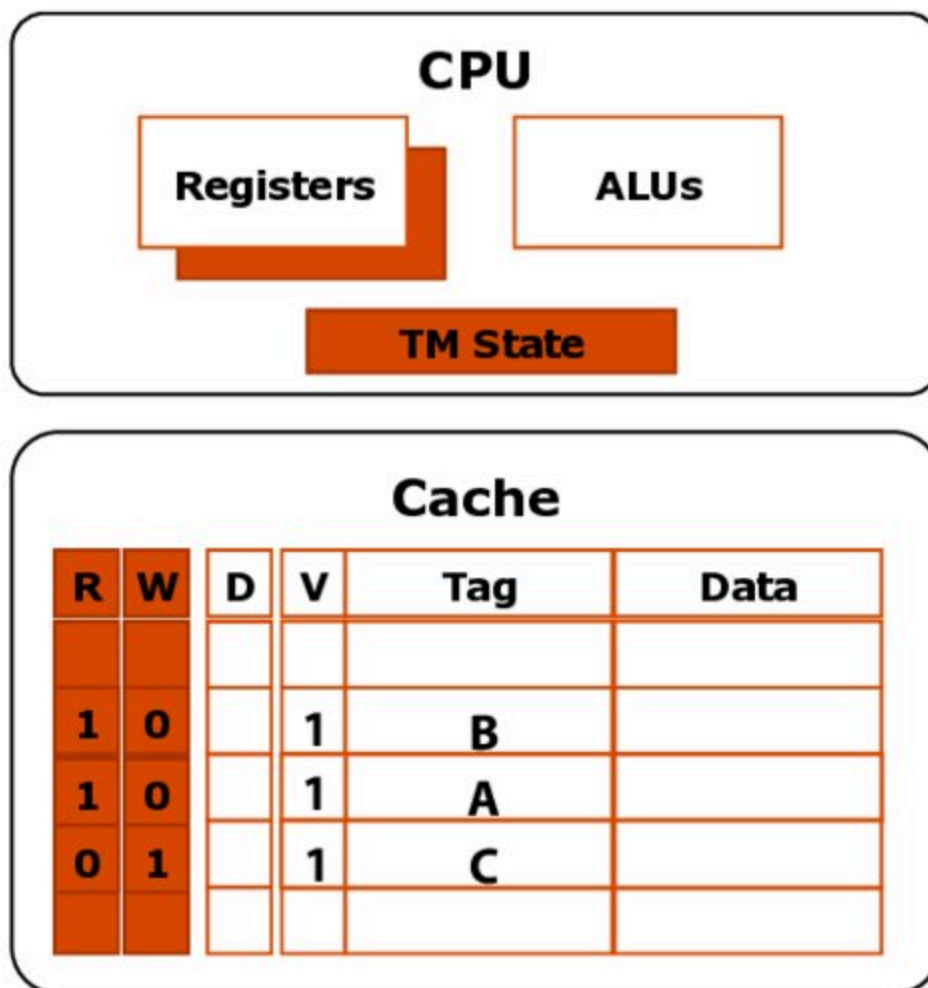
Store C ← 5

Xcommit

Load operation

- Serve cache miss if needed
- Mark data as part of read set

HTM transaction execution



Xbegin

Load A

Load B

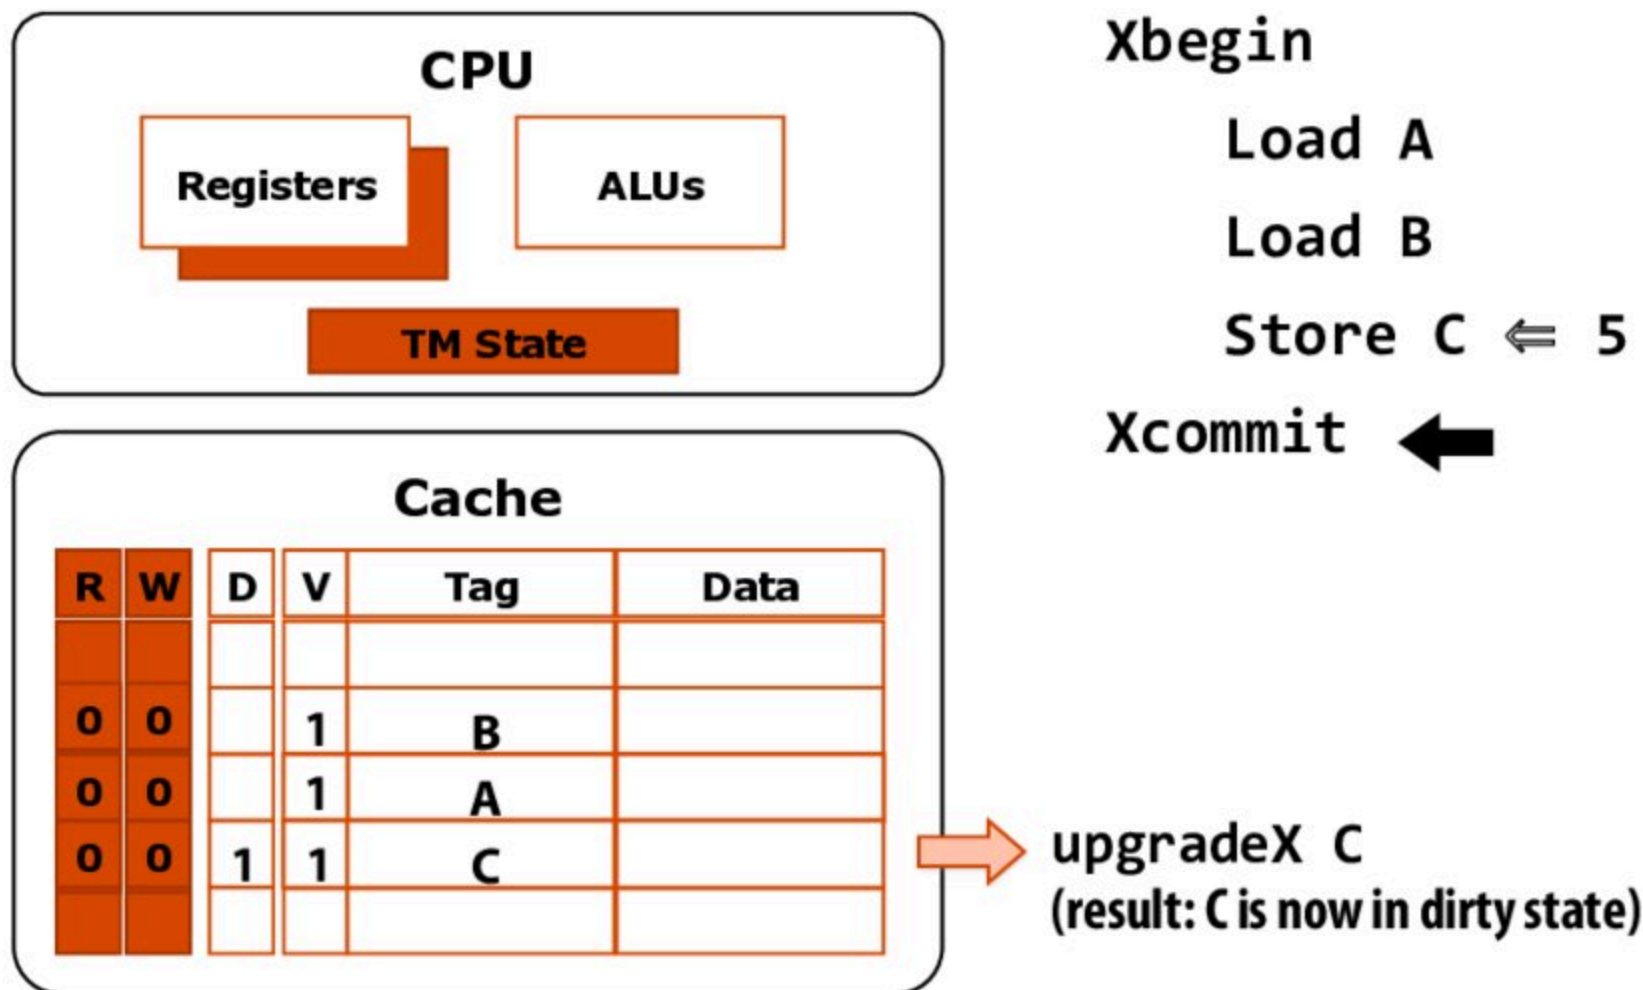
Store C \leftarrow 5 

Xcommit

Store operation

- Service cache miss if needed
- Mark data as part of write set (note: this is not a load into exclusive state. Why?)

HTM transaction execution: commit

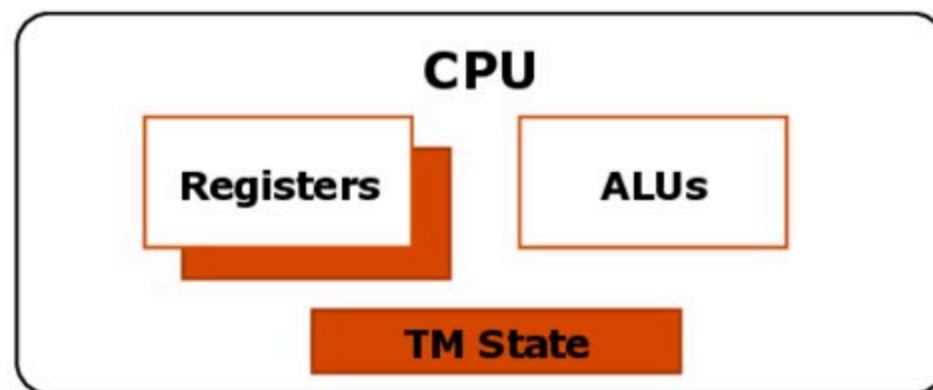


Fast two-phase commit

- **Validate:** request RdX access to write set lines (if needed)
- **Commit:** gang-reset R and W bits, turns write set data to valid (dirty) data

HTM transaction execution: detect/abort

Assume remote processor commits transaction with writes to A and D



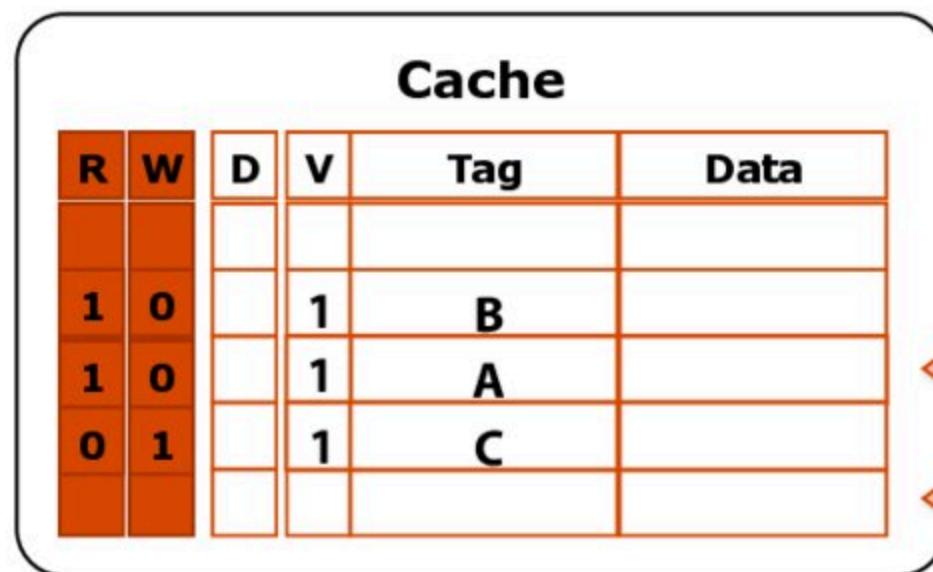
Xbegin

Load A

Load B

Store C \leftarrow 5 

Xcommit



 upgradeX A


 upgradeX D

coherence requests from
another core's commit
(remote core's write of A
conflicts with local read of A:
triggers abort of pending
local transaction)

Fast conflict detection and abort

- Check: lookup exclusive requests in the read set and write set
- Abort: invalidate write set, gang-reset R and W bits, restore to register checkpoint

HTM Performance Example



- 2x to 7x over STM performance
- Within 10% of sequential for one thread
- Scales efficiently with number of processors

Review: Transactional Memory

Atomic construct: declaration that atomic behavior must be preserved by the system

- Motivating idea: increase simplicity of synchronization without (significantly) sacrificing performance

Transactional memory implementation

- Many variants have been proposed: SW, HW, SW+HW
- Implementations differ in:
 - Data versioning policy (eager vs. lazy)
 - Conflict detection policy (pessimistic vs. optimistic)
 - Detection granularity (object, word, cache line)

Software TM systems (STM)

- Compiler adds code for versioning & conflict detection
 - Note: STM barrier = instrumentation code (e.g. StmRead, StmWrite)
- Basic data-structures
 - Transactional descriptor per thread (status, rd/wr set, ...)
 - Transactional record per data (locked/version)

Hardware Transactional Memory (HTM)

- Versioned data is kept in caches
- Conflict detection mechanisms augment coherence protocol

HTM Example: Transactional Coherence and Consistency

Use TM as the coherence mechanism → all transactions all the time

Successful transaction commits update memory and all caches in the system

P1	P2	P3
Begin T1 Read A Write A, 1 Write C, 2 Read D Commit T1	Begin T2 Read A Write E, 3 Commit T2 Begin T3 Write C, 4 Read A Write E, 5 Commit T3	Begin T4 Read E Write B, 6 Write C, 7 Read F Commit T4

Assumptions

- Lazy and optimistic
- One “commit” per execution step across all processors
- When one transaction causes another transaction to abort and re-execute, assume that the transaction “commit” of one transaction can overlap with the “begin” of the re-executing transaction
- Minimize the number of execution steps

HTM Example: Transactional Coherence and Consistency

P1	P2	P3
Begin T1 Read A Write A, 1 Write C, 2 Read D Commit T1	Begin T2 Read A Write E, 3 Commit T2 Begin T3 Write C, 4 Read A Write E, 5 Commit T3	Begin T4 Read E Write B, 6 Write C, 7 Read F Commit T4

P1			P2			P3		
Action	Read set	Write set	Action	Read set	Write set	Action	Read set	Write set
B T1			B T2			B T4		
R A	A:0		R A	A:0		R E	E:0	
W A, 1	A:0	A:1	W E	A:0	E:3	W B, 6	E:0	B:6
W C, 2	A:0	A:1,C:2	C T2	A:0	E:3	B T4		

HTM Example: Transactional Coherence and Consistency

P1	P2	P3
Begin T1 Read A Write A, 1 Write C, 2 Read D Commit T1	Begin T2 Read A Write E, 3 Commit T2 Begin T3 Write C, 4 Read A Write E, 5 Commit T3	Begin T4 Read E Write B, 6 Write C, 7 Read F Commit T4

P1			P2			P3		
Action	Read set	Write set	Action	Read set	Write set	Action	Read set	Write set
B T1			B T2			B T4		
R A	A:0		R A	A:0		R E	E:0	
W A, 1	A:0	A:1	W E	A:0	E:3	W B, 6	E:0	B:6
W C, 2	A:0	A:1,C:2	C T2	A:0	E:3	B T4		
R D	A:0,D:0	A:1,C:2	B T3			R E	E:3	
C T1	A:0,D:0	A:1,C:2	W C, 5		C:5	W B, 6	E:3	B:6

HTM Example: Transactional Coherence and Consistency

P1	P2	P3
Begin T1 Read A Write A, 1 Write C, 2 Read D Commit T1	Begin T2 Read A Write E, 3 Commit T2 Begin T3 Write C, 4 Read A Write E, 5 Commit T3	Begin T4 Read E Write B, 6 Write C, 7 Read F Commit T4

P1			P2			P3		
Action	Read set	Write set	Action	Read set	Write set	Action	Read set	Write set
B T1			B T2			B T4		
R A	A:0		R A	A:0		R E	E:0	
W A, 1	A:0	A:1	W E	A:0	E:3	W B, 6	E:0	B:6
W C, 2	A:0	A:1,C:2	C T2	A:0	E:3	B T4		
R D	A:0,D:0	A:1,C:2	B T3			R E	E:3	
C T1	A:0,D:0	A:1,C:2	W C, 5		C:4	W B, 6	E:3	B:6
			R A	A:1	C:5	W C, 7	E:3	B:6,C:7
			W E, 6	A:1	C:5,E:6	R F	E:3,F:0	B:6,C:7
				A:1	C:5,E:6	C T4	E:3,F:0	B:6,C:7

HTM Example: Transactional Coherence and Consistency

P1	P2	P3
Begin T1 Read A Write A, 1 Write C, 2 Read D Commit T1	Begin T2 Read A Write E, 3 Commit T2 Begin T3 Write C, 4 Read A Write E, 5 Commit T3	Begin T4 Read E Write B, 6 Write C, 7 Read F Commit T4

P1			P2			P3		
Action	Read set	Write set	Action	Read set	Write set	Action	Read set	Write set
B T1			B T2			B T4		
R A	A:0		R A	A:0		R E	E:0	
W A, 1	A:0	A:1	W E	A:0	E:3	W B, 6	E:0	B:6
W C, 2	A:0	A:1,C:2	C T2	A:0	E:3	B T4		
R D	A:0,D:0	A:1,C:2	B T3			R E	E:3	
C T1	A:0,D:0	A:1,C:2	W C, 5		C:5	W B, 6	E:3	B:6
			R A	A:1	C:5	W C, 7	E:3	B:6,C:7
			W E, 6	A:1	C:5,E:6	R F	E:3,F:0	B:6,C:7
				A:1	C:5,E:6	C T4	E:3,F:0	B:6,C:7
			C T3	A:1	C:5,E:6			

Hardware transactional memory support in Intel Haswell architecture

New instructions for “restricted transactional memory” (RTM)

- **xbegin**: takes pointer to “fallback address” in case of abort
 - e.g., fallback to code-path with a spin-lock
- **xend**
- **xabort**
- Implementation: tracks read and write set in L1 cache

Processor makes sure all memory operations commit atomically

- But processor may automatically abort transaction for many reasons (e.g., eviction of line in read or write set will cause a transaction abort)
 - Implementation does not guarantee progress (see fallback address)
- Intel optimization guide (ch 12) gives guidelines for increasing probability that transactions will not abort

Summary: transactional memory

Atomic construct: declaration that atomic behavior must be preserved by the system

- Motivating idea: increase simplicity of synchronization without (significantly) sacrificing performance

Transactional memory implementation

- Many variants have been proposed: SW, HW, SW+HW
- Implementations differ in:
 - Versioning policy (eager vs. lazy)
 - Conflict detection policy (pessimistic vs. optimistic)
 - Detection granularity (object, word, cache line)

Software TM systems

- Compiler adds code for versioning & conflict detection
 - Note: STM barrier = instrumentation code
- Basic data-structures
 - Transactional descriptor per thread (status, rd/wr set, ...)
 - Transactional record per data (locked/version)

Hardware transactional memory

- Versioned data is kept in caches
- Conflict detection mechanisms built upon coherence protocol